

Embedded Systems

Microprocessors

Ver 2.0, 2021

Prof. William Fornaciari

Politecnico di Milano – DEIB

william.fornaciari@polimi.it

Outline

- Introduction
 - Microprocessor general purpose
 - CISC, RISC, Superscalar
 - CISC/RISC, EPIC/VLIW, CISC/VLIW
 - Comparison
 - Dedicated Microprocessors
 - Digital Signal Processors
 - Network Processors
 - Microcontrollers
 - Conclusions

Microprocessors

Introduction

Introduction

A Microprocessor is an electronic circuit that elaborates several instructions

- Part of the embedded applications implement in sw their algorithms
 - Flexibility, TTM, easy upgrade, sometime cost (dep on volume)
- (During design process)
- Flexibility has two main flavors
 - **Maintainability and evolution of the application**
 - Sw development is less complex
 - Faster
 - Verification is less critical (even not trivial)
 - Sw designers are available in «volume»

Not only a blue sky

- Typically an equivalent hw solution has better performance than sw
 - Microprocessor architecture is optimized for flexibility but it is generic, not specialized
 - Performance is not only computing speed
 - Energy/power, Memory footprint, chip area + package, size, security, support by the suppliers, ...

A typical 1st step in designing a system is similar to "reversing a design" that already exists and improving it

Example 1/2

- Sum of two 32-bit integers
 - Compare hw and sw realization, considering the availability of a 32 bit core for the sw execution. Same clock speed of Hw and Sw
 - Software implementation
 - Just one assembly instruction
 - Hardware implementation
 - Need to design and implement the entire 32 bit adder
 - It seems that the sw solution is the favorite one...but

Example 2/2

- New requirements: move from 32 to 33 bits, after a more accurate analysis of the required accuracy
 - Hw solution is almost the same in terms of architecture and performance
 - Sw solution require to use two o presumably three assembly instructions, with a performance penalty of around 100%

Introduction

- The «extreme» example, figures out the differences between a dedicated hw solution and a more flexible sw implementation
 - Designing a sw solution requires a deep knowledge of the architecture of the microprocessors available on the market
 - The characteristics of the problems (algorithm) should drive the designer, beside non functional constraints (performance, cost, development time, power, ...)
 - Before to select a specific processor, the initial step is the choice of the «class» of processor to be used and in which «form» it has to be purchased

Classes of a processor

- Specialized vs generic
 - Specialized: *Application Specific Processor (ASP)*
 - Tailored for specific classes of applications like signal processing, image processing or the management of network data packets
 - Generic: *General Purpose Processor (GPP)*
 - 4-season processor, not optimized for any specific applications. Versatile, it can be used for different type of applications
 - Embedded Systems multi-processors
 - Typical solution is to use GPP for supervision and control of the activities of one or more ASPs Typical for designing mobile applications

In which form is available?

- COTS vs IP

- COTS (Component off the shelf) – just buy a chip off-the-shelf and mount in on a board (PCB) together with the interfaces to the rest of the systems *Easiest way to create a system*
- IP (intellectual property) – the design (description) of the micro is purchased. Several abstraction levels *full description using softwares "Verilog-like"*
 - Soft-macro: the micro is described in HDL at RT-level
 - Hard-macro: description down to the level of layout *Disclosed with lots of designs impossible to modify; they can be integrated with something else*
- Give the diffusion of PLD, the used of IP is becoming more and more popular. *Lots of famous companies (ARM, Samsung, NXT...) adopt this solution*
- Typically the suppliers of CPLD and/or FPGA make available one or more cores (even for free)
 - Micro/Pico Blaze di Xilinx, Nios-II di Altera, Mico32 di Lattice, ARM7 di Arm Ltd. e Actel

Selection process

- Class
 - The main driver is the nature of the algorithm, the operations and data to be elaborated
- Form
 - The main driver is the target architecture of the systems
 - SoC, MPSoC, NoC, PCB, ... The choice can change during the development and availability
 - But performance and power, in addition to royalties, can make the difference

Selection process

- Performance
 - Main metric is the average number of instructions per clock cycle (IPC/CPI) *Average can be misleading: our goal is to look to the peak*
 - It is relative to the clock frequency, hence it needs to be scaled to compare different architectures
 - MIPS- absolute measure of the throughput. It can be misleading when processors have different instruction sets (ISA)
 - MFLOPS is the same but for Floating Point operations
 - For more specific architectures, like DSP or NP, metrics can change
 - DSP, frequently it is used MMACS
 - NP, average number of processed packets in a time unit

Selection process

- Power is probably the most important driver for ES
 - Average power or peak power
 - Useful to roughly estimate the max power or the avg power of the overall systems

Average useful to have an idea regarding durability

Peak useful to know the strenght

- Frequently, especially at the beginning, it is used a joint measure between power and speed (related to performance)
 - For example mW/MHz or, at a higher level of abstraction, MIPS/mW
 - These information are very useful but must be properly used during the design process

Memory

Frequently underestimate, but it's a crucial topic. In Embedded Systems, we generally evaluate memory in terms of Kb. We use single chip apps, with a specific size

- Memory requirements (bandwidth and size) of the application can be a hard constraint
- In some case it could be impossible to use an external memory, only that internal to the micro is available
- Maybe that some bandwidth can be achieved only by integrated memories
- For complex systems, the crucial point could be the addressing space
 - Simple architectures exploits 16-bit addressing (64KB), 20 bits (1MB) or 24 bits (16MB), while only those more complex can arrive to 32 bits (4GB)

Peripherals

Something "hidden". Generally focus is on executing a code, but in Embedded Systems we have to catch signals from the external environment. That's why peripherals are needed

- An embedded systems typ elaborates external signals and control physical apparatus
- It is possible to select and architecture oriented toward the computing, and to design the rest of the system, or to look for a single chip solution integrating computing, interfacing and some peripherals
 - Microprocessor vs. Microcontrollers
 - The choice of the micro require a PCB or a SoC integrating all the peripherals under the responsibility of the designer
 - Using a microcontroller, the designer get rid of the integration problem, paid in terms of reduction in the number of alternatives

Available Software is crucial

- Many embedded applications have a limited legacy code, most of them exploit standard functions and libraries
 - E.g. numerical functions (FFT, Wavelet, ...), Image processing (convolutional filters, motion detection, DCT, ...)
- The availability of libraries simplifies both design and validation of the applications, making feasible in some cases the development of sw solutions otherwise impractical
- The availability of an operating systems and SDK for the specific processor is important, too

Software Development

- It is the cornerstone of the sw development, many differences exist among the various SDKs
 - Important aspects to be evaluated are the software compilation flow, the quality of the code generated and the flexibility and knobs offered to the designer
 - Concerning the analysis tools, the richness, their accuracy and reliability are important. Debuggers should be carefully analyzed
 - Documentation and reference design
 - Community end ecosystem

Something more tangible

- Packaging
 - If the micro is a COTS, different packages are available
 - Size, pinout, material (e.g. plastic, metal), certification (consumer, industrial, aerospace, ...)
- Certifications
 - There exists several certification
 - Consumer, industrial, aerospace, automotive, military, ...
 - Some are specifically tailored for narrow fields
 - » *Automotive: MISRA rules* (<http://www.misra.org.uk/>)
 - The availability of a component with a proper certification could be a must constraint for its adoption

Microprocessors

General Purpose Processors
GPP

GPP

- GPP - generic architecture valid for a range of application fields *Slow interface with the environment (use of sensors)*
 - Common in PC/PDA
 - For ES low end when energy and performance are not crucial and the nature of the application is quite heterogeneous
 - Control and management of a slow interfacing with sensors, interactivity via GUI
- GPPs can have different architectures
 - CISC vs RISC
 - CISC: huge amount of complex instructions
 - RISC: few simple instructions *prepared for Embedded Systems*

GPP

CISC

GPP

- CISC
 - The first GPP were designed by following the approach of a *Complex Instruction Set Computer*
 - Limited amount of memory at that time
 - Need to obtain compact code, with few lines
 - ISA and architectures designed to carry out three operations (*load*, *execute* e *store*) in a single instruction
 - The addressing mode has been enriched
 - Development of the compiler simplified

GPP

- Architetture CISC
 - Addressing modes

Modalità	Esempio	Significato
Assoluto (istruzioni)	<code>jmp a</code>	$PC \leftarrow a$
Relativo (istruzioni)	<code>jmp (a)</code>	$PC \leftarrow PC + a$
Indiretto (istruzioni)	<code>jmp Rs</code>	$PC \leftarrow PC + Rs$
Registro	<code>add Rd, R1, R2</code>	$Rd \leftarrow R1 + R2$
Assoluto	<code>load Rd, a</code>	$Rd \leftarrow M(a)$
Indiretto	<code>load Rd, [Rs]</code>	$Rd \leftarrow M(Rs)$
Base, offset	<code>load Rd, a[Rs]</code>	$Rd \leftarrow M(Rs + a)$
Indice	<code>load Rd, [Rs, Ri]</code>	$Rd \leftarrow M(Rs + Ri)$
Indice, offset	<code>load Rd, a[Rs, Ri]</code>	$Rd \leftarrow M(a + Rs + Ri)$
Indice, offset, scala	<code>load Rd, a[Rs, Ri], s</code>	$Rd \leftarrow M(s \times (a + Rs + Ri))$
Autoincremento	<code>load Rd, +[Rs]</code>	$Rd \leftarrow M(Rs), Rs \leftarrow Rs + 1$
Autodecremento	<code>load Rd, -[Rs]</code>	$Rd \leftarrow M(Rs), Rs \leftarrow Rs - 1$
Push	<code>push Rs</code>	$M(SP) \leftarrow Rs, SP \leftarrow SP + 1$
Push	<code>pop Rd</code>	$SP \leftarrow SP - 1, Rd \leftarrow M(SP)$

GPP

- CISC
 - Ideally, each Arith/Logic operation should be capable to access (read) data and write the result by exploiting any of the available addressing modes
 - Fully *orthogonal* architecture and Instruction Set
 - In the reality, the number of operation-addressing mode combinations has been limited
 - Usually one of the operand refers to the memory, while the other is a register with the role of both source and destination
 - Despite this simplification, CISC instructions are coded in variable length formats
 - More complex fetch and decode units and need of more speed when accessing the memory

GPP

- CISC Architecture
 - The ALU, is implementing a variety of the basic operations
 - Add/Sub, mul, div, logic operations, shift, comparison
 - In some architectures the instructions can be vectorial, working on more registers at the same time
 - *Single Instruction Multiple Data*
 - So many addressing mode and A/L operations, require a complex *data-path* that is hard to make fast

GPP

- CISC Architectures
 - Some flavors of the ADD operation

Istruzione Significato	
add	Istruzione di base
adc	Somma con riporto in ingresso
addi	Somma con una costante
addi	Somma con una costante a 32 bit
addq	Somma con una costante a 3 bit
aaa	Somma di operandi unpacked BCD
daa	Somma di operandi packed BCD
xadd	Somma e scambia gli operandi
addps	Somma di operandi a 128 bit
addss	Somma dei 32 LSB di operandi a 128 bit
paddb	Somma i singoli byte
padds	Somma i singoli byte con saturazione
adawi	Somma operandi allineati alla parola
addb3	Somma tre operandi a 8 bit
addw3	Somma tre operandi a 32 bit
addl3	Somma tre operandi a 64 bit

GPP - CISC

- To increase the *throughput*, the CISC processors implement complex pipeline (e.g. up to 23 stages) that are hard to manage
- Another strategy, still very costly in terms of hw resources, is based on altering the order of execution of the assembly instruction without modifying the semantics of the program
 - For a compiler, when considering so complex pipelines, it is hard to control the status of the execution of the instruction
 - The instructions could arrive to the decoding in an order that is not optimal w.r.t. the status of the pipelines at that time
 - Modern CISCs have dedicated hw units to dynamically reschedule the instructions
 - Out-of-order execution
 - CISC have remarkable performance but, over the time, energy efficiency and cost moved in favor of other architectures

GPP

RISC

GPP

- RISC - *Restricted Instruction Set Computer*, when the architecture has few simple instructions
 - End of '80 the price of DRAMs fall down
 - Size of the program is no longer a wall
 - » CISC operations can be decomposed into RISC instructions to simplify the architecture
 - In the same period the integration scale raised up, dramatically
 - Possibility to realize on a single chip complex architectures
 - Pipeline and quantity of GP registers
 - To exploit pipelining, the execution time of the instructions must be uniform
 - This is possible when instructions are simple

GPP

- RISC
 - Main advantage is simple instructions set and simple architecture
 - Instruction decoding is simplified
 - Frequently RISC instruction have fixed lengths
 - Easy balancing the stages of the pipeline and to increase clock
 - Execution units (ALU and BPU) take advantage from simple inst
 - RISC operates only on registers hence more speed and duration is predictable
 - RISC processor tends to have many registers, so the presence of memory spill is reduced significantly

GPP

- RISC
 - I/O, namely load/store from memory/registers and viceversa
 - Each instruction requires to pre-load of the data and, eventually, to write the result in memory
 - **load/store architecture**
 - Data transfer operation are crucial for performance
 - Fortunately, the instructions have a limited number and simple of addressing modes
 - Aggressive use of memory (cache) hierarchy to reduce the access time to data

GPP

- RISC
 - Addressing modes

Modalità	Esempio	Significato
Relativo (istruzioni)	<code>jmp (a)</code>	$PC \leftarrow PC + a$
Indiretto (istruzioni)	<code>jmp Rs</code>	$PC \leftarrow PC + Rs$
Registro	<code>add Rd, R1, R2</code>	$Rd \leftarrow R1 + R2$
Base, offset	<code>load Rd, a[Rs]</code>	$Rd \leftarrow M(Rs + a)$
Autoincremento	<code>load Rd, +[Rs]</code>	$Rd \leftarrow M(Rs), Rs \leftarrow Rs + 1$
Autodecremento	<code>load Rd, -[Rs]</code>	$Rd \leftarrow M(Rs), Rs \leftarrow Rs - 1$
Multiplo	<code>load {R1, R2, ...}, [Rs]</code>	$R1 \leftarrow M(Rs), R2 \leftarrow M(Rs+1), \dots$

GPP - RISC

- Software development flow
 - Considering the same high level source code, the number of RISC instructions will be higher than CISC
 - RISC architectures have no hw units to solve pipeline conflicts
 - But good predictability of execution allows the compiler to generate efficient code
- Power consumption
 - RISC are simple and inherently low power w.r.t. CISC
 - 20-100 mW/MHz or around mW/MIPS

GPP

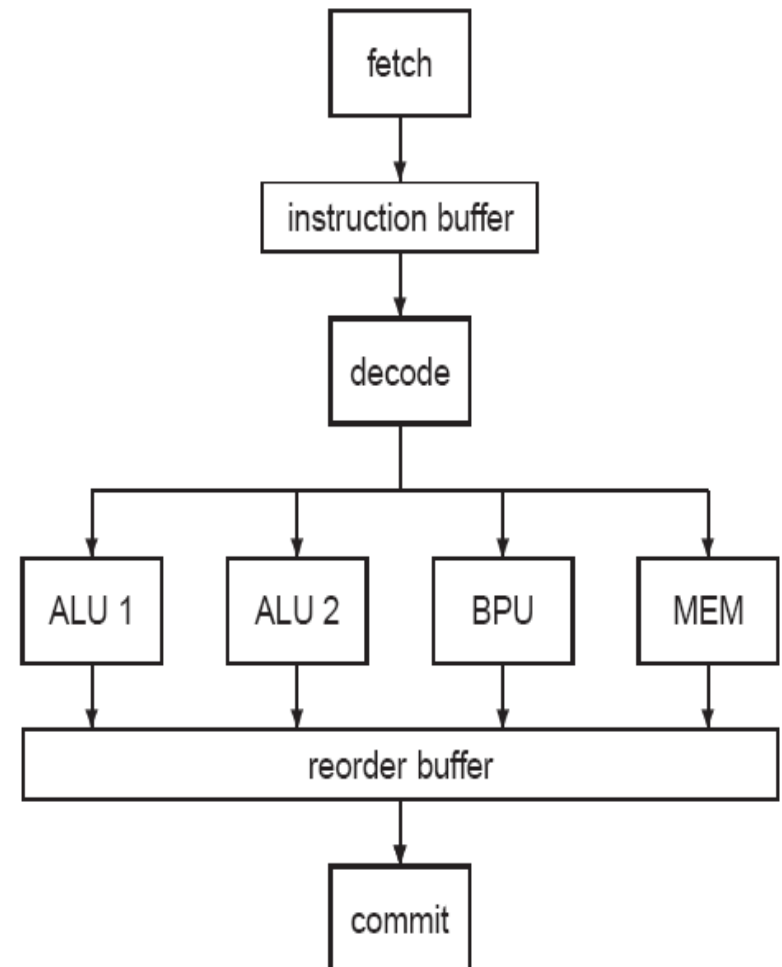
Superscalarity

GPP

- Superscalar architectures
 - There exists a lot of parallelism at the level of assembly instruction
 - Example: $y=a+b$ and $z=c*d$
 - There is no dependency, hence any execution order is fine or they can execute in parallel
 - This consideration spurred the appearing of the superscalar architectures
 - It is superscalar an architecture having more than one executing units, hence more ALUs and or BPU

GPP

Typical super scalar architecture



GPP – super scalar

- The parallelism offered by the execute stage allows to deliver more instruction per clock cycle
 - The complexity of the control logic increases as the impact on clock speed and power consumption
 - Moreover, it is the microprocessor to define dynamically the scheduling of the instructions
 - Pros
 - Compiler has not to work on organizing the code for a specific hardware structure
 - Cons
 - A relevant part of the processor is deserved to implement complex scheduling mechanisms and to manage the consistency

GPP

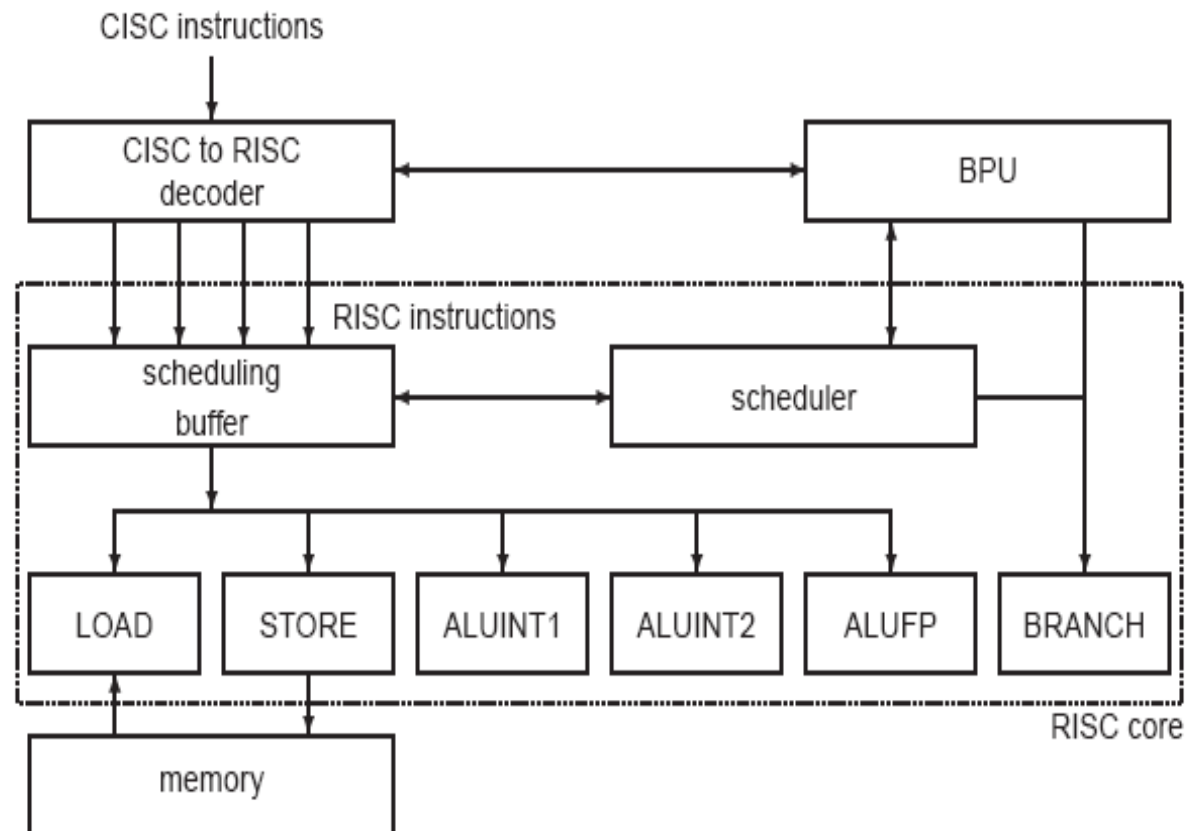
CISC/RISC

GPP

- CISC/RISC
 - Approach combining *out-of-order execution* and *superscalarity* with the goal to remain compatible with code developed in the past
 - Instruction set x86 is the standard de-facto for CISC
 - Each CISC instruction is decomposed in a set of simple instructions easy to be executed by a simple and efficient core
 - There is an hardware cost, but this solution joins the benefit of backward compatibility with the high-performance achievable by a RISC core

GPP

- CISC/RISC
 - Structure of a CISC based on a RISC core



GPP

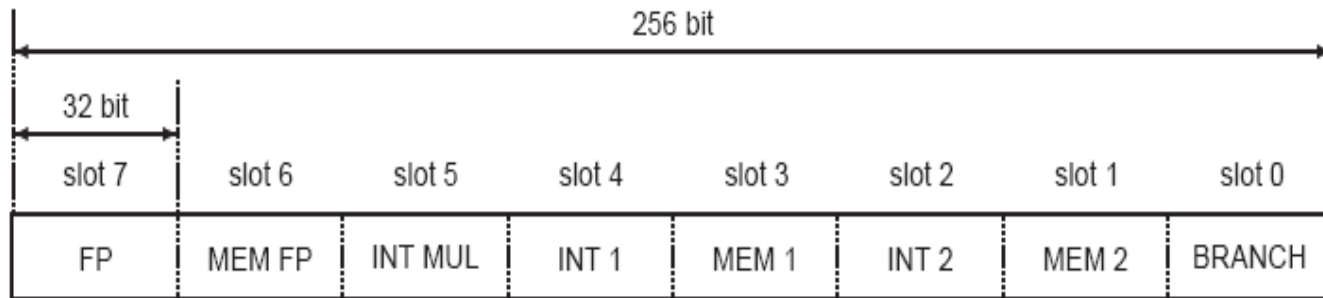
EPIC/VLIW

GPP

- EPIC/VLIW
 - To overcome the limits related to the complexity of the hw scheduling of the instructions, Intel developed in the past a new class of processors named EPIC (or VLIW). Itanium has been the first commercialized
 - These architectures support explicit parallelism
 - Possibility to execute more instructions at the same time under the control of the program
 - A VLIW instruction is the packing of more elementary instructions, typically RISC, in a single wide word, having a fixed structure

GPP

- EPIC/VLIW
 - Structure of a VLIW instruction



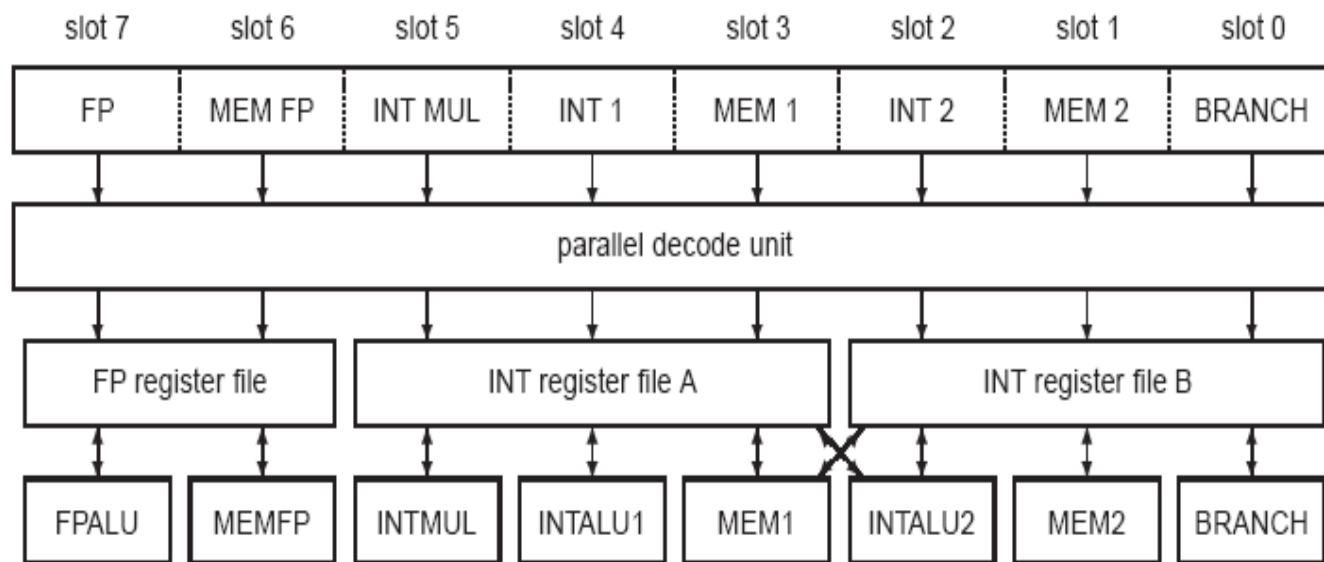
GPP

- EPIC/VLIW
 - The VLIW word, once fetched from the memory, is decoded in parallel
 - At the end of the decoding, the single RISC instructions corresponding to the various slots, are dispatched to the different execution units
 - This fixed structures moves to the compiler the task to perform a complex scheduling
 - The algorithms to be optimized should be well structure to fit well a specific architecture

GPP

- EPIC/VLIW

- To execute in parallel more instructions, the register files should be equipped with multiple read ports and should be connected flexibly to the different units



GPP

- EPIC/VIW
 - EPIC/VIW has the goal to move the complexity of the scheduling and code optimization towards the compiler
 - The architectures are simple, even if of large size given the number of execution units
 - Power consumption is in between simple and superscalar RISC
 - Similar solutions, allowing to achieve peak performance, has been adopted also in non general purpose architecture like the *digital signal processors (DSPs)*

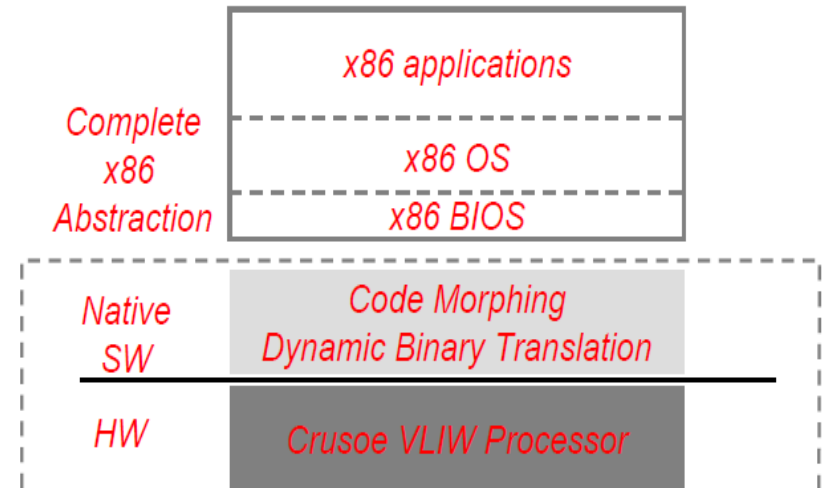
GPP

CISC/VLIW

GPP

- CISC/VLIW
 - Similar to CISC/RISC, performance of VLIW with the goal of backward compatibility
 - CISC instructions are decomposed in terms of basic ones execute by a VLIW (simple and efficient) core
 - Example (one one?)
 - Transmeta Crusoe
 - Today available only as IP

Transmeta Crusoe & Code Morphing



- Crusoe boots “Code Morpher” from ROM at power-up
- Crusoe+Code Morphing == x86 processor
x86 software (including BIOS) cannot tell the difference

GPP

Comparison

	RISC	CISC	Superscalari	EPIC/VLIW
Instruction set	semplice	complesso	complesso	semplice
Modi d'indirizzamento	pochi	molti	molti	pochi
Dimensione istruzioni	fissa	variabile	variabile	fissa, grande
Register file	singolo ⁽¹⁾	singolo ⁽¹⁾	singolo ⁽¹⁾	multiplo
Numero registri	alto	basso	medio	molto alto
Scheduling istruzioni	statico	dinamico	dinamico	statico
Gestione conflitti	statico ⁽²⁾	dinamico	dinamico	statico
Complessità compilatore	media	alta	bassa	molto alta
Parallelismo	assente	assente	implicito	esplicito
Complessità hardware	bassa	alta	molto alta	alta
Complessità pipeline	bassa	alta	molto alta	media ⁽³⁾
Consumo di potenza	molto basso	alto	molto alto	alto
IPC tipico	≈ 1	< 1	> 1	$\gg 1$

⁽¹⁾In alcuni casi disgiunto tra registri interi e registri floating-point.

⁽²⁾I conflitti vengono riconosciuti e risolti unicamente mediante stalli.

⁽³⁾Struttura semplice, ma di grandi dimensioni per via delle molte unità di elaborazione.

Microprocessors

Application Specific

ASP

- ES need to solve problem with high degree of specialization and with a limited number of functions
 - GPP are frequently not the best choice
- Specialized architectures have been designed
 - Differences in terms of instruction sets, architecture, size, power consumption, computing power, ...
 - Digital Signal Processor, Network Processors, Microcontrollers

ASP

Digital Signal Processor

ASP

- Digital Signal Processor
 - DSP are the popular one dedicated processors
 - They can vary a lot among the producers, but they have in common the design tailored to perform numerical computing
 - To understand their structure, let's review the main operations used in the algorithm for numeric processing

Operazione	Forma chiusa	Forma iterativa
Somma vettoriale	$\mathbf{Z} = \mathbf{X} + \mathbf{Y}$	$z_n = x_n + y_n \quad \forall n$
Riscaldamento	$\mathbf{Z} = k\mathbf{X}$	$z_n = k \cdot x_n \quad \forall n$
Somma	$z = \mathbf{X} _1$	$z = \sum_n x_n$
Somma quadratica	$z = \mathbf{X} _2$	$z = \sum_n x_n^2$
Prodotto scalare	$z = \mathbf{X} \times \mathbf{Y}$	$z = \sum_n x_n \cdot y_n$
Convoluzione	$\mathbf{Z} = \mathbf{X} * \mathbf{Y}$	$z_n = \sum_k x_k \cdot y_{N-k} \quad \forall n$
Trasformata di Fourier	$\mathbf{Z} = \mathcal{F}(\mathbf{X})$	$z_n = \sum_k x_k e^{\frac{-2\pi i}{N}nk} \quad \forall n$

ASP

- Digital Signal Processor
 - Widely used operation

$$z_{t+1} = z_t + x * y$$

- It is so important for DSP to worth a specific name
 - ***MAC : Multiply Accumulate***
- A valuable DSP architecture should optimized as much as possible the sum and mul operation, possibly by exploiting a single 3-operand instruction

ASP

- Digital Signal Processor
 - Given the nature of the operations, a relevant part of the algorithms to be executed, contains loops
 - The possibility to optimize loops in a DSP is possible, since
 - The body is typically small (~10 asm instructions)
 - The control loop variable is not altered within the loop body
 - The update of the control loop variable is simple
 - The pattern to access the vector is generally regular and simple
 - Specific Hw to optimize loops
 - Circular buffer to store the loop body
 - Dedicated register where to store the control loop variable
 - Add/Sub to perform inc/dec operation without using ALU

ASP

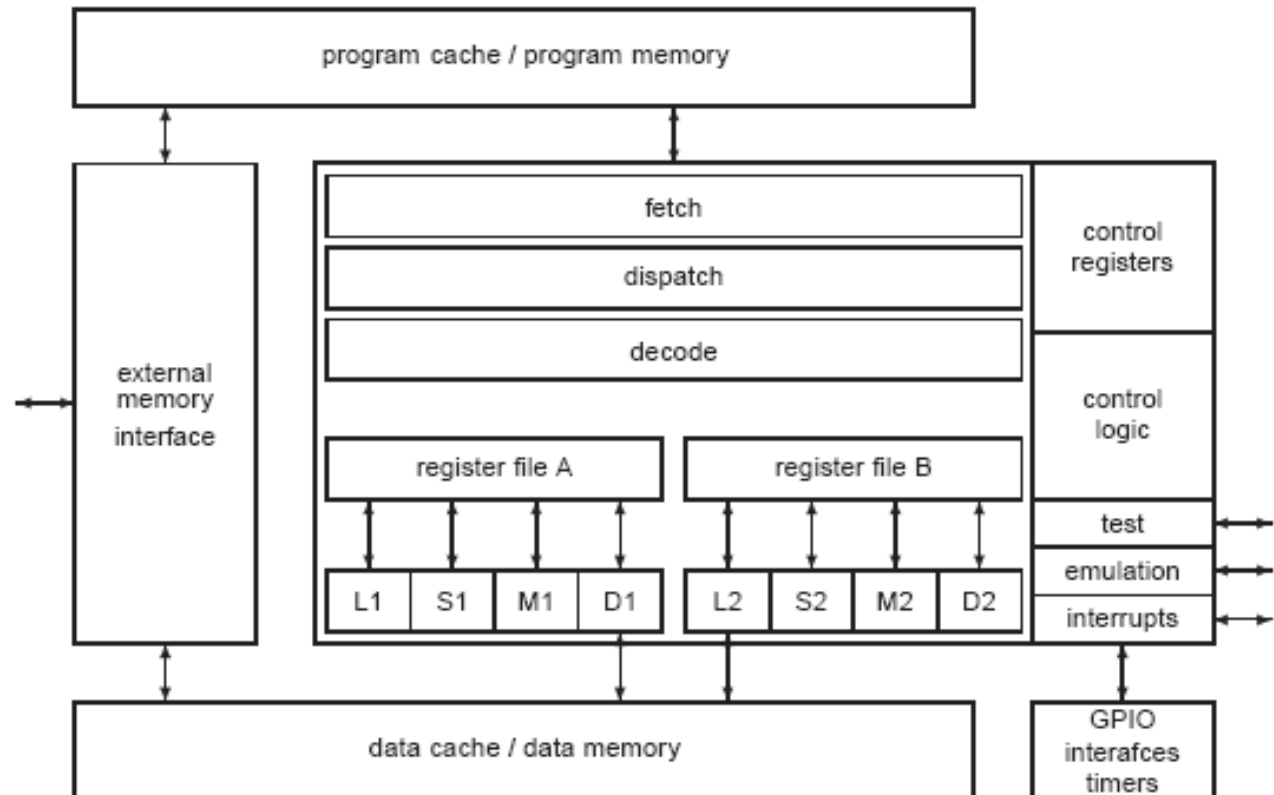
- Digital Signal Processor
 - Access to the data in memory could be a bottleneck
 - High bandwidth interfaces are designed to this purpose
 - Memory hierarchy for DSP
 - High speed bus with word width up to 256 bits
 - Multi-level cache hierarchy
 - Classic
 - » Stanford: unified for data and instructions
 - Specialised
 - » Harvard: cache L0 separate for data and instructions
 - » SHARC: Three L0 caches for data, instructions and constants

ASP

- Digital Signal Processor
 - In addition to the flexibility and computing power offered by the special instructions and by the specific ALUs, modern DSPs are frequently organized by following a VLIW approach
 - In fact, it is suitable for numerical applications, where the control is limited and there exists a high intrinsic parallelisms of operation

ASP

- Digital Signal Processor
 - Example: DSP VLIW by TEXAS



ASP

Network Processors

ASP

- Network Processor
 - Dedicated architectures tailored to process network packets
 - Complex system-on-chip with high parallelism, for network applications (e.g. routers)
 - Functionalities and protocols are evolving, the design of a NP is not trivial
 - Classes of the typical operations of a NP
 - Buffering of input/output packets
 - Remove/add of header at the level of data-link
 - Search for a specific field in an IP header
 - Search for an address in a lookup table
 - Computation of CRC codes
 - Send or delete of packets

ASP

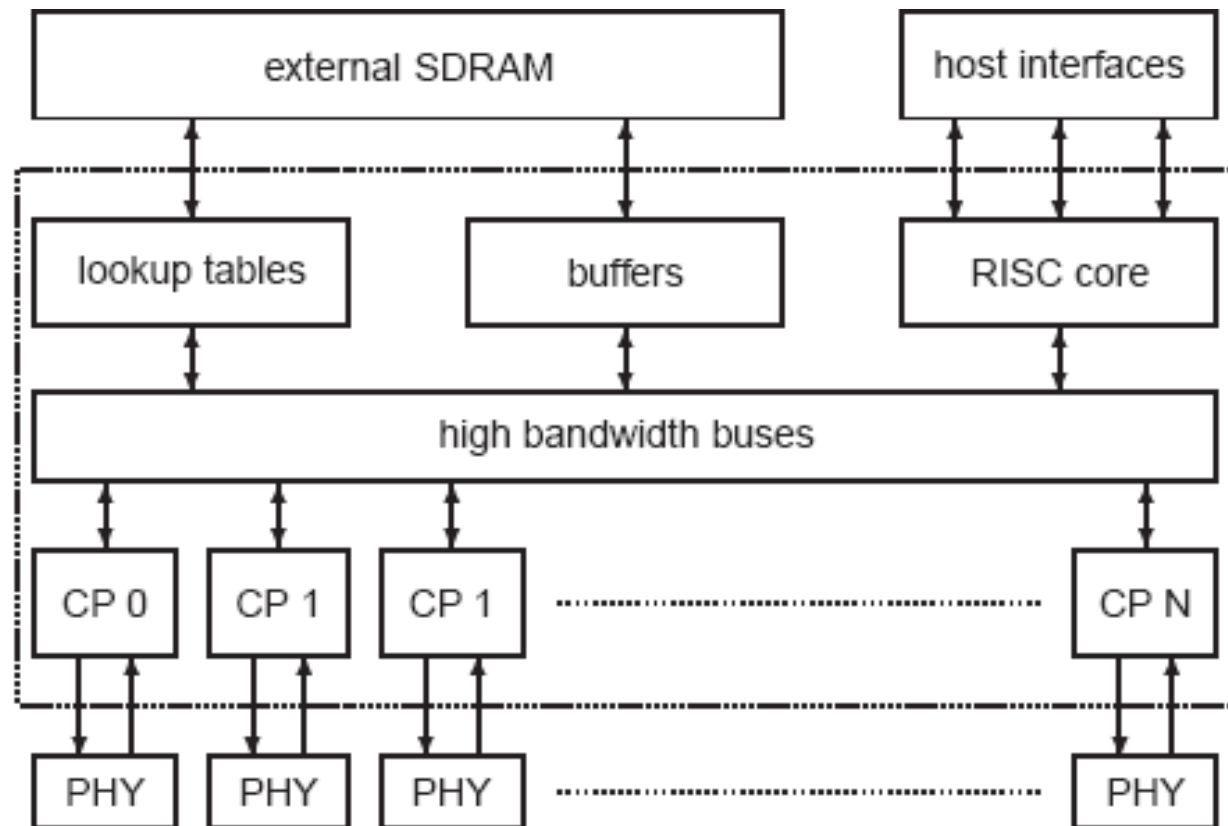
- Network Processor
 - To increase performance, dedicated hw is used
 - Fast I/O interfaces, queue management, crypto-cores, multi RISC cores
 - Frequently packet processing application requires the management of several independent channels who can benefit from parallel processing
 - Each channels is managed by a dedicated hw units frequently named as PP (*Packet Processor*) or CP (*Channel Processor*)

ASP

- Network Processor
 - When the separate flows and not fully disjoint
 - Use of RISC cores, to supervise and for high-level management
 - The typical structure of an application of *packet processing* is a set of simple routines, each working on a huge amount of packets
 - Small programs working on a enormous amount of data

ASP

- Network Processor
 - Typical architecture of a NP



ASP

- Network Processor
 - Which type of programming?
 - Core RISC
 - Development of programs starting from C source code like that of a RISC
 - Single PP
 - Sometime is still carried out in assembly, bunch of simple routines
 - » Programming CP is not well structured and only some time we can use a functional approach based on pattern recognition

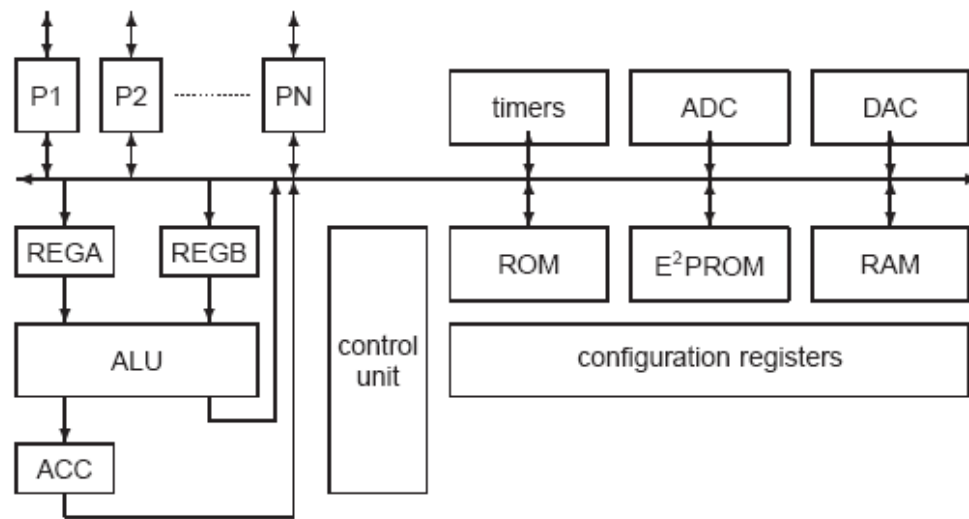
ASP

Microcontrollers

ASP

- Microcontrollers

- (MCU – *Micro Controller Unit*) is a class of microprocessor integrating peripherals and interfaces on a single-chip
- Their SoC nature is fine to deal with not top computing power, but with application having constraints on the use of hw resources and development time



ASP

- Microcontrollers
 - Frequently there is no interface to external memory
 - Programs for MCU are typically small
 - Memory interface increases the pinout
 - Small size is frequently a strict requirement
 - To provide a memory interface limiting the pinout, some architectures use a multiplexing of the I/O ports
 - The control logic must be properly configured to exploit this operating mode

ASP

- Microcontrollers
 - The architectures offer several peripheral and interfaces
 - I2C, SPI, CAN, JTAG, *PWM*, *UART/USART*, *watchdog*, *timer*, analog I/O and comparators, etc...
 - Programming is typically in C, only sometime in assembly
 - SDKs for microcontrollers ranges from simple C compiler up to integrated framework for performance analysis and configuration management