



POLITECNICO
MILANO 1863

DIPARTIMENTO DI ELETTRONICA,
INFORMAZIONE E BIOINGEGNERIA



ADVANCED OPERATING SYSTEMS

Prof. William Fornaciari

Task Scheduling

William Fornaciari

william.fornaciari@polimi.it

Outline

- Introduction
- Classification of scheduling algorithms
- Scheduling algorithms
- Priority-based scheduling
- Multiprocessor scheduling


Introduction

Basic concepts and terminology

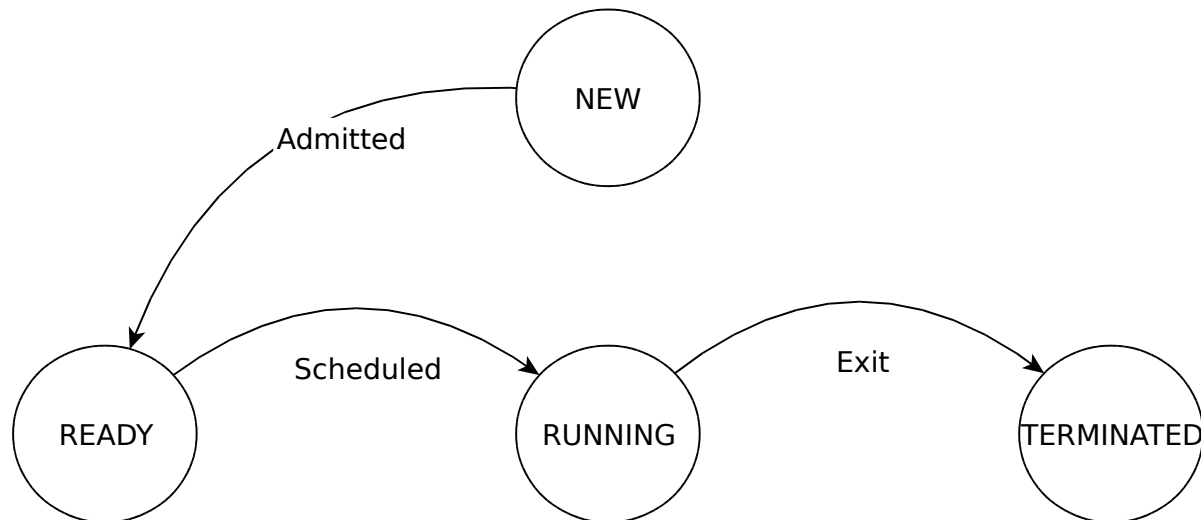
- A task is a computation performed by the processor in a sequential fashion. It is characterized by an execution status:
 - *Running* – the processor is executing the task
 - *Ready* – the task is waiting for the allocation of the processor
 - *Active* – used to refer to both running and ready tasks
- Ready tasks are kept in a ready queue by the operating system
- The scheduler is the OS component in charge of establishing the execution order of the tasks
 - Ordering algorithm for the ready queue (scheduling policy)
 - Scheduling policy invocation can be *periodic* or *event-based* (e.g., arrival or termination of tasks)
- Dispatching : allocating a (the) processor to a task

Introduction

Scheduling overview




Task state



Introduction

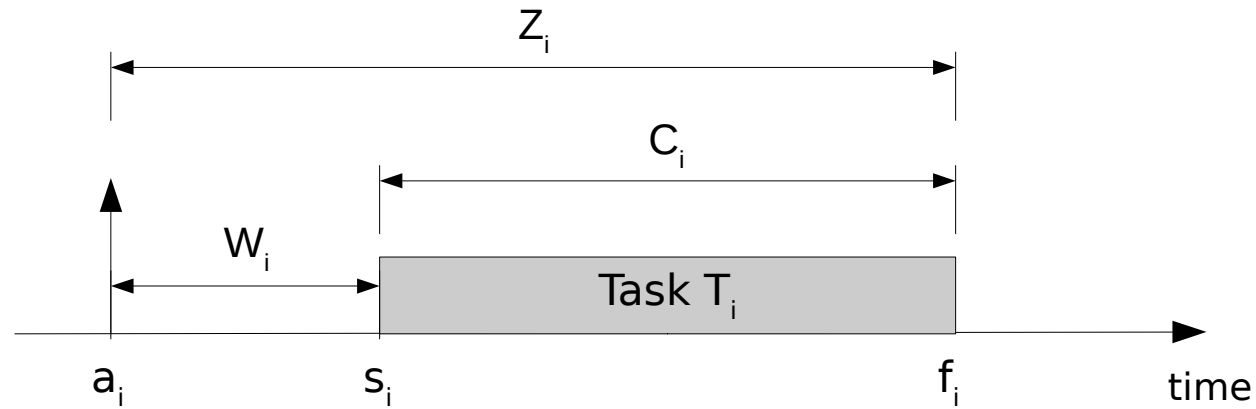
Task parameters



- a_i : Arrival time (or Request time)– time instant at which task is ready for execution (enter the ready queue)
- s_i : Start time – the time instant at which execution starts
- W_i : Waiting time – time spent in the ready queue before being scheduled
- f_i : Finishing time (or Completion time) – time instant at which the execution terminated

Introduction

Task parameters



- C_i : Computation time (or Burst time) – amount of time necessary to the processor to execute the task (without interruption)
- Z_i : Turnaround time – difference between finishing and arrival time ($f_i - a_i$)

Introduction

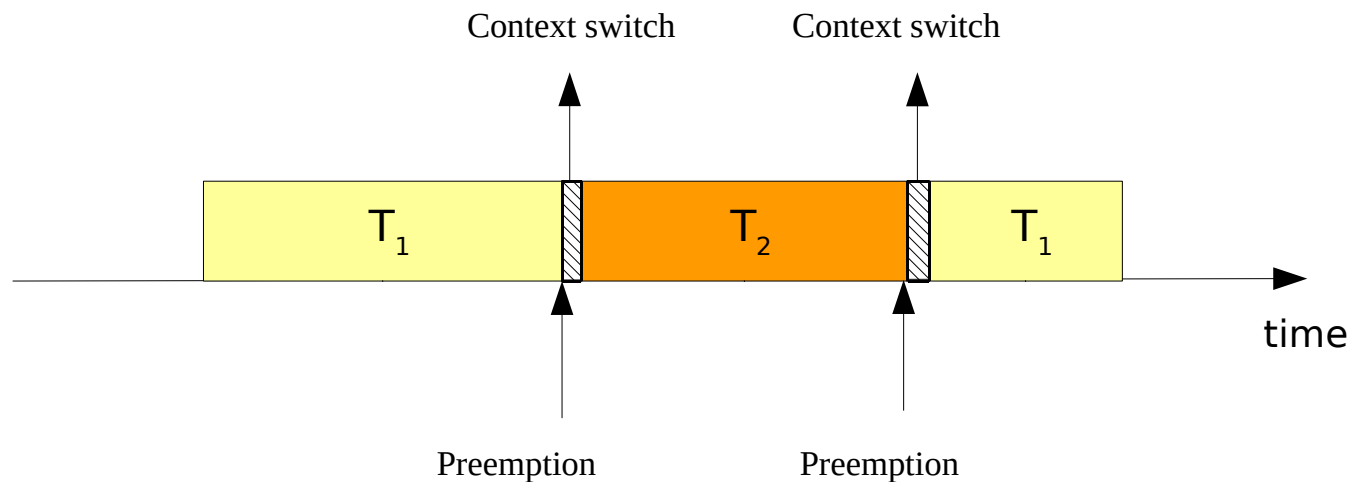
Preemption

- According to the schema shown so far the tasks run in “run-to-completion” fashion, i.e., without being interrupted
- For reasons that we will explain later, operating systems typically need to have the possibility of interrupting a task execution
- Preemption : operation for suspending the execution of a task and allocate the processor to another task
- Context switch : required when a preemption is performed
 - Save the context (stack and registers) of the suspended task
 - Resume the context of the next ready task

Introduction

Preemption

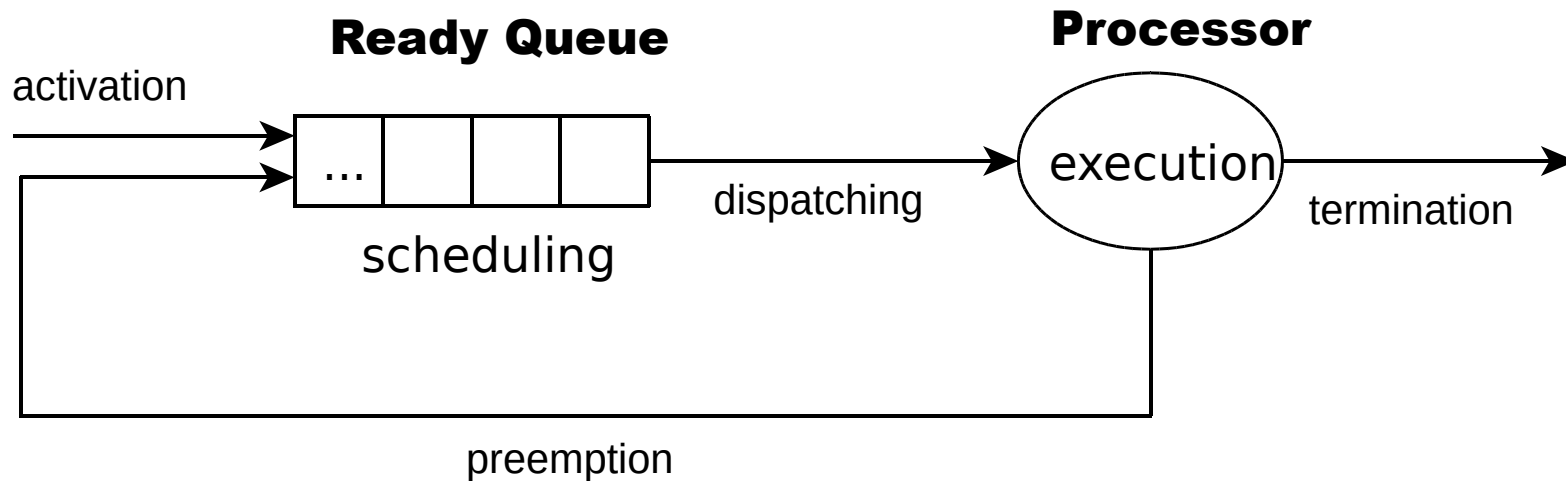
- Example
 - Task 1 runs for a while
 - Task 2 enters the system and the scheduler decides to suspend Task 1 to execute Task 2
 - The execution of Task 1 resumed later



Introduction

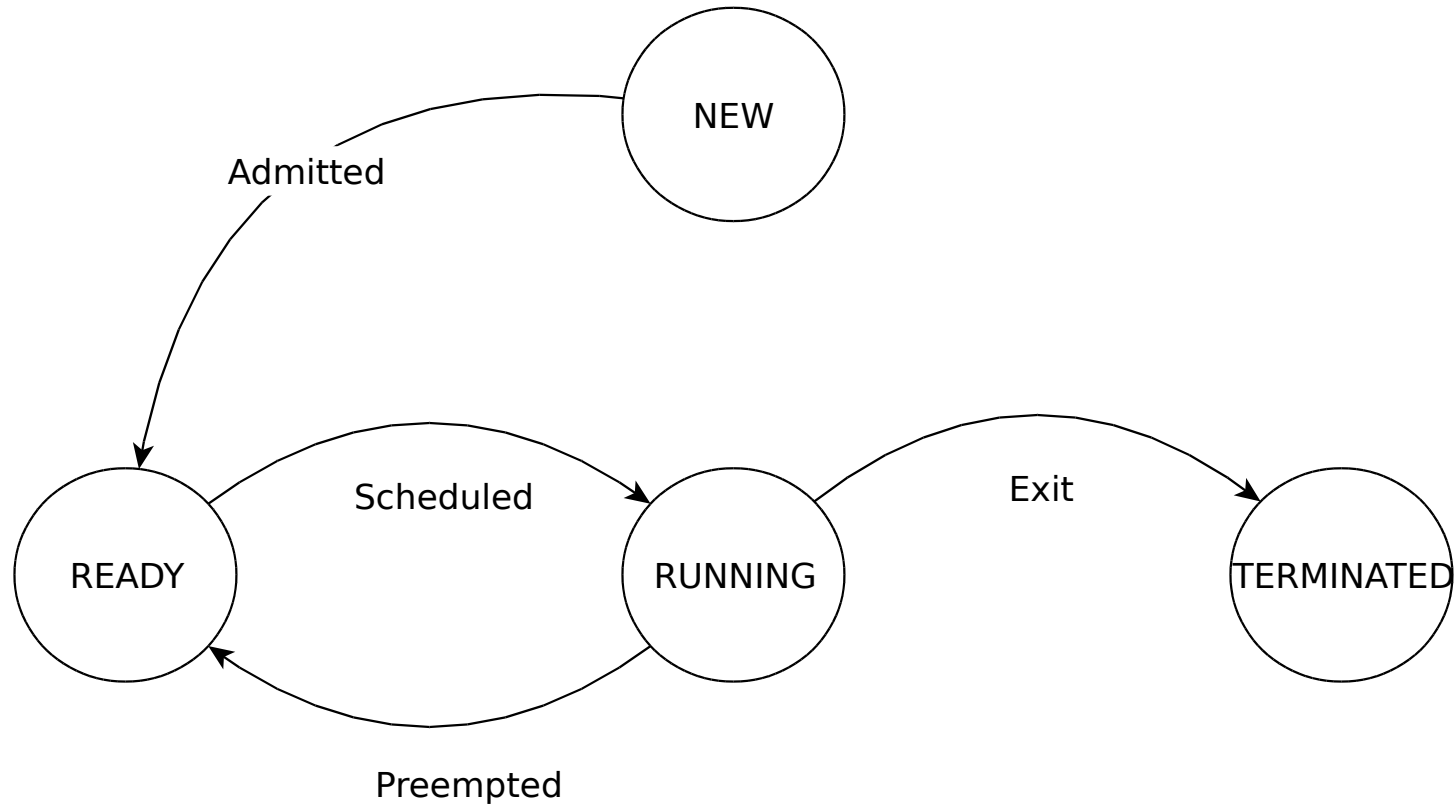
Scheduling overview with preemption

- Preempted tasks are moved back to the Ready queue



Introduction


Task state diagram



Introduction

Task parameters

- If the task can be preempted, this means that we will have several runs of the same task




- R_i : Response time – time between arrival time and the completion of the first run of the task (first response produced)
 - In non-preemptive systems $Z_i = R_i$
- W_i : Waiting time – difference between turnaround time and computation (burst) time: $(Z_i - C_i)$

Introduction

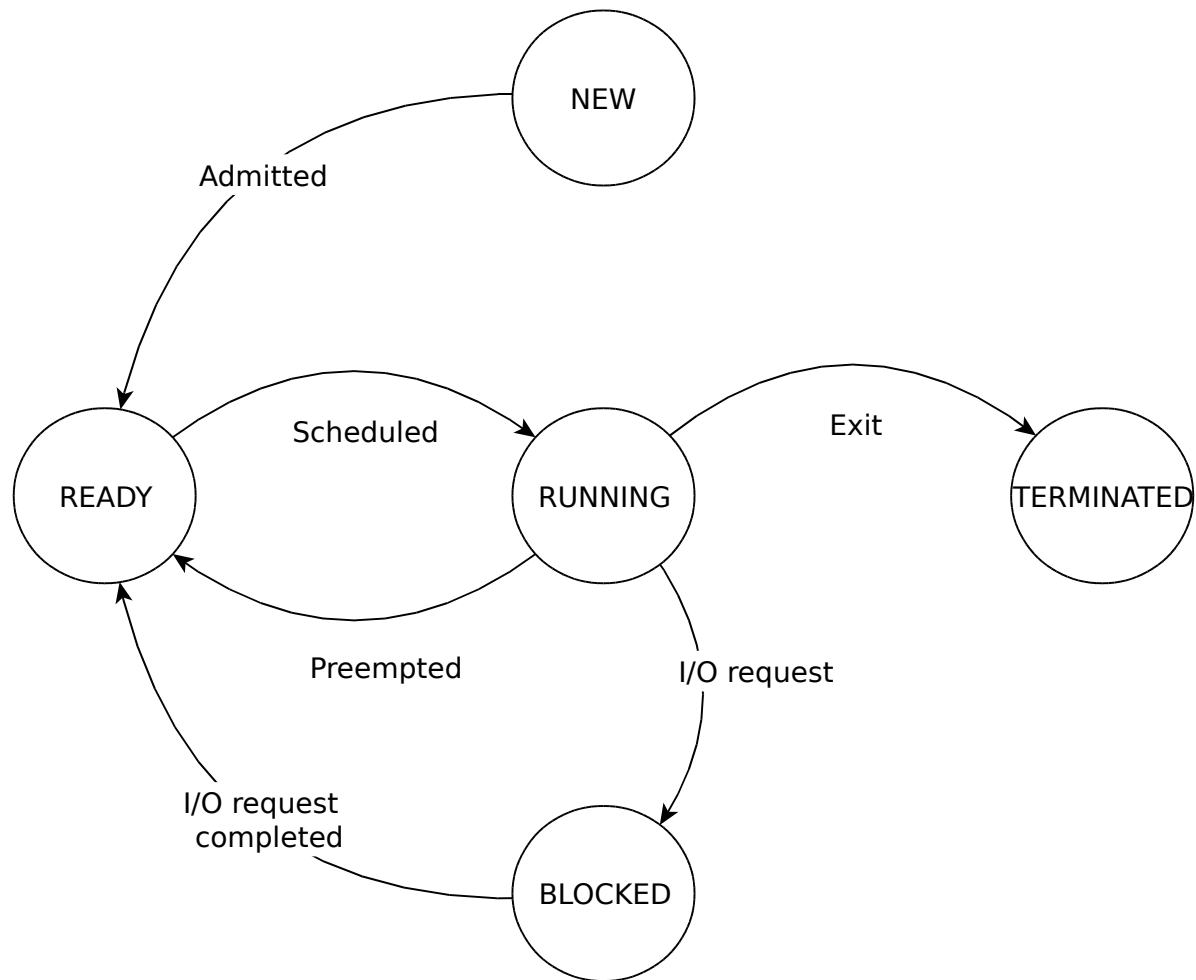
Scheduling and I/O

- Tasks may block their execution because waiting for data coming from a I/O requests (e.g., peripheral access, file read/write, etc...)
- Blocked tasks release the processor for other ready tasks



Introduction


Task state diagram with I/O requests



Introduction

Task boundness

- Depending on the type of operations dominating the lifetime of task, we may identify the “boundness” of a task
- CPU-bound
 - The task spent most of its time actually executing operations
 - Typical of “batch” / “background” tasks



- I/O bound
 - The task spent most of its time waiting for the completion of I/O operations
 - Typical of “interactive” / “foreground” tasks



Introduction

Scheduling metrics

- The scheduler aims at optimizing one (or more) objectives
 - We need metrics in order to evaluate the goodness of the policy
- Processor utilization : percentage of time the CPU is kept busy
- Throughput : number of tasks completing their execution per time unit
- Waiting time (avg) : average time the tasks spent in the ready queue
- Response time (avg) : average time the tasks spent in the ready queue before being served for the first time
- Fairness : do the tasks have a fair allocation of the processor?

Introduction

Problem statement

- Given a set of n tasks : $\mathbf{T} = \{ T_1, T_2, \dots, T_n \}$
- Given a set of processors : $\mathbf{P} = \{ P_1, P_2, \dots, P_m \}$
 - $|\mathbf{P}| = 1 \rightarrow$ uniprocessor systems
 - $|\mathbf{P}| > 1 \rightarrow$ multiprocessor systems
- Given a set of resources : $\mathbf{R} = \{ R_1, R_2, \dots, R_s \}$
- (Optional) Given a set of precedence relationships and constraints
- Define an ordered assignment of processors to tasks...
 - In order to optimize one or more objectives
 - ...and that constraints are not violated
- This problem has been shown to be NP-complete

Introduction


Common scheduling objectives

- Maximize *CPU utilization*, to keep it as busy as possible
- Maximize the *throughput*
- Minimize *turnaround time*
- Minimize *waiting time*
- Minimize *response time*
- Maximize the *fairness*
- ...

Introduction

Starvation

- Whatever is the objective of the scheduler, a good algorithm should guarantee that all the tasks are served
- *Starvation* is the undesirable perpetuated condition in which one (ore more) tasks cannot executed due to the lack of resources
 - e.g., a task is indefinitely postponed by the arrival of new ones



Classification of scheduling algorithms

Outline

- Preemptive vs Non-preemptive
- Static vs Dynamic
- Offline vs Online
- Optimal vs Heuristic

Classification of scheduling algorithms

Preemptive vs Non preemptive

- *Preemptive*
 - Running tasks can be interrupted at any time to allocate the processor to another active task
 - Good if we need a responsive system
 - Necessary for fairness objective

- *Non-preemptive*
 - Once started a task is executed until its completion (“run-to-completion”)
 - Scheduling decisions taken when tasks terminates
 - Good if we aim at minimizing tasks completion time
 - Negative impact on responsiveness
 - A single task may “monopolize” the processor

Classification of scheduling algorithms

Static vs Dynamic

- We can distinguish cases in which tasks parameters are known a priori, from cases in which we need to collect runtime information
- *Static*
 - Scheduling decisions are based on fixed parameters, whose values are known before task activation
 - Stronger assumptions required
- *Dynamic*
 - Scheduling decisions are based on parameters that typically changes at runtime, during the system activity
 - Need runtime feedback

Classification of scheduling algorithms

Offline vs Online

- *Offline*
 - The scheduler is executed offline on a set of known tasks (before their activation)
 - The outcome is stored into a data structure (e.g. a table) that is processed at runtime by a dispatcher
 - We need to know a priori the system workload (the set of tasks)
 - Necessary if we must provide some guarantee, by performing some preliminary check

- *Online*
 - The scheduler is executed at runtime
 - The system workload is variable, with tasks activated and terminated at random instant times

Classification of scheduling algorithms

Optimal vs Heuristic

- *Optimal*
 - The scheduler is based on an algorithm optimizing a given cost function, defined over the task set
 - The algorithm may be characterized by a not negligible complexity

- *Heuristic*
 - Algorithms based on a heuristic function
 - Tending to optimal scheduling, but without any guarantee about achieving it
 - Generally much faster than optimal algorithms

Scheduling algorithms

Outline

- First In First Out (FIFO)
- Shortest Job First (SJF)
- Shortest Remaining Time First (SRTF)
- Highest Response Ratio Next (HRRN)
- Round Robin (RR)

Scheduling algorithms

First In First Out (FIFO)


- Very simple algorithm
- Non-preemptive scheduling class
- Also known as “First Come First Served (FCFS)”
- Tasks are dispatched according to the arrival order
- When current running task terminates, the oldest one is selected
- Not good for responsiveness
 - Long tasks may monopolize the processor, penalizing short ones

Scheduling algorithms

First In First Out (FIFO)

- *Example*

Tasks	a	C	s	f	R
1	0	7	0	7	7
2	2	4	7	11	9
3	4	1	11	12	8
4	5	4	12	16	11



→ Avg. waiting time = $(0 + 5 + 7 + 7) / 4 = 4.75$

T₃ has to wait 7 time units before executing for 1

→ Avg. response/turnaround time = $(7 + 9 + 8 + 11) / 4 = 8.75$

Scheduling algorithms

Shortest Job First (SJF)


- Also known as “Shortest Job Next (SJN)”
- Non-preemptive scheduling class
 - Tasks executed in run-to-completion mode
- Tasks are scheduled in ascending order of computation time (C_i)
- Optimal algorithm
 - Minimize the average waiting time of the task set
 - Useful for benchmarking
- Starvation may occur!
 - What if small tasks enter the system while a long one has not been scheduled yet?

Scheduling algorithms

Shortest Job First (SJF)

- Example

Tasks	a	C	s	f	R
1	0	7	0	7	0
2	2	4	8	12	6
3	4	1	7	8	3
4	5	4	12	16	7



→ Avg. waiting time = $(0 + 6 + 3 + 7) / 4 = 4$

→ Avg. response/turnaround time = $(7 + 10 + 4 + 11) / 4 = 8$

Scheduling algorithms

Shortest Remaining Time First (SRTF)


- Preemptive variant of SJF
- Preempted tasks had already spent part of the computation time, so...
- Compare the *remaining time* instead of the initial computation time C_i
- Improved responsiveness
- Anyway.... starvation may still occur!
 - We can still end in a condition for which a task execution is indefinitely postponed because tasks with smaller C_i are continuously coming in the ready queue

Scheduling algorithms

Shortest Remaining Time First (SRTF)

- Example

Tasks	a	C	s	f	R	W
1	0	7	0			
2	2	4	2			
3	4	1				
4	5	4				




- At t_2 task T_2 is ready, $C_2 = 4$ while remaining time $RC_1 = 5$
Task T_2 si scheduled

Scheduling algorithms

Shortest Remaining Time First (SRTF)

- Example

Tasks	a	C	s	f	R	W
1	0	7	0			
2	2	4	2			
3	4	1	4	5		
4	5	4				



→ At t_4 task T_3 is ready, $C_3 = 1$ while remaining time $RC_1 = 5$ and $RC_2 = 2$


Task T_3 is scheduled and terminates after 1 time unit

Scheduling algorithms

Shortest Remaining Time First (SRTF)

- Example

Tasks	a	C	s	f	R	W
1	0	7	0			
2	2	4	2	7		
3	4	1	4	5		
4	5	4				




- At t_5 task T_4 is ready, $C_4 = 4$, $RC_1 = 5$, $RC_2 = 2$
Task T_2 is scheduled and terminates after 2 time units

Scheduling algorithms

Shortest Remaining Time First (SRTF)

- Example

Tasks	a	C	s	f	R	W
1	0	7	0	16		
2	2	4	2	7		
3	4	1	4	5		
4	5	4	7	11		




- At t_7 : $C_4 = 4$, $RC_1 = 5 \rightarrow$ Schedule T_4
- At t_{11} : T_4 terminates, T_1 is the last task left to schedule

Scheduling algorithms

Shortest Remaining Time First (SRTF)

- Example

Tasks	a	C	s	f	R	W	Z
1	0	7	0	16	2	9	16
2	2	4	2	7	2	1	5
3	4	1	4	5	1	0	1
4	5	4	7	11	6	2	6



- Avg. waiting time = $(9 + 1 + 0 + 2) / 4 = 3.0$
- Avg. response time = $(2 + 2 + 1 + 6) / 4 = \mathbf{2.75}$
- Avg. turnaround time = $(16 + 5 + 1 + 6) / 4 = 7$

Scheduling algorithms

Highest Response Ratio Next (HRRN)

- Non-preemptive scheduling class
- Evolution of Shortest Job First
- Select the task with the highest Response Ratio:

$$RR_i = \frac{W_i + C_i}{C_i}$$

- Short tasks (small computation time C) OK but...
- Waiting time (W) must be taken into account
- This has the effect of preventing starvation


Scheduling algorithms

Highest Response Ratio Next (HRRN)

- Example

Tasks	a	C	s	W	RR
1	0	7	0	0	1
2	2	4			
3	4	1			
4	5	4			

- T_1 is immediately scheduled, meanwhile... other three tasks enter the ready queue




Scheduling algorithms

Highest Response Ratio Next (HRRN)

- Example

Tasks	a	C	s	W	RR
1	0	7	0	0	1
2	2	4	...	5	2.25
3	4	1	...	3	4.00
4	5	4	...	2	1.50

- Once T_1 terminates ($t=7$), we compute the response ratio and pick T_3 , which has the highest ratio




Scheduling algorithms

Highest Response Ratio Next (HRRN)

- Example

Tasks	a	C	s	W	RR
1	0	7	0	0	-
2	2	4	...	6	2.50
3	4	1	7	3	-
4	5	4	...	3	1.75

- Once T_3 terminates ($t=8$), we compute the HRR and pick T_2 , which has the highest ratio




Scheduling algorithms

Highest Response Ratio Next (HRRN)

- *Example*

Tasks	a	C	s	W	RR
1	0	7	0	0	-
2	2	4	8	6	-
3	4	1	7	3	-
4	5	4	12	7	-

- T_4 is the last task in the queue and will start once T_2 finishes at $t=12$




Scheduling algorithms

Highest Response Ratio Next (HRRN)

- *Example*

Tasks	a	C	s	W	f
1	0	7	0	0	7
2	2	4	8	6	12
3	4	1	7	3	8
4	5	4	12	7	16



- Avg. waiting time = $(0 + 6 + 3 + 7) / 4 = 4$
- Avg. response/turnaround time = $(7 + 10 + 4 + 11) / 4 = 8$

Scheduling algorithms

Round Robin (RR)

- Preemptive scheduling class
- *Time sharing* approach
 - Task are scheduled to process for a given time quantum (q) or (time slice)
 - When the time quantum expires the task is preempted and moved back to the ready queue
- For a known number of tasks n , the maximum waiting time is bound to $(n - 1) * q$
- Good for managing a *fair* allocation of the processor
- Good for managing the system *responsiveness*
- Turn-around time usually worst than SJF
- No starvation


Scheduling algorithms

Round Robin (RR)

▪ Example 1

- Time quantum $q=2$

Tasks	a	C	s	f	R	W
1	0	7	0			
2	2	4	2			
3	4	1				
4	5	4				



- t_0 : Ready queue = { **T1** }
- t_2 : Ready queue = { **T2**, T1 }
- t_4 : Ready queue = { **T1**, T3, T2 }


Scheduling algorithms

Round Robin (RR)

▪ Example 1

- Time quantum $q=2$

Tasks	a	C	s	f	R	W
1	0	7	0			
2	2	4	2	9		
3	4	1	6	7		
4	5	4	9			



- t_6 : Ready queue = { **T3**, T2, T4, ¹T1 }
- t_7 : Ready queue = { **T2**, T4, T1 }
- t_9 : Ready queue = { **T4**, T1 }


Scheduling algorithms

Round Robin (RR)

▪ Example 1

- Time quantum $q=2$

Tasks	a	C	s	f	R	W
1	0	7	0	16		
2	2	4	2	9		
3	4	1	6	7		
4	5	4	9	15		



- t_{11} : Ready queue = { **T1**, T4 }
- t_{13} : Ready queue = { **T4**, T1 }
- t_{15} : Ready queue = { **T1** }


Scheduling algorithms

Round Robin (RR)

Example 1

- Time quantum $q=2$

Tasks	a	C	s	f	R	W	Z
1	0	7	0	16	2	9	16
2	2	4	2	9	2	3	7
3	4	1	6	7	3	2	3
4	5	4	9	15	6	6	10



- Avg. waiting time = $(9 + 3 + 2 + 6) / 4 = 5$
- Avg response time = $(2 + 2 + 3 + 6) / 4 = 4.33$
- Avg. turnaround time = $(16 + 7 + 3 + 10) / 4 = 9$

Scheduling algorithms

Round Robin (RR)

- *Example 1*
 - Time quantum $q=2$

Tasks	a	C	s	f	R	W	Z
1	0	7	0	16	2	9	16
2	2	4	2	9	2	3	7
3	4	1	6	7	3	2	3
4	5	4	9	15	6	6	10

Let's see what happens if we change the time quantum for example by setting **q=3**


Scheduling algorithms

Round Robin (RR)

▪ Example 2

- Time quantum $q=3$

Tasks	a	C	s	f	R	W
1	0	7	0		3	
2	2	4	3		4	
3	4	1				
4	5	4				



- t_0 : Ready queue = { **T1** }
- t_3 : Ready queue = { **T2**, T1 }
- t_6 : Ready queue = { **T1**, T3, T4, T2 }


Scheduling algorithms

Round Robin (RR)

▪ Example 2

→ Time quantum
 $q=3$

Tasks	a	C	s	f	R	W
1	0	7	0		3	
2	2	4	3		4	
3	4	1	9	10	6	
4	5	4	10		8	



→ t_9 : Ready queue = { **T3**, T4, T2, T1 }

→ t_{10} : Ready queue = { **T4**, T2, T1 }

→ t_{13} : Ready queue = { **T2**, T1, T4 }


Scheduling algorithms

Round Robin (RR)

▪ Example 2

→ Time quantum
 $q=3$

Tasks	a	C	s	f	R	W
1	0	7	0	15	3	
2	2	4	3	14	4	
3	4	1	9	10	6	
4	5	4	10	15	8	



- t_{14} : Ready queue = { **T1**, T4 }
- t_{15} : Ready queue = { **T4** }
- t_{16} : Ready queue = { }


Scheduling algorithms

Round Robin (RR)

▪ Example 2

→ Time quantum
 $q=3$

Tasks	a	C	s	f	R	W	Z
1	0	7	0	15	3	8	15
2	2	4	3	14	4	8	12
3	4	1	9	10	6	5	6
4	5	4	10	16	8	7	11



→ Avg. waiting time = $(8 + 8 + 5 + 7) / 4 = 7$

→ Avg response time = $(3 + 4 + 6 + 8) / 4 = 5.25$

→ Avg. turnaround time = $(15 + 12 + 6 + 10) / 4 = 11$

For $q=3$ the responsiveness of the system results to be decreased

Scheduling algorithms

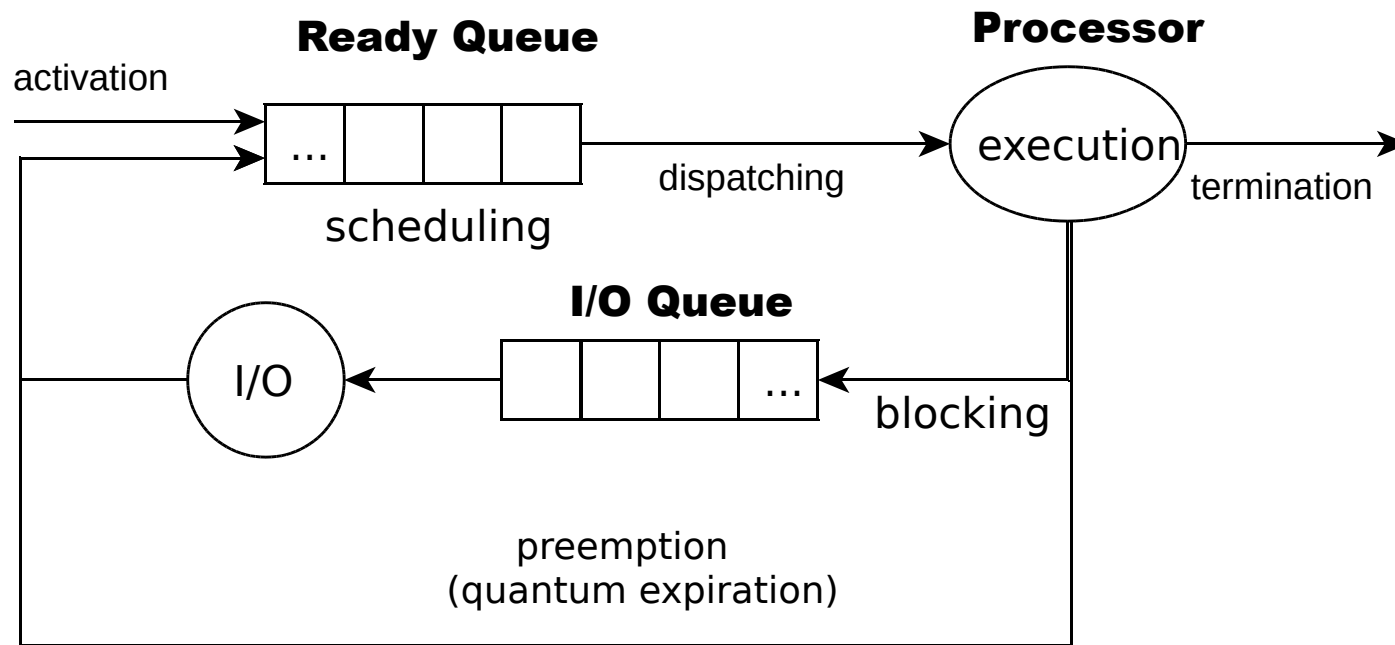
Time sharing

- The length of the time quantum is a critical choice
- “*Long*” quantum :
 - Scheduler performance → FIFO
 - Favor CPU-bound tasks
 - Less context switches → lower overhead
- “*Short*” quantum :
 - Reduced average waiting time
 - Good for responsiveness
 - Easier to achieve a fair assignment of the processor
 - Favor I/O-bound tasks
 - More context switches → higher overhead

Scheduling algorithms

Time sharing

- Scheduling with I/O occurrences and time sharing



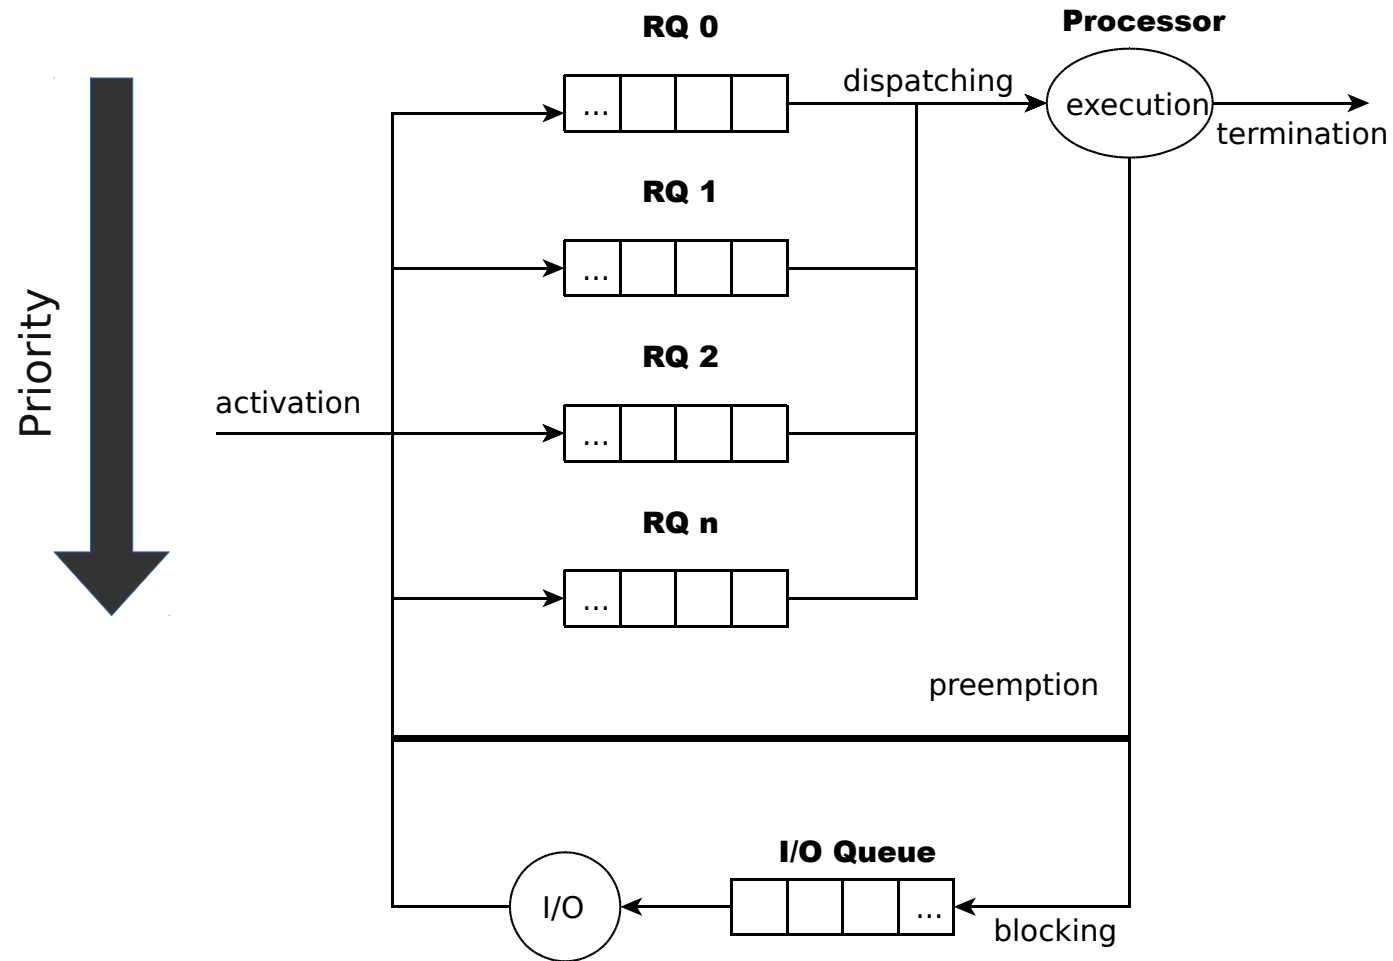
Priority based scheduling

Priority based scheduling

- *Priority* is a task parameter through which we can specify the importance of a task
 - ➔ **Fixed** or **dynamic** (runtime variable)
- Conventionally, the priority is expressed with an integer value
 - ➔ The lower the integer value the higher the priority
 - ➔ The higher the integer value the lower the priority
- Priority-based scheduler takes into account this parameter...
 - ➔ To drive the **dispatching time** (which task to schedule first)
 - ➔ To define the **length of the time quantum** assigned to each task
- Priority-based scheduler typically need to manage multiple ready queues
 - ➔ Each ready queue is associated to a different priority level

Priority based scheduling

Multi-level Queue Scheduling



Priority based scheduling

Multi-level Queue Scheduling

- For each queue we can specify a different scheduling algorithm (e.g., RR or FCFS)
- The first task to schedule is picked from the topmost non-empty queue (highest priority)
- Tasks cannot be moved from one ready queue to another
- Possible to assign a different time quantum per queue
 - *Foreground (interactive)* tasks and *background (batch)* tasks use CPU time differently
 - We may give high priority to interactive tasks
- Risk of **starvation!**
 - While highest priority queues are populated by new tasks, the scheduling of tasks in lower priority queues is delayed

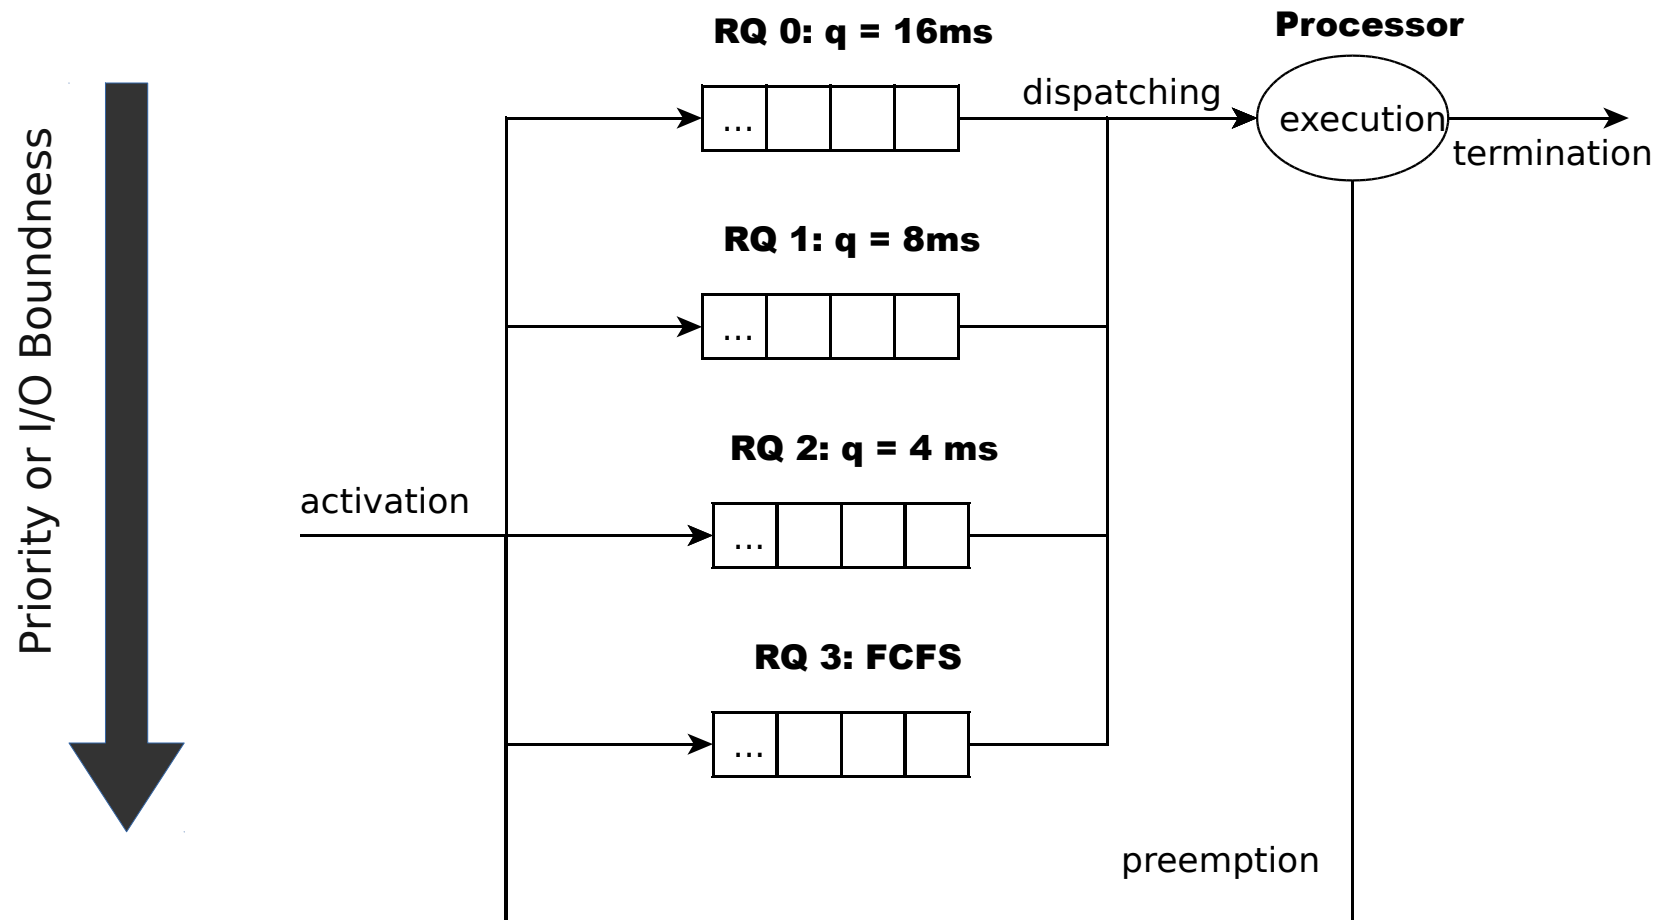
Priority based scheduling

Multi-level Feedback Queue Scheduling

- How to determine if a task is interactive or batch?
- How to practically infer at runtime if a task is CPU-bound or I/O-bound?
- Most of the scheduling algorithm previously introduced rely on information known “a priori”
 - How to know the *computation time* C_i in advance?
 - Often unrealistic assumptions
- Practical implementations need runtime feedback about how the scheduled tasks use the processor
- The Multi-level Feedback Queue Scheduling aims at overcome the limitations above

Priority based scheduling

Multi-level Feedback Queue Scheduling



Priority based scheduling

Multi-level Feedback Queue Scheduling

- A quantum length is associated to each queue
- A task enters a queue, depending on the priority
- Tasks are moved among queues
 - Priority is dynamic no longer fixed
- When scheduled, if the task does not terminate in the time quantum allocated, it is moved to another queue (with a different quantum length)
 - CPU-bound (batch) tasks are progressively moved in queues with longer time quantum
 - I/O-bound (interactive) tasks are progressively moved in queue with shorter time quantum

Priority based scheduling


Multi-level Feedback Queue Scheduling

- What about starvation?
- We can exploit the possibility of moving the task among the queues (i.e., changing the priority at runtime)
- **Aging**
 - The priority of the task is increased as long as it spend time in ready queues (it gets “older”)
 - Prevent a task from being indefinitely postponed by new coming higher priority tasks

Priority based scheduling

Aging


- Example
 - Three tasks at different priorities in the ready queues



Priority based scheduling

Aging

- Example
 - T1 is scheduled, so the priority 0 queue becomes empty



Priority based scheduling

Aging


- Example
 - T2 is scheduled
 - T4 “ages” and its priority is increased to 1
 - T3 (priority 1) arrives, but this does not prevent T4 from running



Multi-processor scheduling

Problem overview

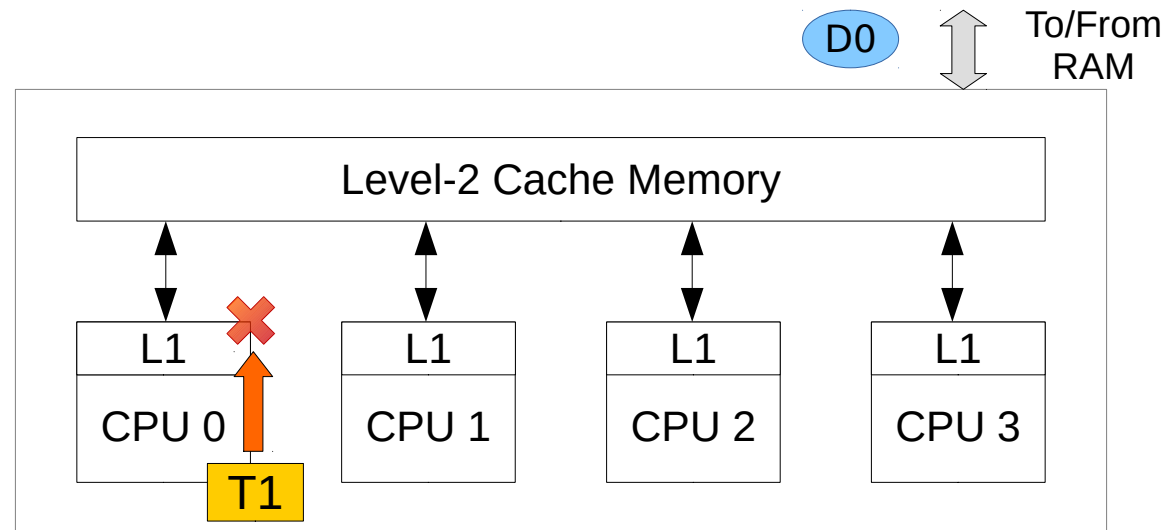
- The scheduler must choose not only the task to execute, but also to **which** processor dispatch it
- This introduces further aspect to consider....



Multi-processor scheduling

Processor affinity

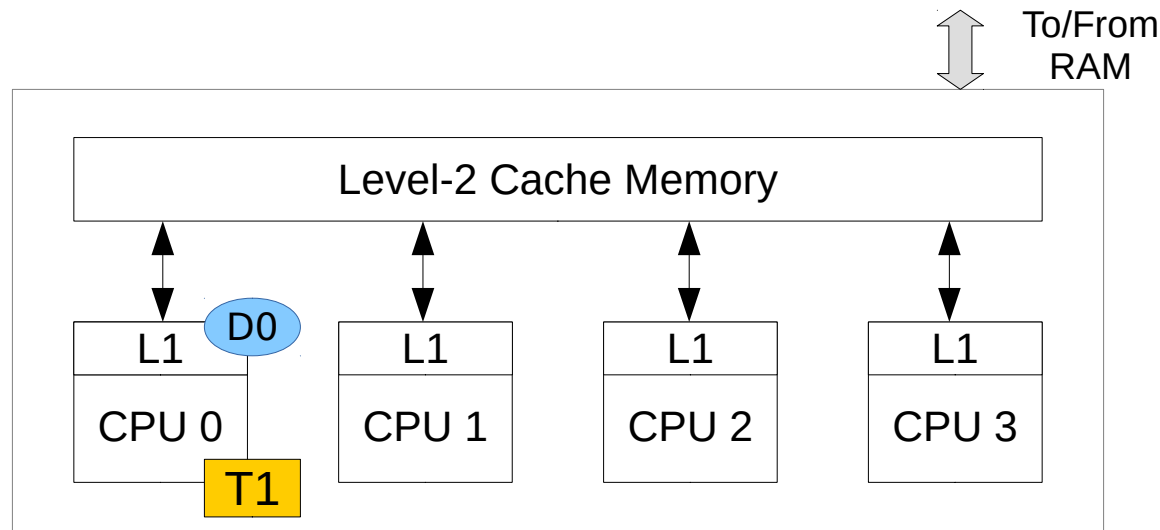
- Cache memory in modern multi-core processor dramatically affect performance
- Data should stay as closer to the CPU as possible
 - Less CPU cycles to complete load/store



Multi-processor scheduling

Processor affinity

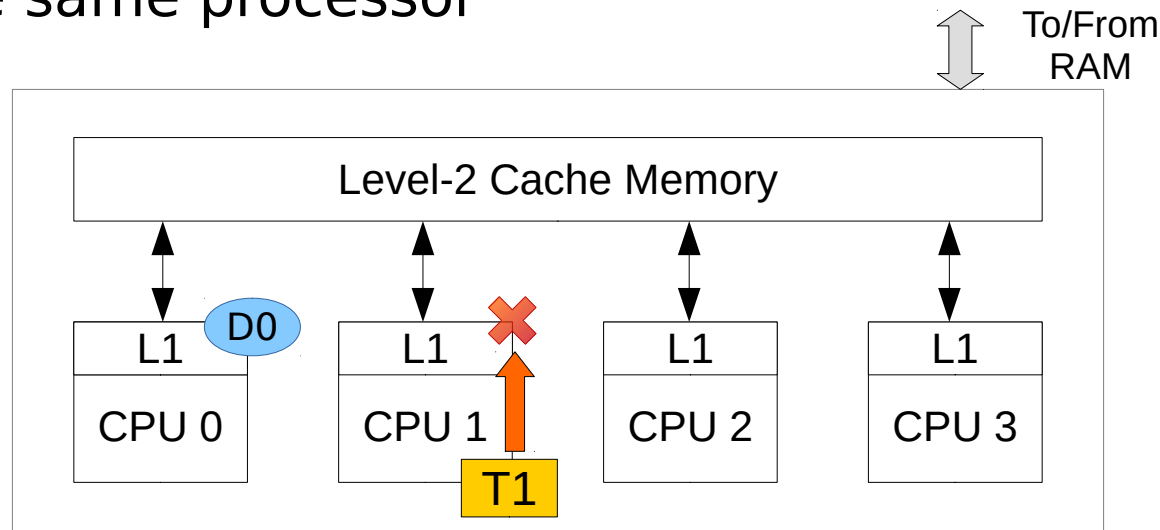
- First time a task looks for a data in L1-cache a “cache missed” is experienced
 - Lot of CPU cycles lost, waiting for the data being retrieved from higher memory levels
- Once data is retrieved we want it to remain in L1-cache



Multi-processor scheduling

Processor affinity

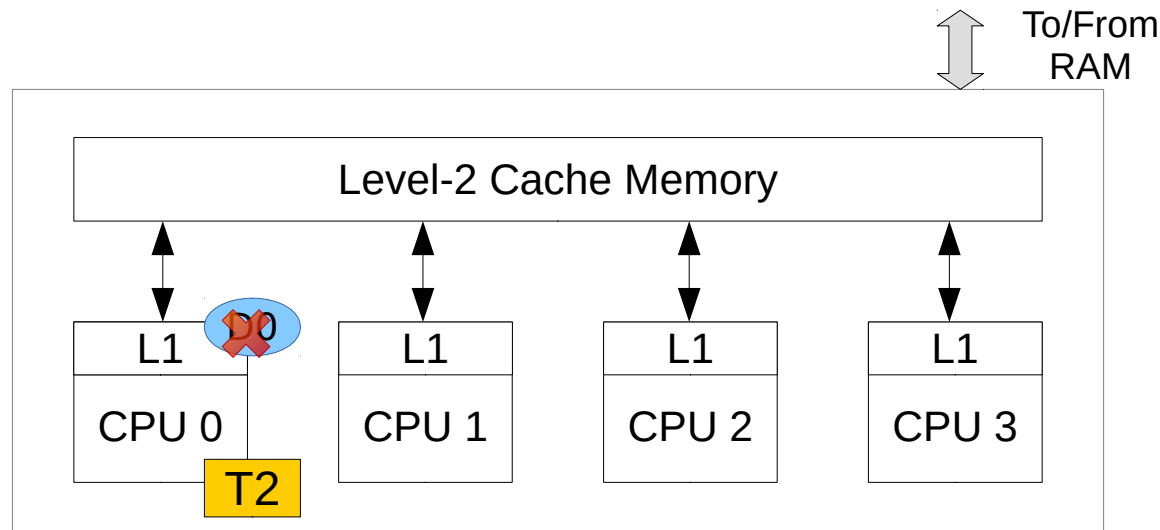
- If (in a preemptive scheduler) the next run of a task is scheduled on a different processor, new L1-cache misses are experienced
 - Loss of performance
- Processor affinity should drive the algorithm to reschedule the task on the same processor



Multi-processor scheduling

Processor affinity

- If a different task is scheduled on the processor running task T_1 (CPU 0), the cache is invalidated
 - Once task T_1 is rescheduled in CPU 0, cache misses are experienced again
- The ordering of the scheduling choices should take this into account



Multi-processor scheduling


Scheduler design

- Single queue vs Multiple queues
- Single scheduler vs Multiple per-processor schedulers

Multi-processor scheduling

Single queue vs Multiple queues

- *Single queue*
 - All the ready tasks waits in the same global queue



Multi-processor scheduling


Single queue vs Multiple queues

- *Single queue*
 - All the ready tasks waits in the same global queue
 - Simple design
 - Good for fairness
 - Good for managing CPU utilization
 - Scalability limitations
 - Scheduling code can run in whatever processor (it may require a lock to safely access the ready queue)

Multi-processor scheduling

Single queue vs Multiple queues


- *Multiple queues*
 - A ready queue for each processor



Multi-processor scheduling

Single queue vs Multiple queues


- *Multiple queues*
 - A ready queue for each processor
 - More scalable approach
 - Potentially more overhead (more data structures to handle)
 - Easier to exploit data locality (processor affinity)
 - Possibility of adopting a single global scheduler or per-CPU schedulers
 - Need of **load balancing**



Multi-processor scheduling

Load balancing


- Unbalanced ready queues have a negative impact from several point of views
- CPU utilization
 - Processors may be idle (due to empty queue), while other queues have waiting tasks
- Performance
 - Waiting times and response times can be reduced by moving the task in a queue with a lower number of tasks



Multi-processor scheduling

Load balancing


- Unbalanced ready queues have a negative impact on many aspects...
- Thermal management
 - By balancing the activity of the processors we can level the temperature distribution
 - Temperature has a direct impact on power consumption, reliability and system lifetime



Multi-processor scheduling

Task migration

- Load balancing is typically performed via task migration
 - Task are moved more to less loaded ready queue
 - Implications on processor/cache affinity




- *Push model* – the OS periodically checks queues length periodically moves tasks if balancing is required
- *Pull model* – each processor notifies an empty queue condition to the OS, or picks tasks from other queues

Multi-processor scheduling

Task migration

- *Work stealing* is an example of pull model based approach
 - ➔ Each per-CPU scheduler can “steal” a task from the other in case of empty queue



- ➔ Scalable in theory but in practice we must consider the overhead...

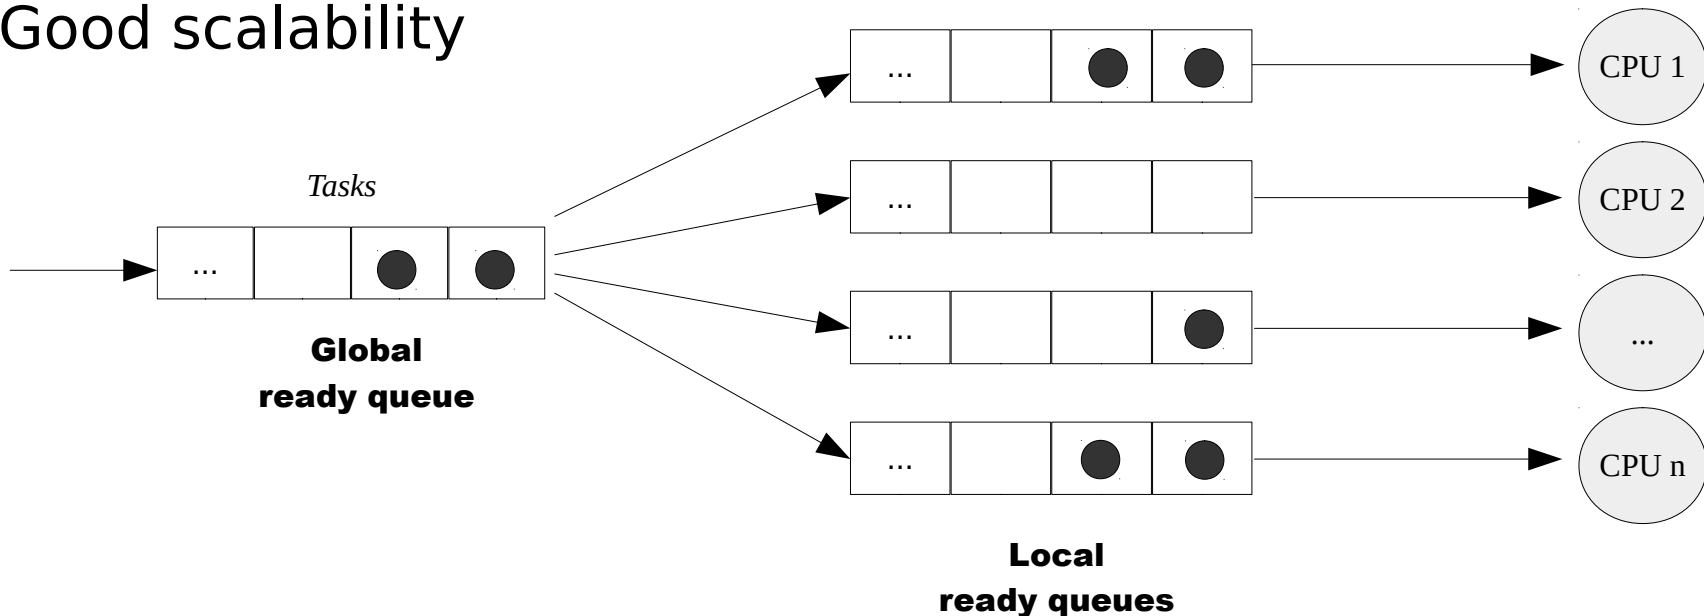
We need to protect the access to ready queues (global) from concurrent accesses

How to identify the shortest queue? Do we need to keep a per-length order?

Multi-processor scheduling

Hierarchical Queues

- A global queue dispatching the tasks in the local ready queues
 - Hierarchy of schedulers
- Better control over CPU utilization / load balancing
- Good for managing processor affinity
- Good scalability



Questions?

Contacts

- William Fornaciari
<https://home.deib.polimi.it/fornacia>
william.fornaciari@polimi.it
- **HEAPLab**
DEIB Politecnico di Milano
Building 21, Floor 1
Phone: 9613 / 9623
- **Slides credits:**
 - Giuseppe Massari <giuseppe.massari@polimi.it>

