



POLITECNICO
MILANO 1863

DIPARTIMENTO DI ELETTRONICA,
INFORMAZIONE E BIOINGEGNERIA



EMBEDDED SYSTEMS
Prof. William Fornaciari

Real-time Systems

V2.0

William Fornaciari

[<william.fornaciari@polimi.it>](mailto:william.fornaciari@polimi.it)

Real-time Systems

Outline

- Introduction to Real-time systems
- Predictability
- Real-time Operating Systems (RTOS)
- Real-time Scheduling

Introduction to Real-time Systems

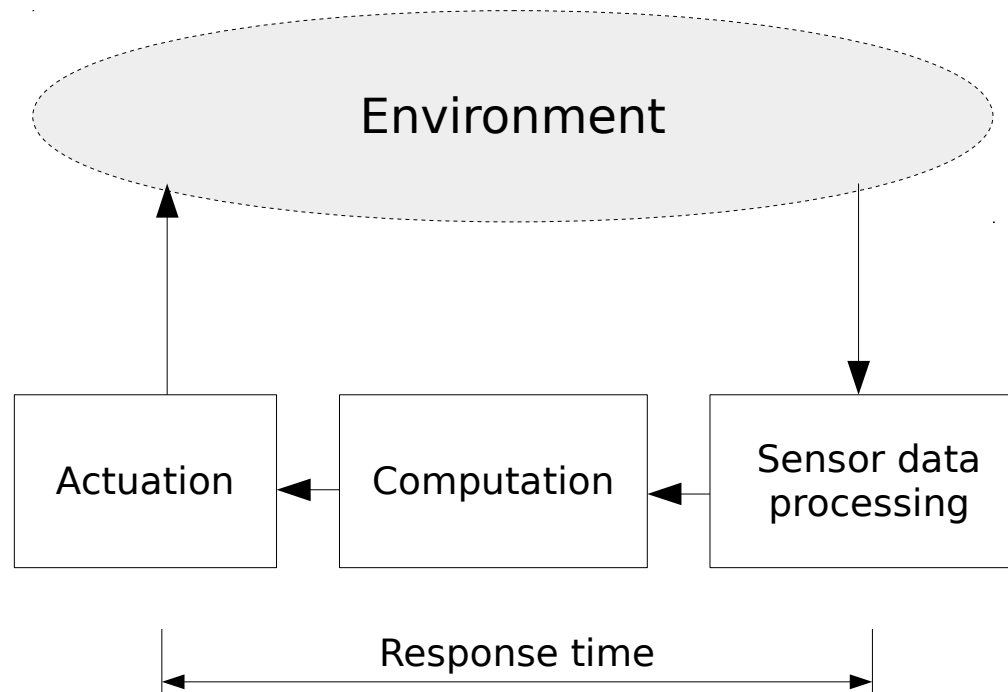
What is a real-time system?

- *Definition 1* – A real-time system is a system in which the correctness of its behavior depends on
 - The **logical correctness** of the tasks' output
 - The **time** within the output are produced
- *Definition 2* – A real time system is a system which has to respond to externally generated input stimuli within a finite and specified period
 - Reactive system with timing constraints
- A wide class of embedded systems is made by real-time systems

Introduction to Real-time Systems


Structure

- Real-time systems are often designed to address dynamic control problems
 - Timing constraints in the control-loop



Introduction to Real-time Systems

Example: *Airbag*



- When a collision is detected by the crash sensor, the inflation of the air bag must be triggered
- The trigger must be within 10-20ms from the time instant in which the collision occurred (has been detected)

Classification

Timing constraints

- The execution of tasks in real-time systems is typically characterized by stringent timing constraints
 - i.e. tasks execution should complete within a **time deadline**
- Depending on the consequences of missing the deadline, we can distinguish among three criticality classes of real-time
 - **Hard** – missing the deadline may lead to catastrophic events on persons or on the system under control
 - **Firm** – missing the deadline does not cause system damages, but the output value is not valid
 - **Soft** – missing the deadline does not affect the validity of the output, but it results to be a performance degradation

Classification

Example: *Autonomous Emergency Braking System*

- Detection of obstacles (pedestrian, stopped cars,...)
- Trigger the braking action
 - Consider the current car speed, distance of the obstacle...
 - How much time to trigger the brake?
- Consequences of a deadline miss?



Classification

Example: *Video player application*

- What happens if the application processes a frame in 100ms (10 FPS) instead of 40ms (25 FPS) ?



Properties

It's all about time

- *Timeliness*
 - System must include time management mechanisms
- *Predictability*
 - Possibility of predicting a priori if the timing constraints can be guaranteed or not
 - The temporal behavior of the system must be analyzable at design-time
- *Determinism*
 - Provide timing guarantees on the execution of tasks at run-time, also in case of occurrence of external events to handle

Properties

Misconceptions

- “Real-time computing is equivalent to fast computing.” (NO!)
- *Fast computing* objective is to minimize the average response time of a task set
- *Real-time computing* objective is to meet the individual timing requirement of each task!

Predictability


Predictability: potential enemies

- Processor
- Cache
- Direct Memory Access (DMA)
- Interrupts
- System calls
- Memory management
- Locking mechanisms
- Programming language

Predictability

Processor

- Modern processor architectures include hardware mechanisms to improve average performance
 - *Instruction / data Prefetching*
 - *Pipeline length and status*
 - *Speculation branch prediction,*
 - ...




- **OK for average performance, but they are also source of non-determinism!**

Predictability

Cache

- Memory access instructions (LOAD/STORE) may require an unpredictable number of cycles depending on the occurrence of a cache *hit* or *miss*




Predictability

Processors and Cache

- *Possible solution*
 - Micro-controllers have a much simple design, which provides more guarantees in terms of predictability
 - Adopting specific processors for real-time applications
e.g., ARM Cortex R family

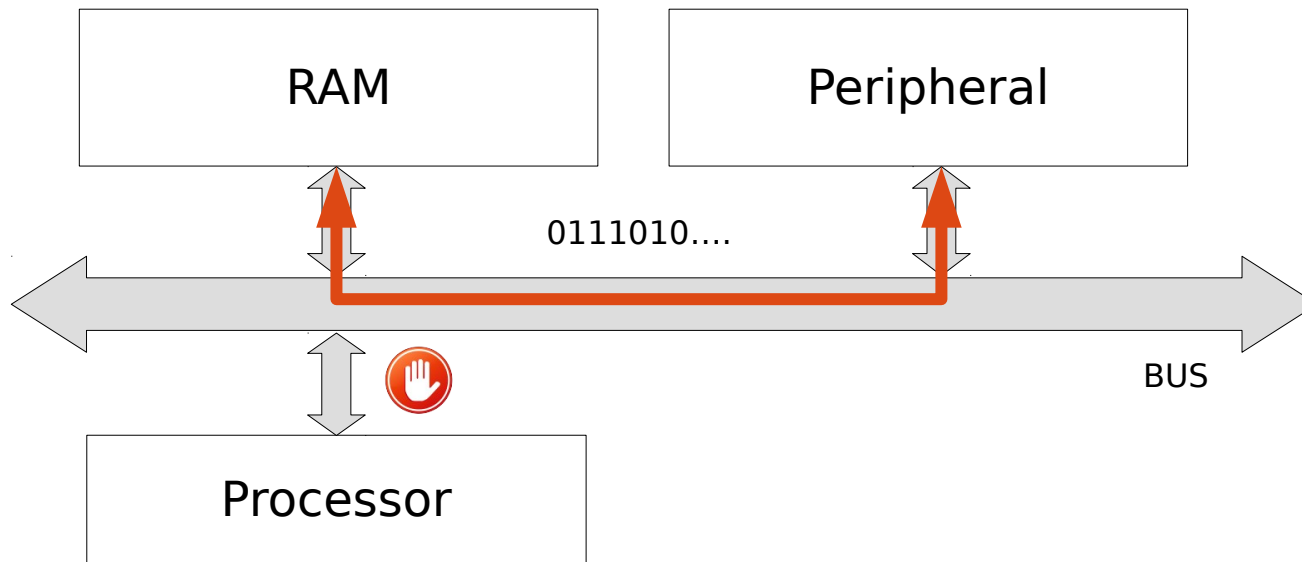
ARM Architecture: For Diverse Embedded Processing Needs



Predictability

Direct Memory Access (DMA)

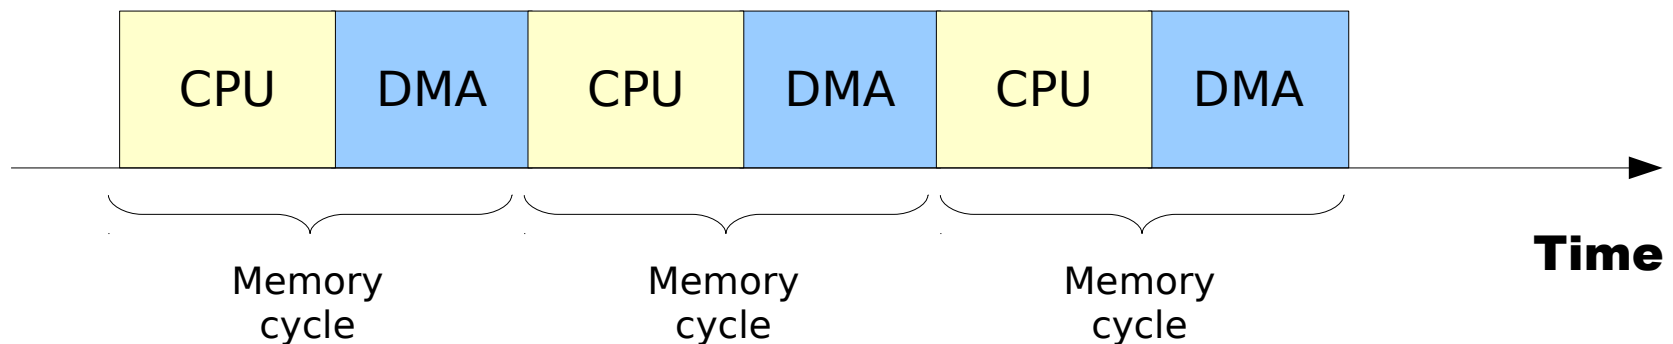
- DMA is used to transfer data between devices and main memory, without involving the CPU
 - Since the bus is shared, DMA could steal cycles to the CPU to perform data transfers (*cycle stealing*)



Predictability

Direct Memory Access (DMA)


- *Possible solution*
 - **Time-slice** method: split a memory cycle into separate time slots (for CPU and DMA), allowing predictive accesses to the shared bus
 - More expensive but more predictive solution
 - CPU and DMA do not conflict
 - Response time does not increase due to DMA



Predictability

Interrupts

- Serving an I/O request may delay the termination of the current task, with risk of missing the deadline




- *Possible solutions*
 - **(A) Disable interrupts (except the timer interrupt)**
 - Low processor efficiency on I/O operations
 - Tasks must manage I/O operations by themselves via polling and direct access to peripheral registers

Predictability

Interrupts

- Serving an I/O request may delay the termination of the current task, with risk of missing the deadline




- *Possible solutions*
 - **(B) I/O requests are periodically handled by OS routines**
Interrupts from devices are disabled but...
Dedicated kernel routines periodically handle I/O operations

Predictability

Interrupts

- Serving an I/O request may delay the termination of the current task, with risk of missing the deadline

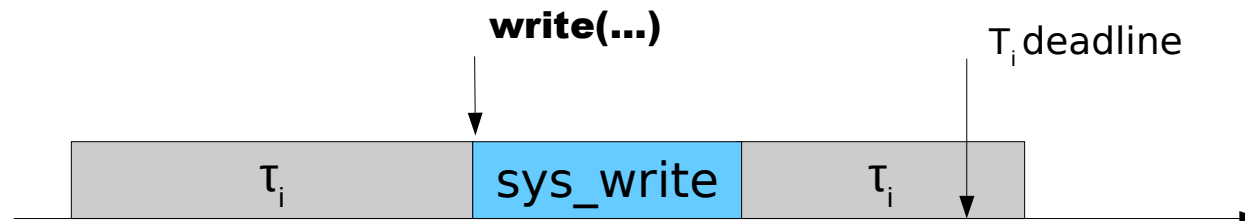


- *Possible solutions*
 - **(C) Minimal interrupt service routine**
 - The routines just set a flag and create a normal task for handling the request
 - The I/O handler task is scheduled as any other task

Predictability

System calls

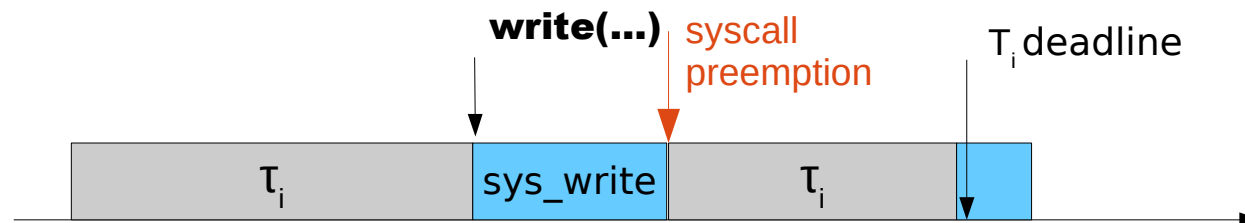
- Serving an I/O request may delay the termination of the current task, with risk of missing the deadline



- *Possible solution*

- ➔ **Interruptible kernel routines**

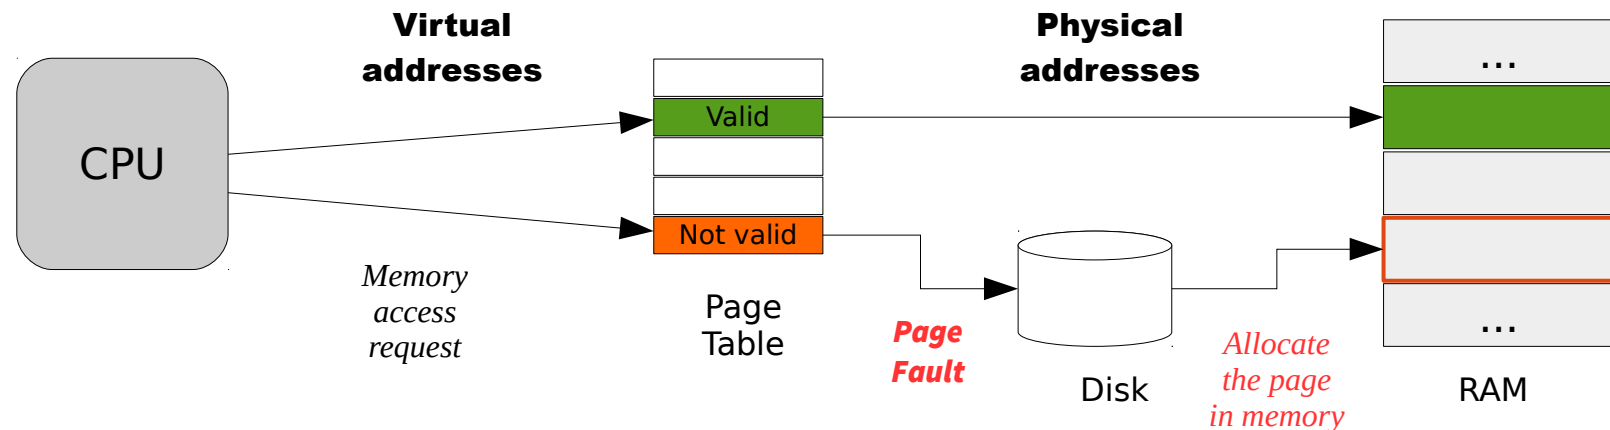
A high priority task approaching the deadline should preempt also system calls



Predictability

Memory management

- On-demanding memory page loading leads to unpredictable delays
 - Page faults similar to cache misses with higher latencies




- *Possible solutions*
 - Fixed memory management schema
 - Use static allocation

Predictability

Locking mechanisms

- Consider the *Priority Inversion* problem (see slides on “Multi-tasking Synchronization...”)
 - High-priority tasks blocked by low-priority tasks
 - Non deterministic delays due to locks



- Possible solution*
 - Use locking data structure implementing suitable **Resource Access Protocols** (e.g., Priority Ceiling Protocol)

Predictability

Programming languages

- Most programming languages DO NOT provide...
 - Constructs to express timing constraints
 - Protocols for shared data accesses
 - ...
- Some constructs are source of non determinism
 - e.g. the *switch-case* prevents the performing of Worst-Case Execution Time (WCET) analysis
- *Possible solutions*
 - Adopting specific programming languages for real-time
 - Avoid dynamic data allocations
 - Avoid recursion
 - Maximum number of loops iterations must be known a priori

Real-time Operating Systems (RTOS)

Outline

- Overview
- Features
- Performance

Real-time Operating Systems (RTOS)

What is a RTOS?

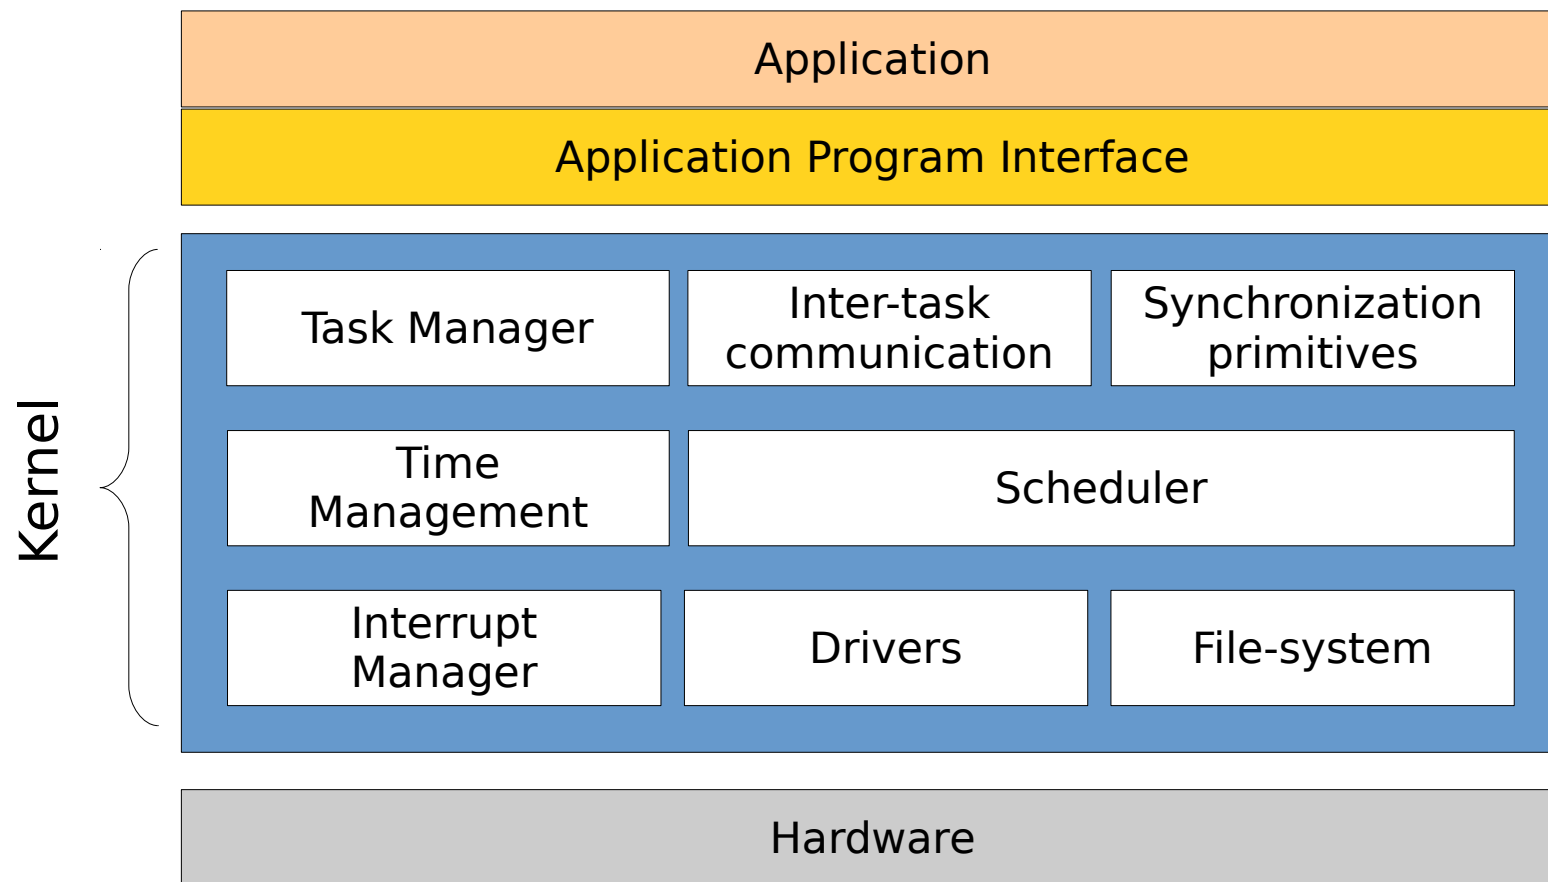
- Operating system specifically designed to address the problem of guaranteeing real-time requirements

Why to use a RTOS?

- Part of the complexity managed by the RTOS (e.g., access to peripheral registers, ...)
- Faster development cycle w.r.t *bare-metal* programming
- Application timing constraints managed by real-time specific schedulers
- Part of sources of non determinism are managed by the RTOS (e.g, predictive ISR, deterministic memory management, resource access protocols, ...)

Real-time Operating Systems (RTOS)

Overview



Real-time Operating Systems (RTOS)

Features

- Multi-tasking / Concurrency support
 - Task creation and execution support
 - Specific real-time scheduling policies
 - Effective inter-task synchronization and communication mechanisms

- Low-level user control
 - Fine-grained control over task priority assignment
 - Control over processes and pages that must reside in memory
 - Selection of scheduling policies
 - Selection of power management strategies (if any)
 - ...

Real-time Operating Systems (RTOS)

Features

- Reliability and robustness support
 - Manage system failures in such a way to preserve as much data and capability as possible
 - Multiple attempts to improve data consistency
 - For some real-time systems a minimal service level must be guaranteed in any case

Real-time Operating Systems (RTOS)

Typical characteristics

- Small memory footprint
 - RTOS usually requires a few KB of memory
- Predictive duration of system calls and kernel services
 - Preemptive kernel
- Responsiveness
 - Bounded *context-switch* and *interrupt management* latencies
- Determinism
 - External events must be handled such that no deadline misses occur

Real-time Operating Systems (RTOS)

Examples

- BeRTOS
- Contiki
- ChibiOS/RT
- ERIKA Enterprise
- FreeRTOS
- Mbed-rtos
- Miosix
- TinyOS
- VxWorks
- ...


Real-time tasks

Outline

- Task parameters
- Task periodicity

Real-time tasks

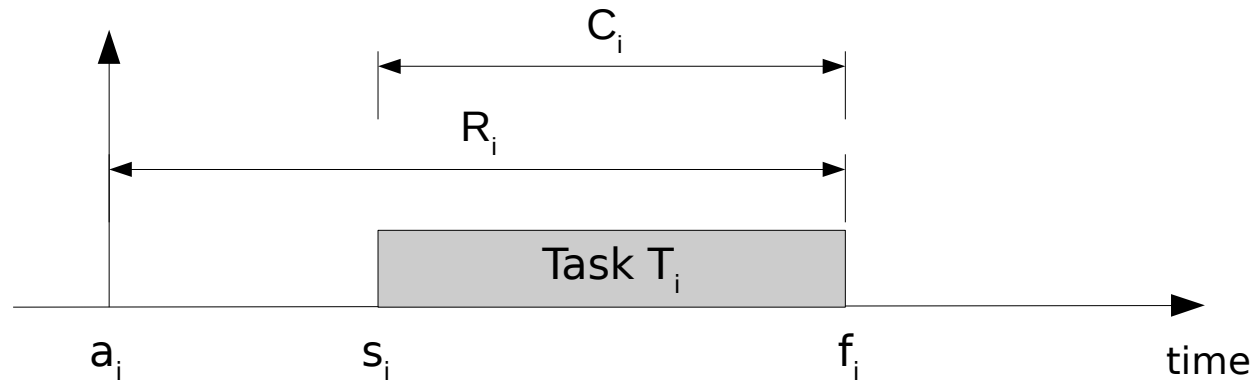
Real-time task parameters



- (See “Task scheduling” slides) ...
- R_i : Response time – time between arrival time and the completion of the task
- C_i : Computation time (or Burst time) – amount of time necessary to the processor to execute the task (without interruption)

Real-time tasks


Real-time task parameters



- Real-time tasks are characterized by further parameters, related to their *timing constraints*

Real-time tasks

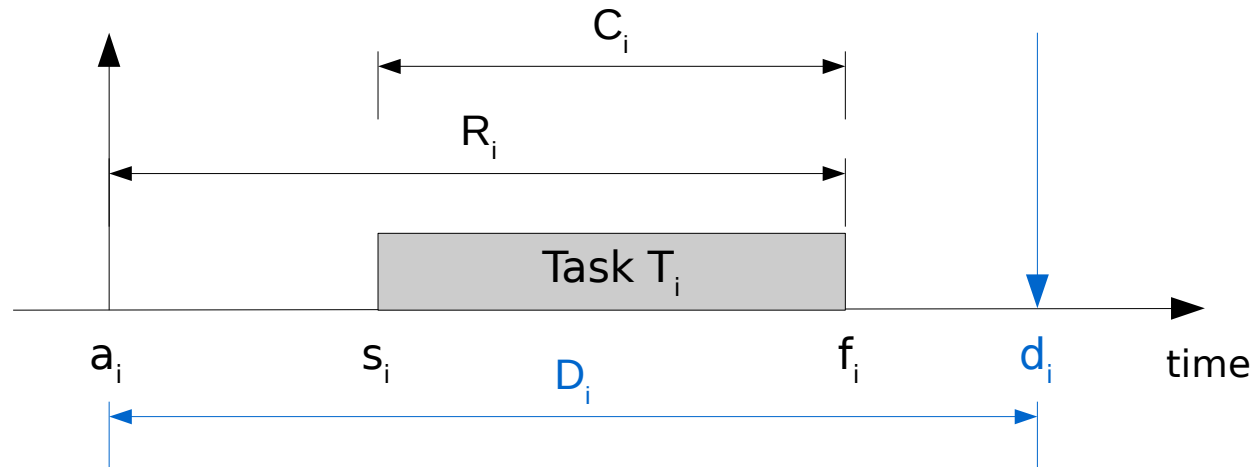
Real-time task parameters



- d_i : Absolute deadline – time before which the task should complete before violating the constraint
- D_i : Relative deadline – difference between absolute deadline and arrival time ($D_i = d_i - a_i$)

Real-time tasks

Real-time task parameters




- L_i : Lateness – delay of a task completion w.r.t. its deadline
($L_i = f_i - d_i$)
- X_i : Laxity (or Slack time) – the maximum amount of time a task can be delayed to complete within its deadline
($X_i = d_i - a_i - C_i$)


Real-time tasks

Task periodicity

- Periodic tasks




- Aperiodic tasks



Real-time tasks

Periodic tasks

- Most of the tasks of a control systems need to execute on a periodical time basis
 - Sensor data sampling
 - System monitoring
- Each periodic task τ_i is characterized by its period length T_i




Real-time tasks

Periodic tasks

- A task τ_i is characterized by several execution instances $\tau_{i,j}$
 - $r_{i,j}$ – the *arrival* (or *release*) time of the instance (when the task instance is ready to run)
 - ϕ_i – the release time of the first instance ($\phi_{i,j} = r_{i,0}$)
 - $s_{i,j}$ and $f_{i,j}$ are the *starting* and *finishing* times of the instances

$r_{i,j}$ and $s_{i,j}$ may differ, depending on the scheduler decisions

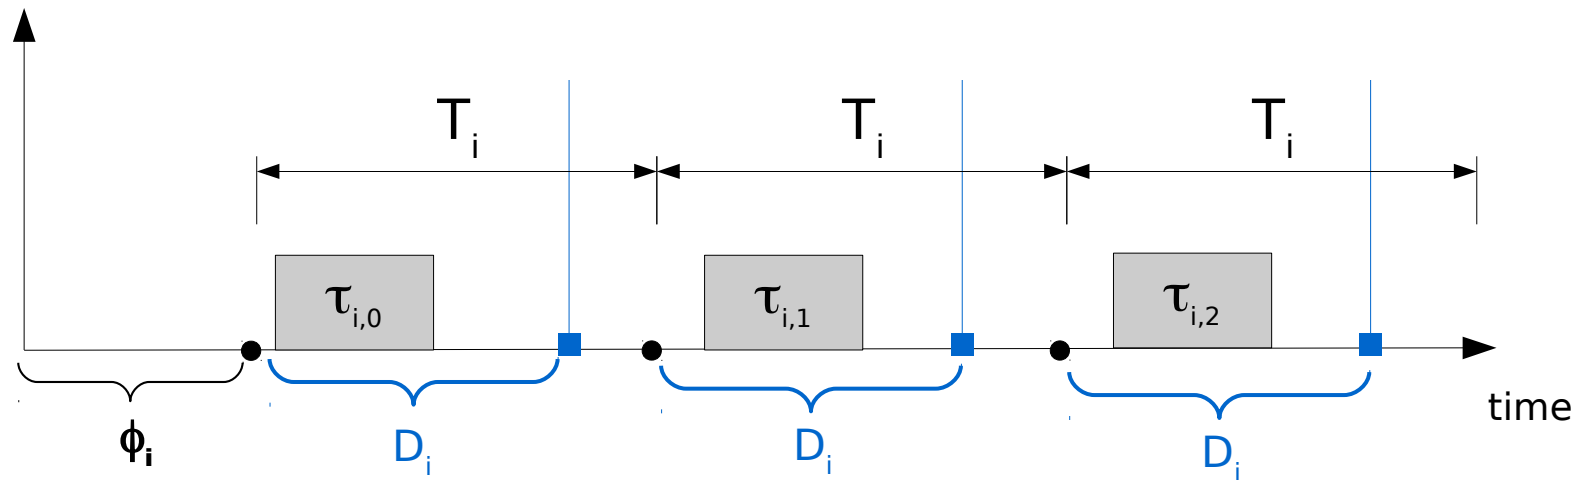


Real-time tasks

Periodic tasks

- A task τ_i is characterized by several execution instances $\tau_{i,j}$
 - $\mathbf{D_i}$ – relative deadline of the task within a period (the same for each instance)
 - $\mathbf{d_{i,j}}$ – absolute deadline (given a instance $\tau_{i,j}$):

$$d_{i,j} = \Phi_i + (j-1)T_i + D_i$$




Real-time tasks

Periodic tasks

- A task instance $\tau_{i,j}$ is also called *job*

→ Job response time: $R_{i,j} = f_{i,j} - r_{i,j}$

- Task response time: $R_i = \max_j R_{i,j}$



Real-time tasks


Periodic tasks

- *Relative start time jitter* – maximum deviation of start time between two consecutive instances/jobs

$$RRJ_i = \max_j |(s_{i,j} - r_{i,j}) - (s_{i,j-1} - r_{i,j-1})|$$

- *Absolute start time jitter* – maximum deviation... among all instances

$$ARJ_i = \max_j (s_{i,j} - r_{i,j}) - \min_j (s_{i,j} - r_{i,j})$$



Real-time tasks


Periodic tasks

- *Relative finishing time jitter* – maximum deviation of finishing time between two consecutive instances/jobs

$$RFJ_i = \max_j |(f_{i,j} - r_{i,j}) - (f_{i,j-1} - r_{i,j-1})|$$

- *Absolute finishing time jitter* – maximum deviation... among all instances

$$AFJ_i = \max_j (f_{i,j} - r_{i,j}) - \min_j (f_{i,j} - r_{i,j})$$



Real-time tasks

Periodic tasks

- When the system includes several concurrent periodic tasks with timing constraints, the scheduler provided by the operating system has to guarantee that...
 - Each task instance is activated (task is ready) at the proper rate
 - The task instance is completed within the deadline D_i



- This is what a real-time scheduler is intended for

Real-time scheduling

Outline

- Definitions and Assumptions
- Algorithms

Real-time scheduling

Definitions

- Task scheduling plays a key role in real-time systems
- Real-time scheduling means that the algorithms must take into account the *timing constraints* of each task
 - A couple of definitions are worth to be mentioned...
- A schedule is *feasible*, if all the tasks of the given set can be completed without timing constraints violations
- A set of tasks is *schedulable* if there exists at least one algorithm capable of producing a feasible scheduling

Real-time scheduling

Assumptions

- [1] The instances of a periodic task τ_i are regularly activated at constant rate
- [2] All instances of a periodic task τ_i have the worst-case execution time C_i
- [3] All instances of a periodic task τ_i have the same deadline D_i equal to the period ($D_i = T_i$)
- [4] All the tasks in a set Γ are independent (no precedence relations and no resource constraints)
- [5] Tasks cannot suspend themselves (e.g., on I/O operations)
- [6] Release time coincides with the arrival of the task
- [7] Kernel overheads (e.g., context switch) are assumed to be zero


Real-time scheduling

Processor utilization factor

- The fraction of the processor time spent in the execution of the tasks from a set Γ

$$U = \sum_{i=0}^n \frac{C_i}{T_i}$$

- A measure of the computational load of the processor



Real-time scheduling

Processor utilization factor

- U is useful to check the feasibility of a schedule
- There exists an **upper bound value** (U_{ub}) above which a schedule is not feasible
- The upper bound depends on the task set Γ and the scheduling algorithm: $U_{ub}(\Gamma, A)$
- Therefore, a task set Γ_i is schedulable with A iff

$$U_i \leq U_{ub}(\Gamma_i, A)$$


- Given the task set Γ the fully utilization of the processor is achieved for $U = U_{ub}(\Gamma, A)$

Real-time scheduling

Processor utilization factor

- Given an algorithm A, we can also define the **least upper bound** $U_{lub}(A)$ as the minimum utilization factor over all the task sets that fully utilize the processor


$$U_{lub}(A) = \min_{\Gamma} U_{ub}(\Gamma, A)$$



Real-time scheduling

Processor utilization factor

- For each task set Γ_i
 - If $U_i < U_{lub}$ $\Rightarrow \Gamma_i$ schedulable by the algorithm
 - If $U_{lub} < U_i \leq 1$ $\Rightarrow \Gamma_i$ schedulable if periods suitably related
 - If $U_i > 1$ $\Rightarrow \Gamma_i$ not schedulable



Real-time scheduling


Scheduling algorithms

- Timeline Scheduling (TS)
- Rate Monotonic (RM)
- Earliest Deadline First (EDF)
- Deadline Monotonic (DM)

Real-time scheduling

Timeline Scheduling

- Also known as *Cyclic Executive* scheduling
- Very used in periodic tasks scheduling in military and traffic control systems
- The requirement is to execute a task at a given rate
- The temporal axis is divided into equal length slots
 - One or more tasks allocated per slot
 - A time synchronizes the activation of the tasks at the beginning of the slot



Real-time scheduling

Timeline Scheduling

- *Example*

Task	Rate (Hz)	Period (ms)
A	40	25
B	20	50
C	10	100

- The optimal slot length would be $T_A = 25\text{ms}$ (the Greatest Common Divisor)
 - Task A will execute in every slot
 - Task B will execute every two slots
 - Task C will execute every four slots


Real-time scheduling

Timeline Scheduling

- *Example*

Task	Rate (Hz)	Period (ms)
A	40	25
B	20	50
C	10	100


→ Possible solution:



Real-time scheduling

Timeline Scheduling

- The duration of a time slot is also called minor cycle
 - *Greatest Common Divisor (GCD)* of the activation periods
- A major cycle instead is the time interval after which a schedule repeats itself (also called *hyper-period*)
 - *Least Common Multiple (lcm)* of the activation periods



Real-time scheduling

Timeline Scheduling

- Schedulability
 - Verify that the sum of the worst-execution times (WCET) of the tasks within each time slot is less than or equal to the Minor Cycle
 - Remember assumption [2]: $WCET_i = C_i$

$$\begin{cases} C_A + C_B \leq 25\text{ ms} \\ C_A + C_C \leq 25\text{ ms} \end{cases}$$

Real-time scheduling

Timeline Scheduling

- PROs
 - *Simple implementation*

Timer interrupt set to the minor cycle length and a main program calling the tasks in the given order
 - *Very low overhead*

The task execution order is not decided by a scheduler but by the program
 - *No jitter*

Always the same sequence of task executions
Start times and response times not subject to noticeable variations

Real-time scheduling

Timeline Scheduling

- CONs

- *Not robust*

- In overloaded conditions, a deadline miss causes a domino effect, breaking the schedule

- *Sensitivity to application changes*

- Updates of computation time or activation rate may invalidate the current schedule

- *Hard to handle aperiodic activities*

- May require a change in the task execution sequence

Real-time scheduling

Rate Monotonic

- Assign a priority to the task according to the **activation (or release) rate**
 - Higher rates (shorter periods) → higher priorities
- *Fixed-priority* assignment
 - Period length is constant periods
 - Priority assigned when task is ready and no longer changed
- *Preemptive* (intrinsically)
 - A new task preempts the currently running one if characterized by a shorter period
- Optimal algorithm among all the fixed-priority based scheduling algorithms
 - Liu and Layland (1973)

Real-time scheduling

Rate Monotonic

- Schedulability
 - It can be demonstrated that for RM the least upper bound utilization factor converges to

$$U_{lub} = n(2^{1/n} - 1)$$

$$n = 2 : \quad U_{lub} = 0.828$$

$$\lim_{n \rightarrow \infty} U_{lub}(n) = \ln(2) = 0.693$$

- The task set is schedulable if

$$U_i \leq U_{lub}$$

Real-time scheduling

Rate Monotonic

- Example

Task	Period T	WCET C
A	10	2
B	60	10
C	30	5

- Scheduler will be executed every minor cycle
 - Minor cycle = $\text{GCD}(10, 60, 30) = 10$
- The schedule length is the given by the major cycle
 - Major cycle = $\text{lcm}(10, 60, 30) = 60$

Real-time scheduling

Rate Monotonic

- Example

Task	Period T	WCET C
A	10	2
B	60	10
C	30	5

- Schedulability

- $U_{lub} = 3(2^{1/3} - 1) = 0.7798$
- $U_i = (2/10) + (10/60) + (5/30) = 0.533$
- $U_i \leq U_{lub}$ **OK!**

Real-time scheduling

Rate Monotonic

- Example

Task	Period T	WCET C
A	10	2
B	60	10
C	30	5


- Reminding that the shorter the period, the higher the priority...
 - $p_A > p_C > p_B$

Real-time scheduling

Rate Monotonic

- Example
 - Tasks are scheduled according to the priority: $\tau_A > \tau_C < \tau_B$
 - $t=10$: τ_B is preempted by τ_A
 - $t=12$: τ_A terminates and τ_B can be resumed to complete its execution

Task	Period T	WCET C
A	10	2
B	60	10
C	30	5




Real-time scheduling

Rate Monotonic

- Example
 - $t=20$: τ_A must be executed again
 - τ_B and τ_C has already completed the execution in the respective periods

Task	Period T	WCET C
A	10	2
B	60	10
C	30	5




Real-time scheduling

Rate Monotonic

- Example
 - $t=30$: τ_A must be executed again
 - $t=32$: τ_C must be executed again

Task	Period T	WCET C
A	10	2
B	60	10
C	30	5




Real-time scheduling

Rate Monotonic

- Example
 - τ_A will be executed again every 10 time units
 - After 60 time units (mayor cycle) the schedule will repeat itself

Task	Period T	WCET C
A	10	2
B	60	10
C	30	5



Real-time scheduling

Rate Monotonic

- Properties
 - Easy to implement algorithm
 - Lower overhead w.r.t EDF
 - Fixed priorities → no need of updating
 - Stability is easier to achieve w.r.t other scheduler, in case of system problems(e.g. due to transient errors),
 - Actually exploited in industrial solutions
 - Versions with relaxed assumptions available (we will not shown them)

Real-time scheduling

Earliest Deadline First (EDF)

- *Dynamic-priority* assignment scheduler
- Task to execute is selected according to the **absolute deadline**
 - Higher priorities to earlier deadline tasks

$$d_{i,j} = \Phi_i + (j-1)T_i + D_i$$

- *Preemptive*
 - Currently running task preempted when a periodic instance with earlier deadline is activated
- Agnostic w.r.t to the task periodicity
 - Both periodic or aperiodic tasks can be scheduled

Real-time scheduling

Earliest Deadline First (EDF)

- EDF is optimal (for single core processors)

Theorem

*If a set of jobs is schedulable by an algorithm A,
then it is schedulable by EDF*

Real-time scheduling

Earliest Deadline First (EDF)

- Schedulability
 - A set of periodic tasks is schedulable with EDF iff

$$U = \sum_{i=0}^n \frac{C_i}{T_i} \leq 1$$

i.e., the processor utilization is not greater than 100%

Real-time scheduling

Earliest Deadline First (EDF)

- Example

- Release times and deadlines coincide ($T = D$)

Task	Relative deadline D	WCET C
A	4	1
B	6	2
C	12	4

- Schedulability

$$U = (1/4) + (2/6) + (4/12) = \mathbf{0.9167} < \mathbf{1} \rightarrow \mathbf{YES!}$$

- (no deadline misses expected)

Real-time scheduling


Earliest Deadline First (EDF)

- Example

- Let's annotate release times and deadlines

- $\text{lcm}(4,6,12) = \mathbf{12}$

Task	Relative deadline D	WCET C
A	4	1
B	6	2
C	12	4




Real-time scheduling

Earliest Deadline First (EDF)

- Example

- **t=0**: Closest deadline is d_A → schedule τ_A

Task	Relative deadline D	WCET C
A	4	1
B	6	2
C	12	4




Real-time scheduling

Earliest Deadline First (EDF)

- Example

- $t=1$: τ_A already executed
- Closest deadline $d_B \rightarrow$ schedule τ_B

Task	Relative deadline D	WCET C
A	4	1
B	6	2
C	12	4



Real-time scheduling


Earliest Deadline First (EDF)

- Example

- **t=3** :

- τ_C is the only ready task → Schedule τ_C

Task	Relative deadline D	WCET C
A	4	1
B	6	2
C	12	4




Real-time scheduling

Earliest Deadline First (EDF)

- Example
 - $t=4$: τ_C preempted by τ_A

Task	Relative deadline D	WCET C
A	4	1
B	6	2
C	12	4



Real-time scheduling


Earliest Deadline First (EDF)

- Example

- **t=5** :

- τ_C is resumed as it is the only ready task

Task	Relative deadline D	WCET C
A	4	1
B	6	2
C	12	4



Real-time scheduling


Earliest Deadline First (EDF)

- Example

- **t=6** :

- τ_C is kept in execution, since τ_B has the same distance from deadline

Task	Relative deadline D	WCET C
A	4	1
B	6	2
C	12	4



Real-time scheduling


Earliest Deadline First (EDF)

- Example

- **t=8** :

- τ_A and τ_B are active and scheduled in sequence since $d_A = d_B$

Task	Relative deadline D	WCET C
A	4	1
B	6	2
C	12	4



Real-time scheduling


Earliest Deadline First (EDF)

- Example

- **t=8** :

- τ_A and τ_B are active and scheduled in sequence since $d_A = d_B$

Task	Relative deadline D	WCET C
A	4	1
B	6	2
C	12	4




Real-time scheduling

Earliest Deadline First (EDF)

- Example

- $\text{lcm}(4,6,12) = 12$
- As expected, the schedule repeat itself after $t=12$

Task	Relative deadline D	WCET C
A	4	1
B	6	2
C	12	4



Real-time scheduling

Earliest Deadline First (EDF)

- With EDF we can achieve a higher processor utilization w.r.t. Rate Monotonic
- Less preemptions w.r.t Rate Monotonic, due to dynamic priorities
- But also... higher overhead w.r.t. Rate Monotonic
 - Task priority is dynamic → it may change online
- Hard and expensive to implement on micro-controller based embedded systems
 - Require very frequent invocations of the scheduler
 - Hard to predict overload conditions of the system
- Seldom adopted in industrial applications

Real-time scheduling

Deadline Monotonic (DM)

- Proposed a generalization of Rate Monotonic by Audsley (1990)
 - Assumption [3] ($D_i = T_i$) is relaxed
- *Fixed priority* assignment
 - Relative deadline is constant
 - Priority inversely proportional to relative deadline D_i
- *Preemptive*
 - Currently running task preempted when a new task with shorter relative deadline arrives
- DM is optimal
 - If a task set is schedulable by some fixed priority static algorithm, it is also schedulable by DM
 - If a task set is NOT schedulable by DM, it is not schedulable by any other fixed priority static algorithm

Real-time scheduling

Deadline Monotonic (DM)

- Schedulability
 - A “pessimistic” test derived from Rate Monotonic (processor utilization overestimated)

$$\sum_{i=0}^n \frac{C_i}{D_i} \leq n(2^{(1/n)} - 1)$$

In DM must consider relative deadlines instead of activation periods


- A more efficient (*sufficient and necessary*) test has been proposed by Audsley (1993) called Response Time Analysis

Real-time scheduling

Deadline Monotonic (DM)

- Example
 - Assumption [3] relaxed
 - $\text{lcm}(5,10,15) = 30$
 - Let's annotate release times and deadlines

Task	Period T	Relative deadline D	WCET C
A	5	3	1
B	10	7	3
C	15	10	5



Real-time scheduling


Deadline Monotonic (DM)

- Example

- Fixed priorities assignment (based on D_i):

$$p_A > p_B > p_C$$

Task	Period T	Relative deadline D	WCET C
A	5	3	1
B	10	7	3
C	15	10	5



Real-time scheduling

Deadline Monotonic (DM)

- Example


- Priorities:

$$p_A > p_B > p_C$$

- $t=0$: schedule τ_A

- $t=1$: schedule τ_B

Task	Period T	Relative deadline D	WCET C
A	5	3	1
B	10	7	3
C	15	10	5



Real-time scheduling

Deadline Monotonic (DM)

- Example


- Priorities:

$$p_A > p_B > p_C$$

- $t=4$: schedule τ_C

- $t=5$: τ_C is preempted by the activation of τ_A

Task	Period T	Relative deadline D	WCET C
A	5	3	1
B	10	7	3
C	15	10	5



Real-time scheduling

Deadline Monotonic (DM)


- Example

- Priorities:

$$p_A > p_B > p_C$$

- $t=6$: schedule τ_C which terminates the instance execution

Task	Period T	Relative deadline D	WCET C
A	5	3	1
B	10	7	3
C	15	10	5



Real-time scheduling

Deadline Monotonic (DM)

▪ Example


→ Priorities:

$$p_A > p_B > p_C$$

→ $t=10$: τ_A and τ_B
activated

→ τ_A scheduled due to
higher priority

Task	Period T	Relative deadline D	WCET C
A	5	3	1
B	10	7	3
C	15	10	5



Real-time scheduling

Deadline Monotonic (DM)

- Example


- Priorities:

$$p_A > p_B > p_C$$

- $t=11$: schedule τ_B

- $t=14$: no active instances → nothing to schedule

Task	Period T	Relative deadline D	WCET C
A	5	3	1
B	10	7	3
C	15	10	5



Real-time scheduling

Deadline Monotonic (DM)


- Example

- Priorities:

$$p_A > p_B > p_C$$

- $t=14$: no active instances → nothing to schedule

Task	Period T	Relative deadline D	WCET C
A	5	3	1
B	10	7	3
C	15	10	5



Real-time scheduling

Deadline Monotonic (DM)

- Example


- Priorities:

$$p_A > p_B > p_C$$

- $t=15$: τ_A and τ_C activated

- τ_A scheduled due to higher priority

Task	Period T	Relative deadline D	WCET C
A	5	3	1
B	10	7	3
C	15	10	5



Real-time scheduling

Deadline Monotonic (DM)


- Example

- Priorities:

$$p_A > p_B > p_C$$

- $t=16$: τ_C scheduled since τ_B already executed

Task	Period T	Relative deadline D	WCET C
A	5	3	1
B	10	7	3
C	15	10	5



Real-time scheduling

Deadline Monotonic (DM)


- Example

- Priorities:

$$p_A > p_B > p_C$$

- $t=20$: τ_C preempted by τ_A

Task	Period T	Relative deadline D	WCET C
A	5	3	1
B	10	7	3
C	15	10	5



Real-time scheduling

Deadline Monotonic (DM)


- Example

- Priorities:

$$p_A > p_B > p_C$$

- $t=21$: τ_B scheduled

Task	Period T	Relative deadline D	WCET C
A	5	3	1
B	10	7	3
C	15	10	5



Real-time scheduling

Deadline Monotonic (DM)


- Example

- Priorities:

$$p_A > p_B > p_C$$

- $t=24$: τ_C resumed

Task	Period T	Relative deadline D	WCET C
A	5	3	1
B	10	7	3
C	15	10	5



Real-time scheduling

Deadline Monotonic (DM)

▪ Example


→ Priorities:

$$p_A > p_B > p_C$$

→ $t=25$: τ_A activated and scheduled

→ Mayor cycle completed

Task	Period T	Relative deadline D	WCET C
A	5	3	1
B	10	7	3
C	15	10	5



Real-time scheduling

Deadline Monotonic (DM)

- Observation
 - DM can schedule periodic tasks without the need of knowing the release times a priori (differently from Rate Monotonic)
Priority does not depend on T_i

Questions?

Contacts

- William Fornaciari
<https://home.deib.polimi.it/fornacia>
william.fornaciari@polimi.it
- **HEAPLab**
DEIB Politecnico di Milano
Building 21, Floor 1
Phone: 9613 / 9623
- **Slides credits:**
 - Giuseppe Massari <giuseppe.massari@polimi.it>

