



POLITECNICO
MILANO 1863



POLITECNICO
MILANO 1863

Software Engineering 2

Introduction to Architectural Styles

Client-server

Interface Design and Documentation

Designing servers to handle multiple requests

Other styles: Multi-tier, Event-driven



POLITECNICO
MILANO 1863

Software Design

Introduction to Architectural Styles


Client-server

Architectural Style

- “an architectural style determines the **vocabulary** of **components** and **connectors** that can be used in instances of that style, together with a set of **constraints** on how they can be combined.
These can include topological constraints on architectural descriptions (e.g., no cycles). Other constraints—say, having to do with execution semantics—might also be part of the style definition.” [Garland & Shaw]

[Garland & Shaw] David Garlan and Mary Shaw, “An Introduction to Software Architecture”, Jan 1994, CMU-CS-94-166. http://www.cs.cmu.edu/afs/cs/project/able/ftp/intro_softarch/intro_softarch.pdf

Client-server



- Two component roles
 - Client issues requests
 - Server provides responses

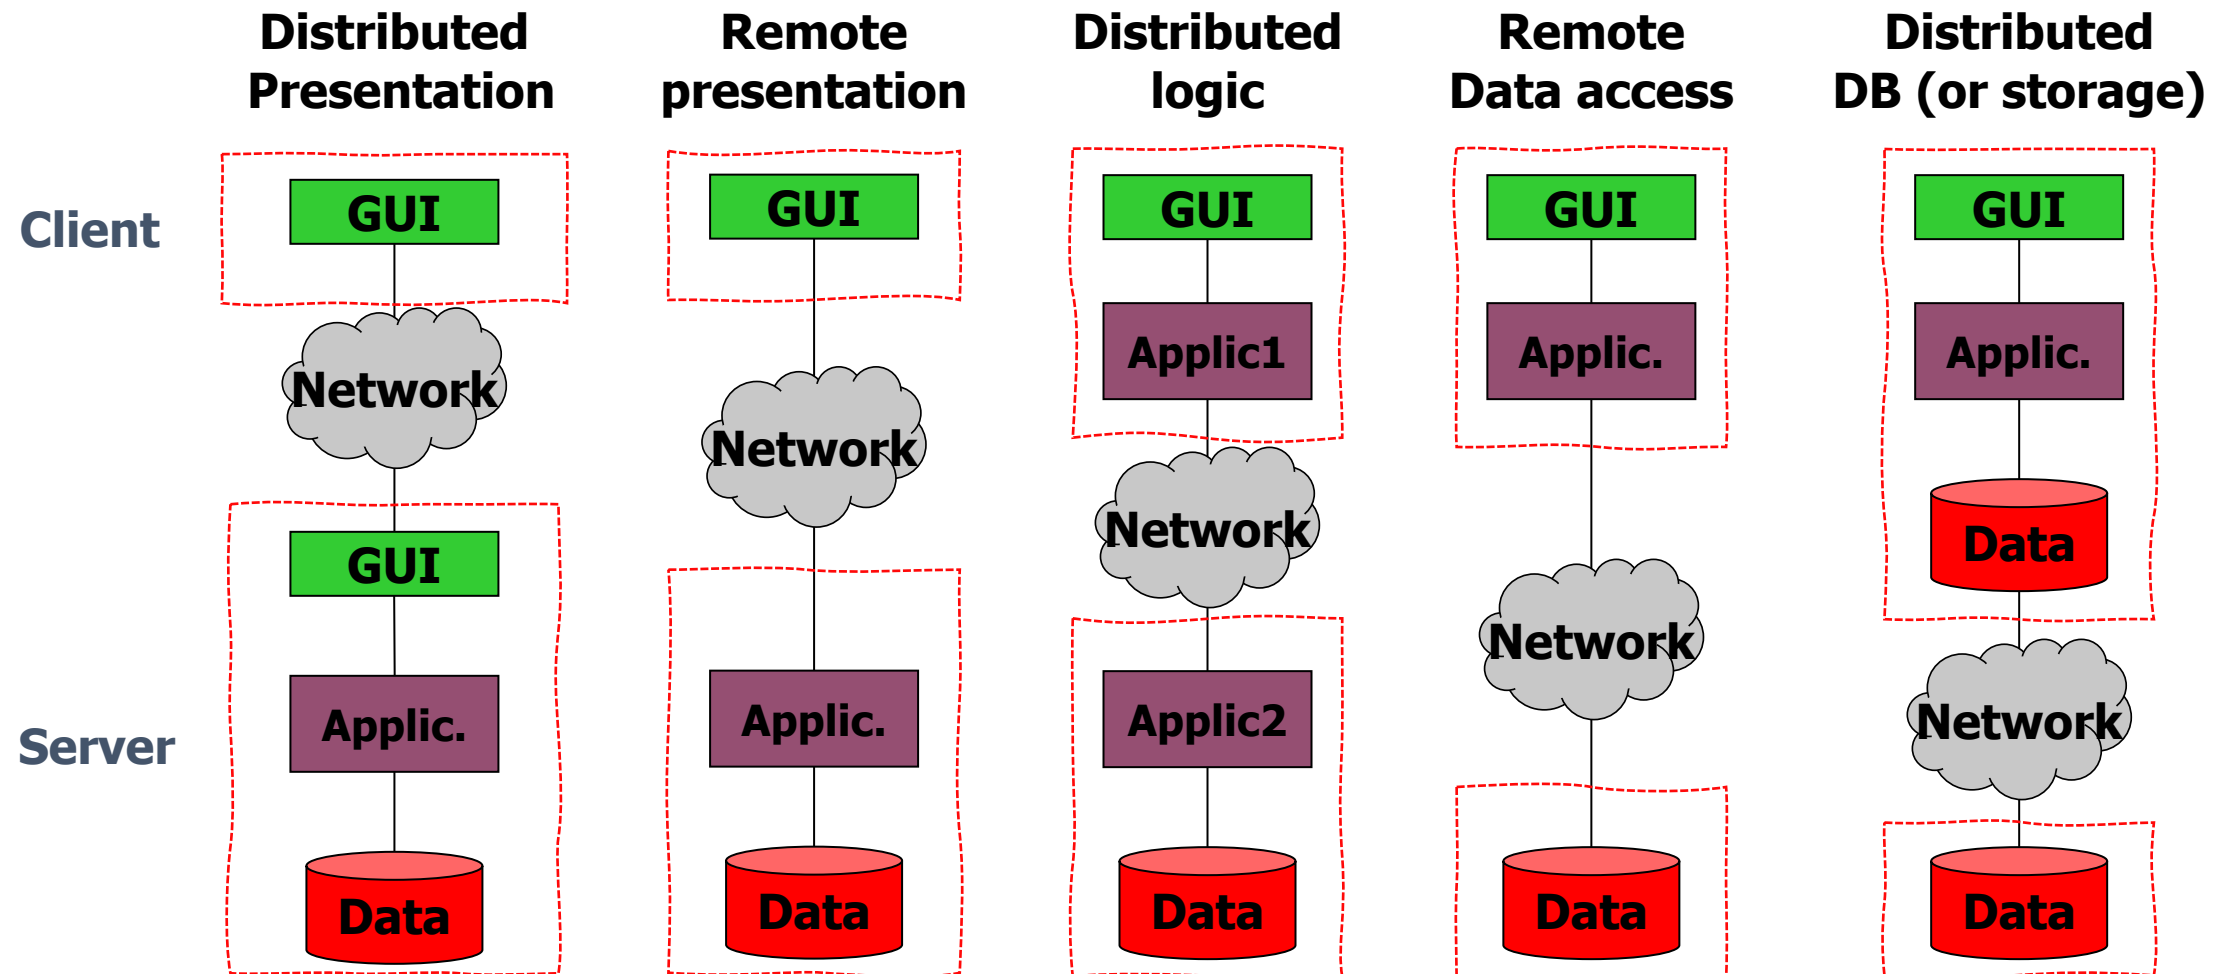
Why do I use it?

- **Multiple users** need to access a **single resource**
 - e.g., database
- There is a preexisting software we must **access remotely**
 - e.g., email server
- It is convenient to organize the system around a **shared piece of functionality** used by multiple components
 - e.g., authentication/authorization server

Organization of Client-Server software: thin vs fat clients



POLITECNICO
MILANO 1863





Client-server: main technical issues

- **Design** and **document** proper interfaces for our server
- Ensure the server is able to **handle multiple simultaneous requests**
 - Forking vs thread pooling



POLITECNICO
MILANO 1863

Software Design

Interface Design



Interface design

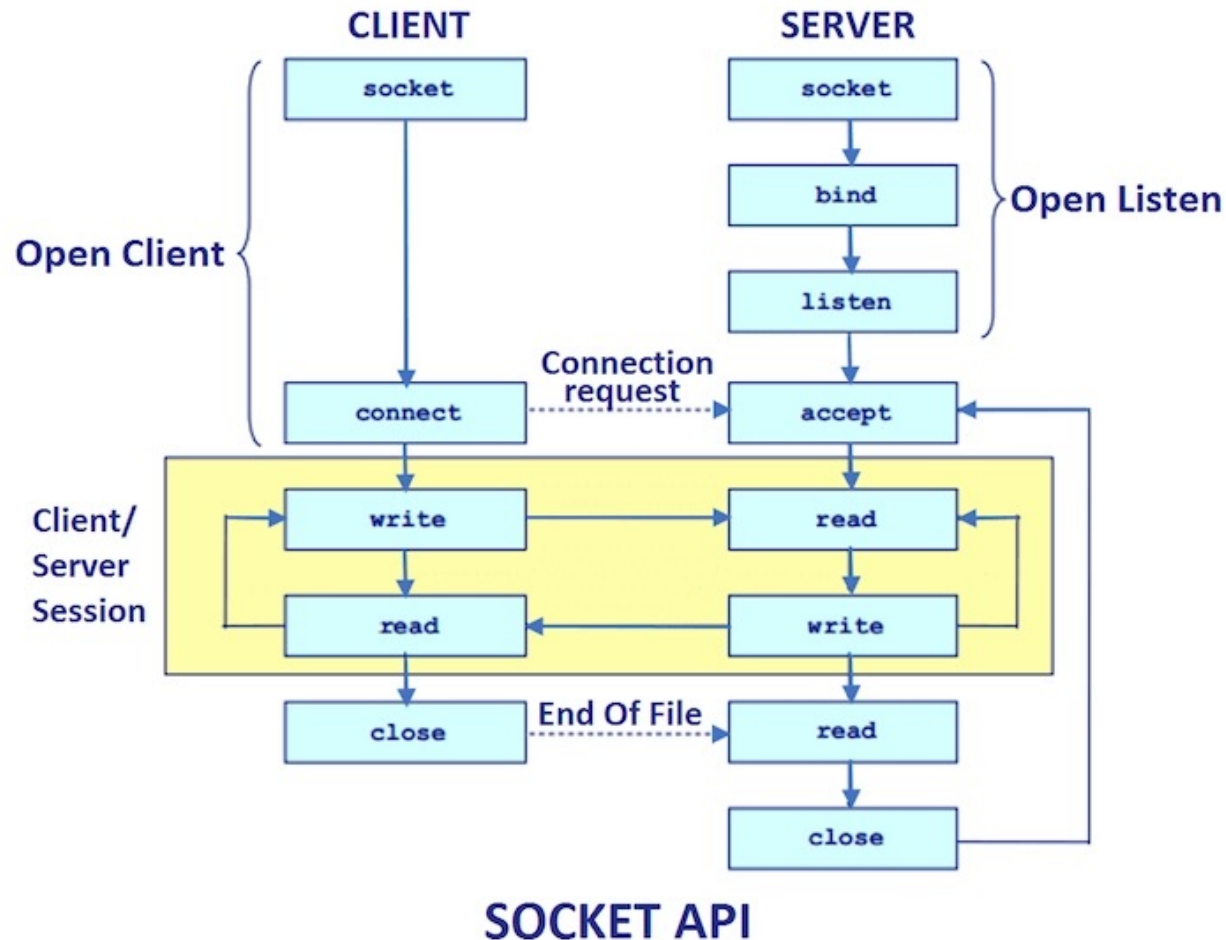
- An interface is a **boundary** across which components interact
- **Proper definition** of interfaces is an architectural concern
 - Impacts maintainability, usability, testability, performance, integrability
- Two important **guiding principles**
 - **Information hiding**
 - **Low coupling**
- An interface shall encapsulate a component implementation so that this can be changed without affecting other components

Aspects to consider during interface design

- **Contract principle**: Any resource (operation, data) added to an interface implies a commitment to maintaining it
- **Least surprise principle**: Interfaces should behave consistently with expectations
- **Small interfaces principle**: Interfaces should limit the exposed resources to the minimum
- Important elements to be defined:
 - **Interaction style** (e.g., sockets, RPC, REST)
 - **Representation** and structure of exchanged data
 - **Error handling**

Sockets

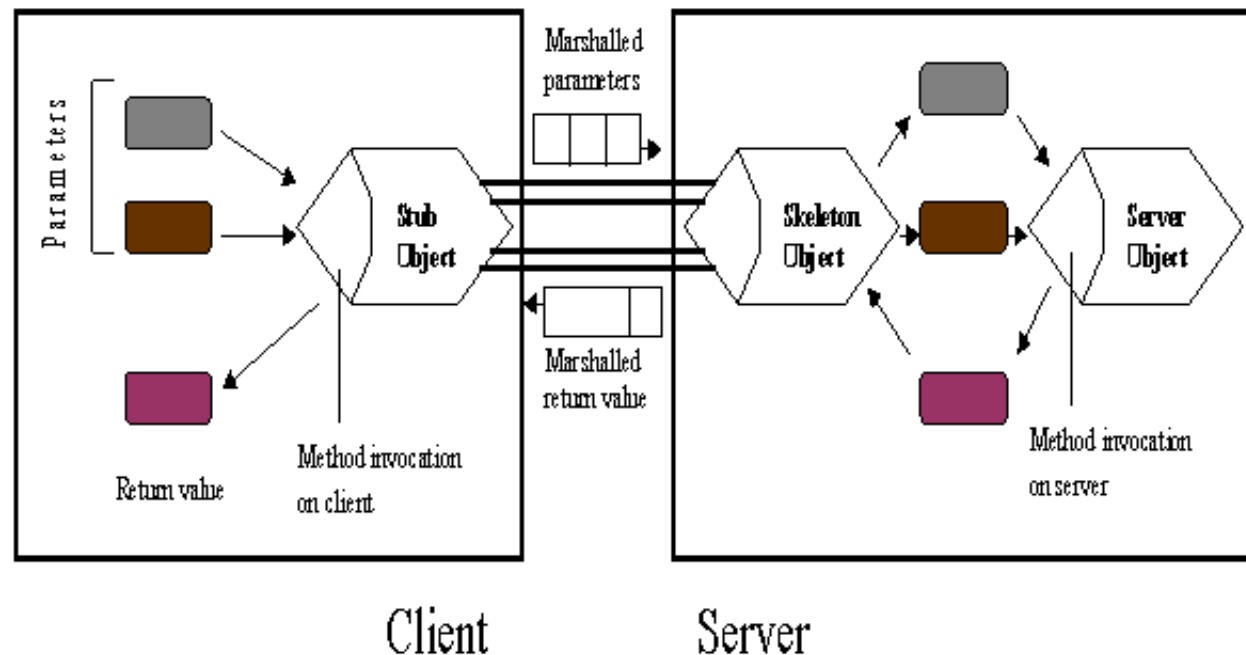
<https://www.codingninjas.com/studio/library/socket-programming>



- After connection establishment communication is bidirectional
- Both parties must agree on the same protocol (sequence of message exchange)

RPC/RMI

http://infolab.stanford.edu/CHAIMS/Doc/Details/Protocols/rmi/rmi_description.html



- Resembles procedure/method call in a centralized setting
- Stubs and skeletons needed to transform procedure/method calls in messages and vice versa

Communication in RMI: back to the banking example



POLITECNICO
MILANO 1863

github.com/matteocamilli/bank_example_rmi

```
import java.rmi.*;

public class Client {

    public static void main (String[] args) {

        try {
            //...

            String name = "Jack B. Quick";
            Account account = manager.open(name);
            float balance = account.balance();
            //...
        } catch (Exception e) {e.printStackTrace();}

    }
}
```

manager and account are local stubs
offering the same interface as the
corresponding remote objects

```
public class AccountManagerImpl ... {

    public Account open(String name) ... {
        //... }

    public static void main(String[] args)
    { //...

        AccountManagerImpl server = new
            AccountManagerImpl(_hostname);

        Naming.rebind("//" + _hostname +
            "/AccountManager", server);

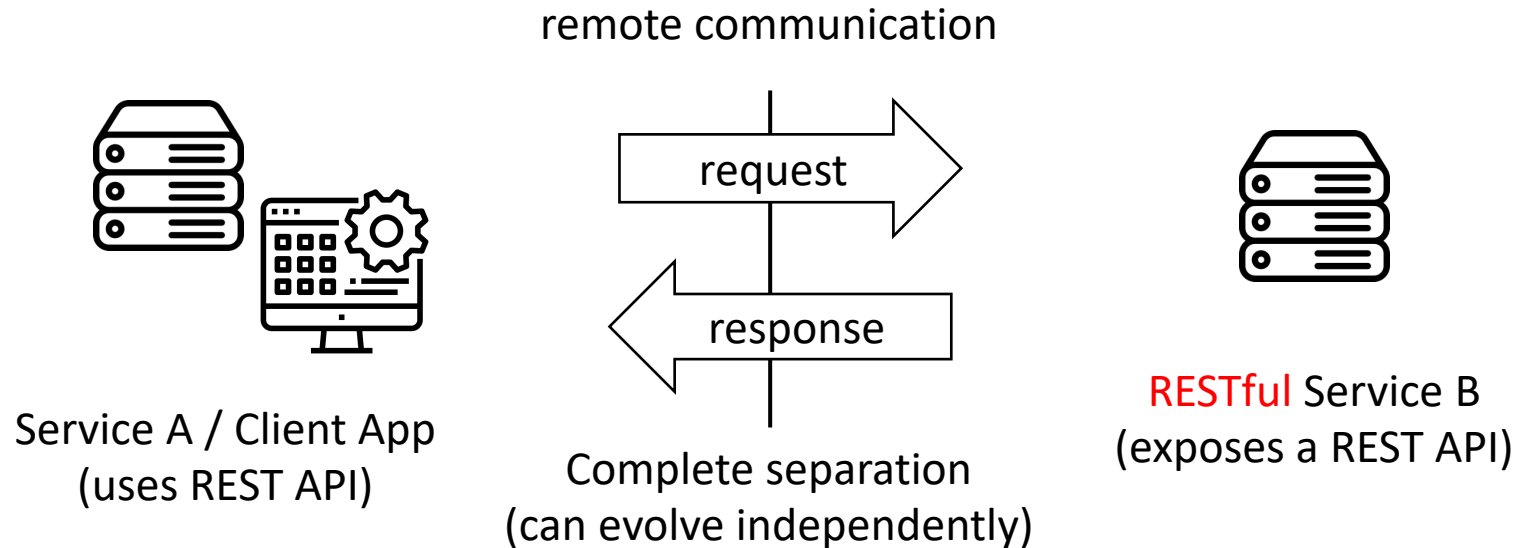
        //...
    }
}
```

server is the remote object that will execute
the open operation

REST: REpresentational State Transfer



POLITECNICO
MILANO 1863



- **REST = REpresentational State Transfer**
 - Specific standardized architectural style for Application Programming Interfaces (APIs)
 - Realizes clear separation between distributed, heterogeneous systems/components

Roy T. Fielding and Richard N. Taylor. 2000. Principled design of the modern Web architecture. ICSE '00. ACM, New York, NY, USA, 407–416. <https://doi.org/10.1145/337180.337228>

REST APIs

- **Characteristics**

- Simple and **standardized** → Developers do not have to worry about
 - Communication protocols (HTTP)
 - how to format the data (JSON)
 - how to format requests (request/response encoding)
 - Ready-to-use industry standard
- **Stateless** → does not keep track of states across server and all clients
 - Lightweight and **scalable**
 - Supports **caching** → high performance


REST APIs: requests

Known endpoint (public URL that identifies a resource)

`https://petstore.swagger.io/v2/pet/findByStatus`

API v2 of
the endpoint

Resource



- **Types of requests (CRUD)**

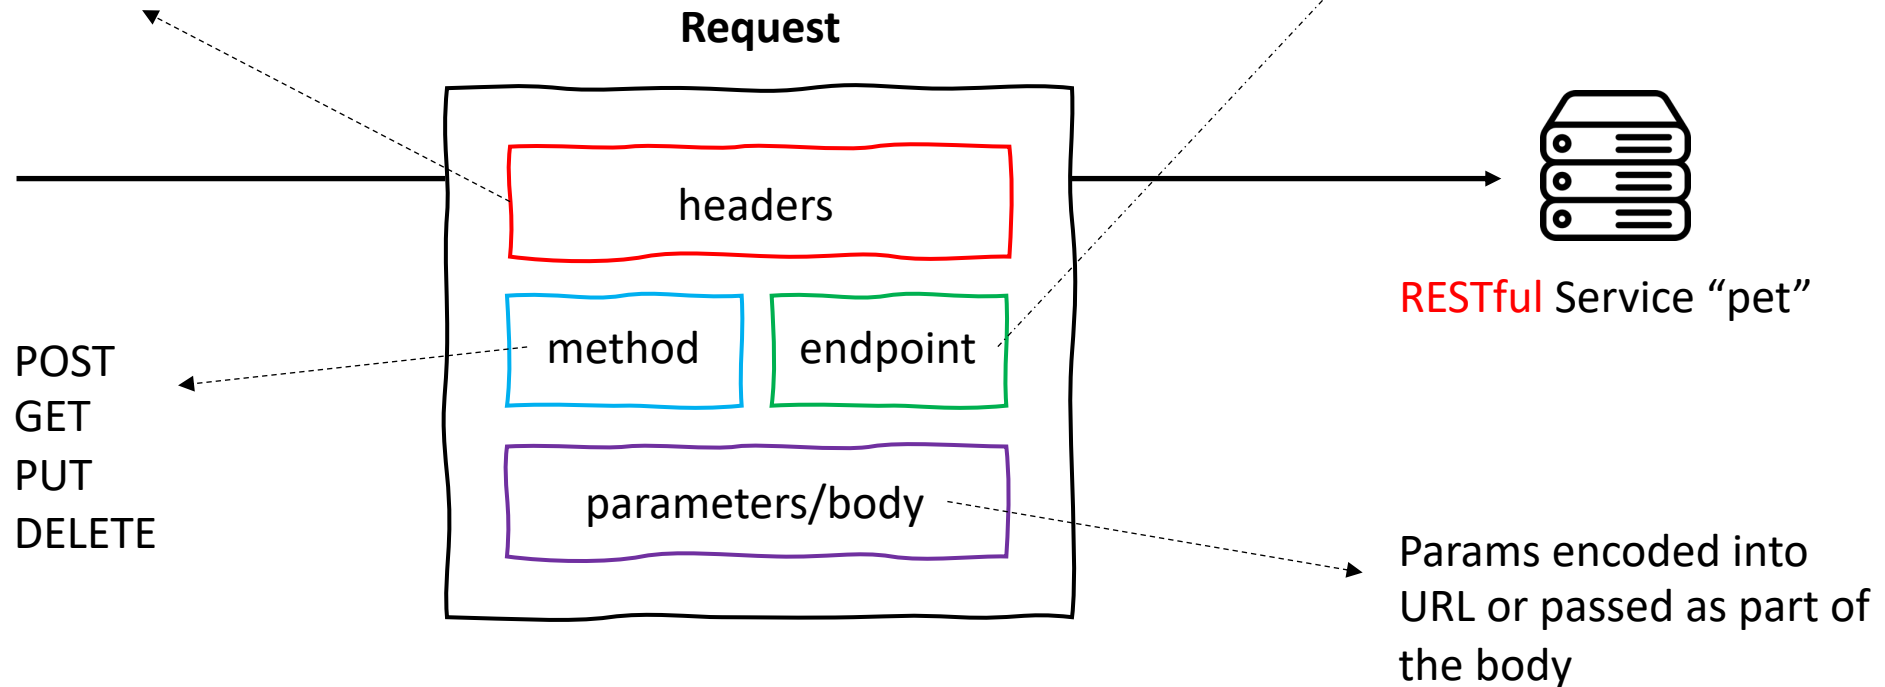
- Create (create a new resource) → POST
- Read (read an existing resource) → GET
- Update (modify a resource) → PUT
- Delete (remove a resource) → DELETE

standard HTTP
methods

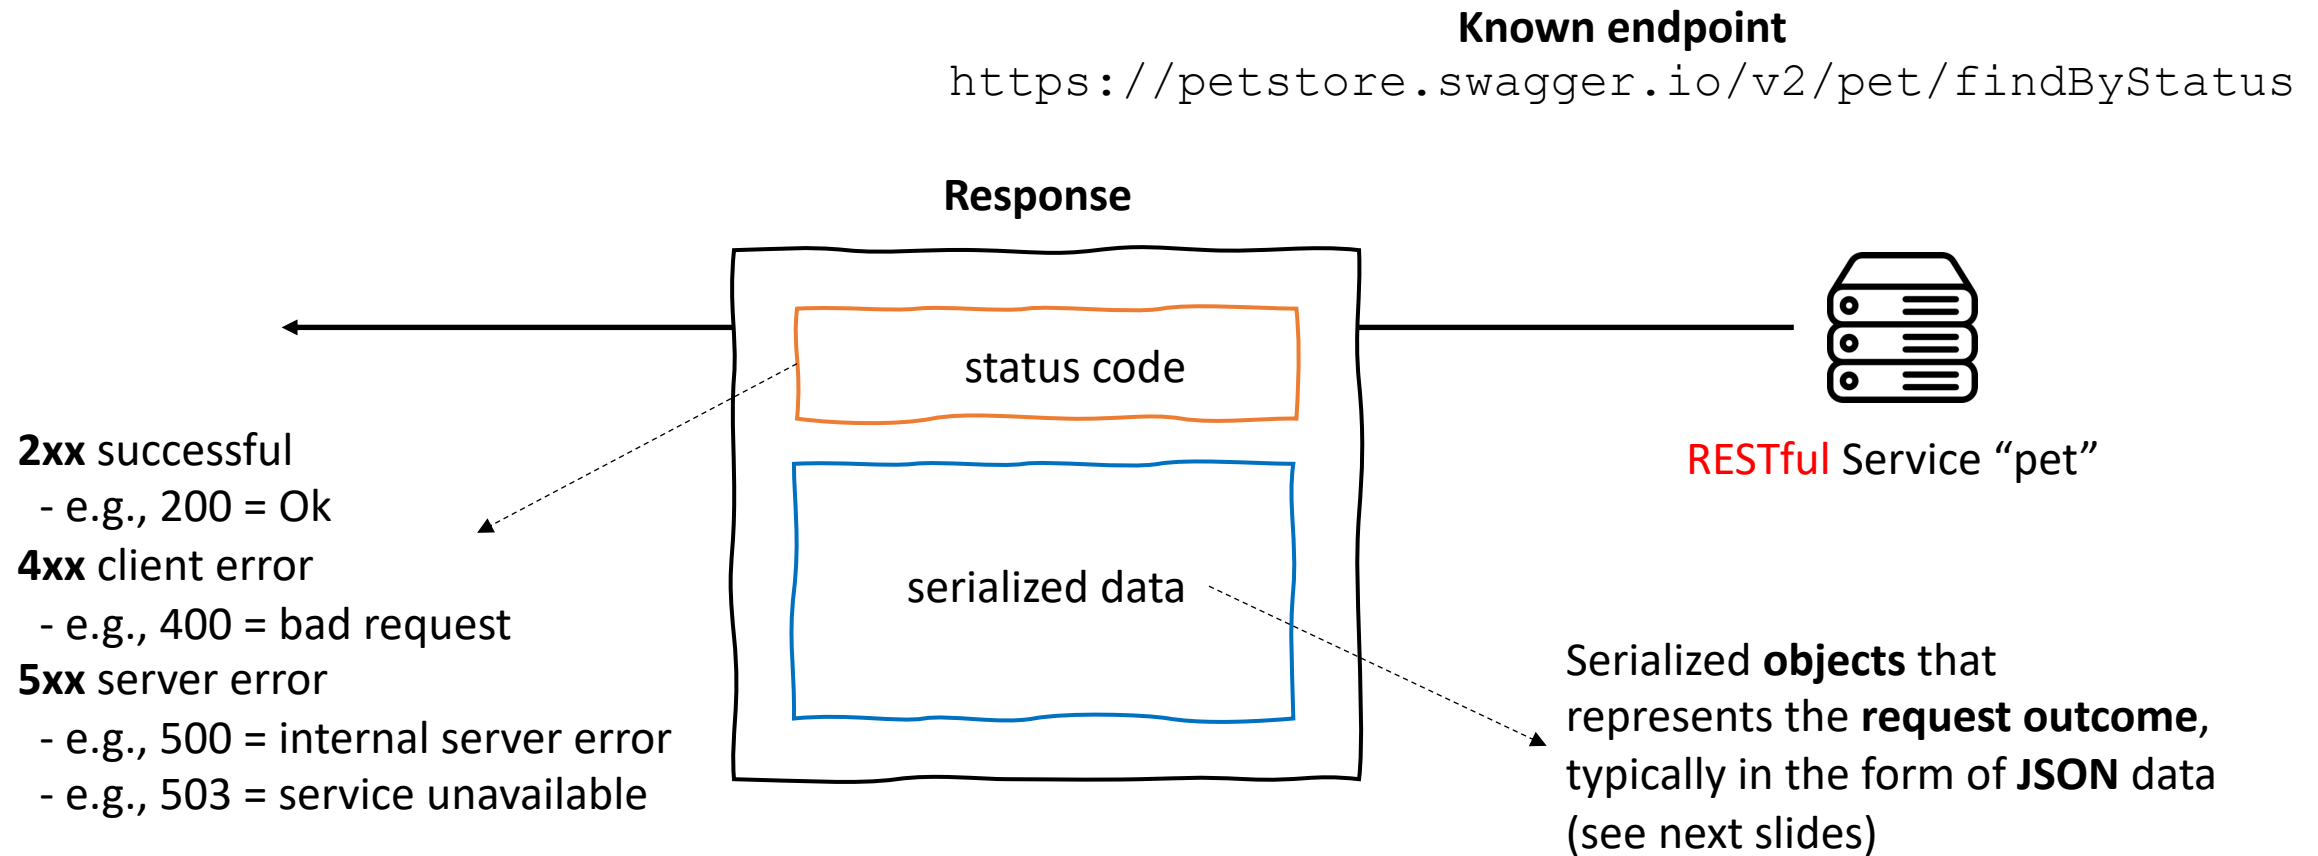
REST APIs: requests

“Special part” contains additional info like auth data and API key

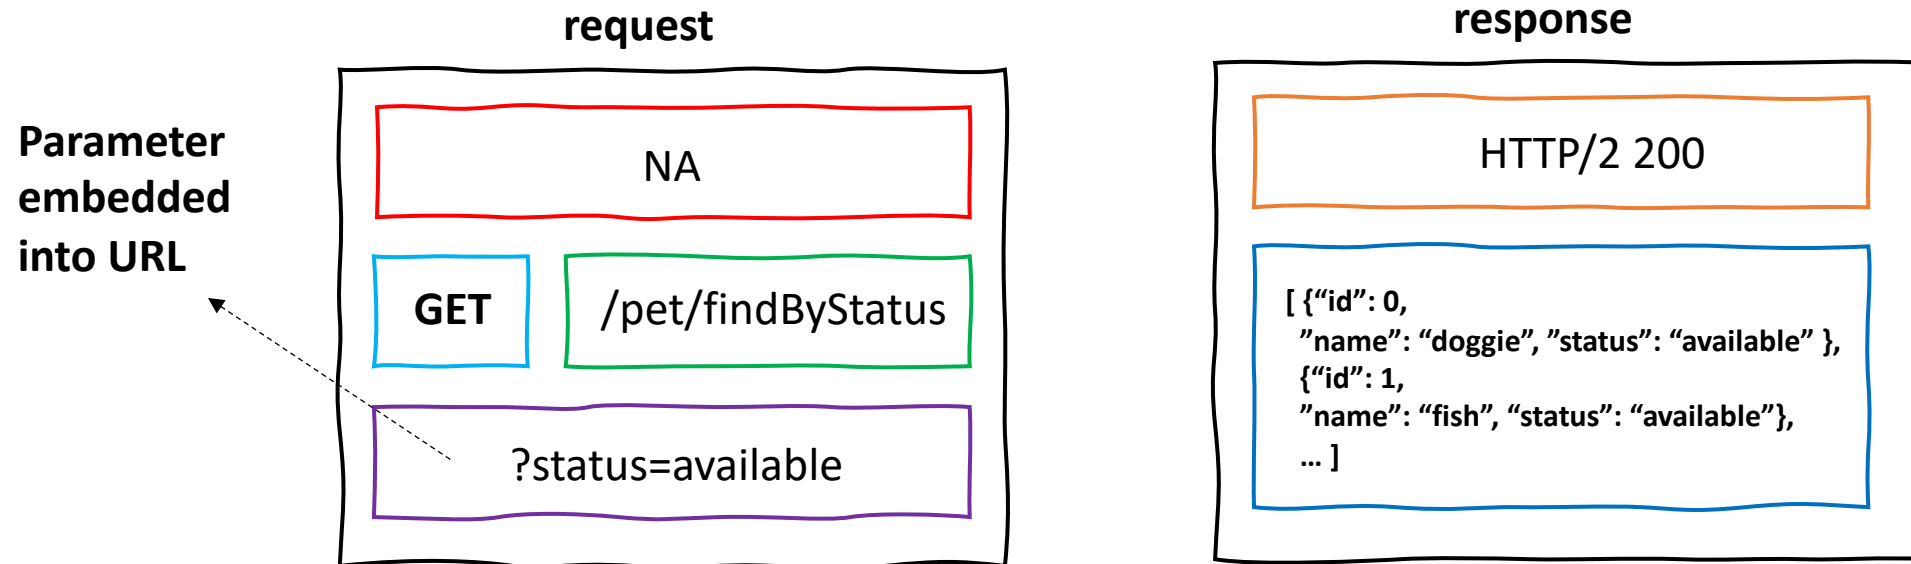
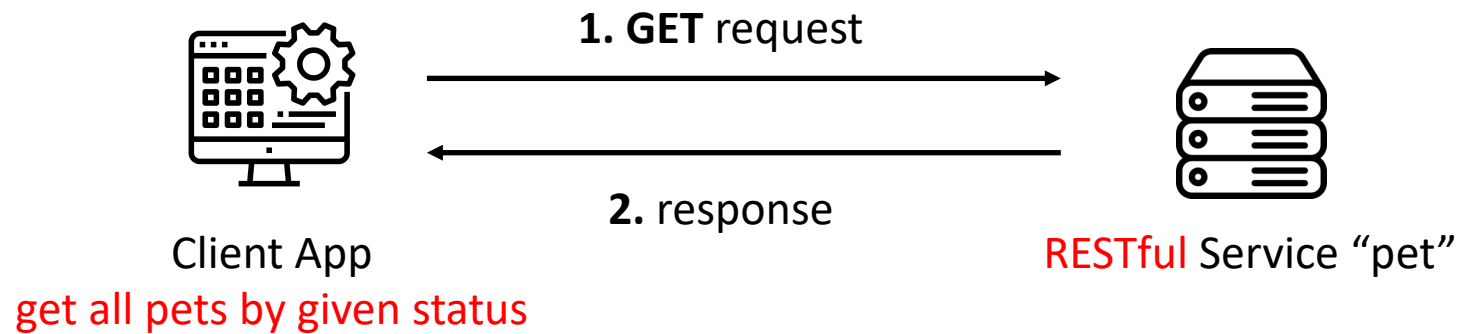
Known endpoint
`https://petstore.swagger.io/v2/pet/findByStatus`



REST APIs: response



REST APIs: request/response scenario (1)



REST APIs: request/response scenario (1)

Request

```
curl -X 'GET' \
  'https://petstore.swagger.io/v2/pet/findByStatus?status=available' \
  -H 'accept: application/json'
```

HTTP method endpoint parameters

Header: body format is JSON

Command line request
(curl utility)

Response

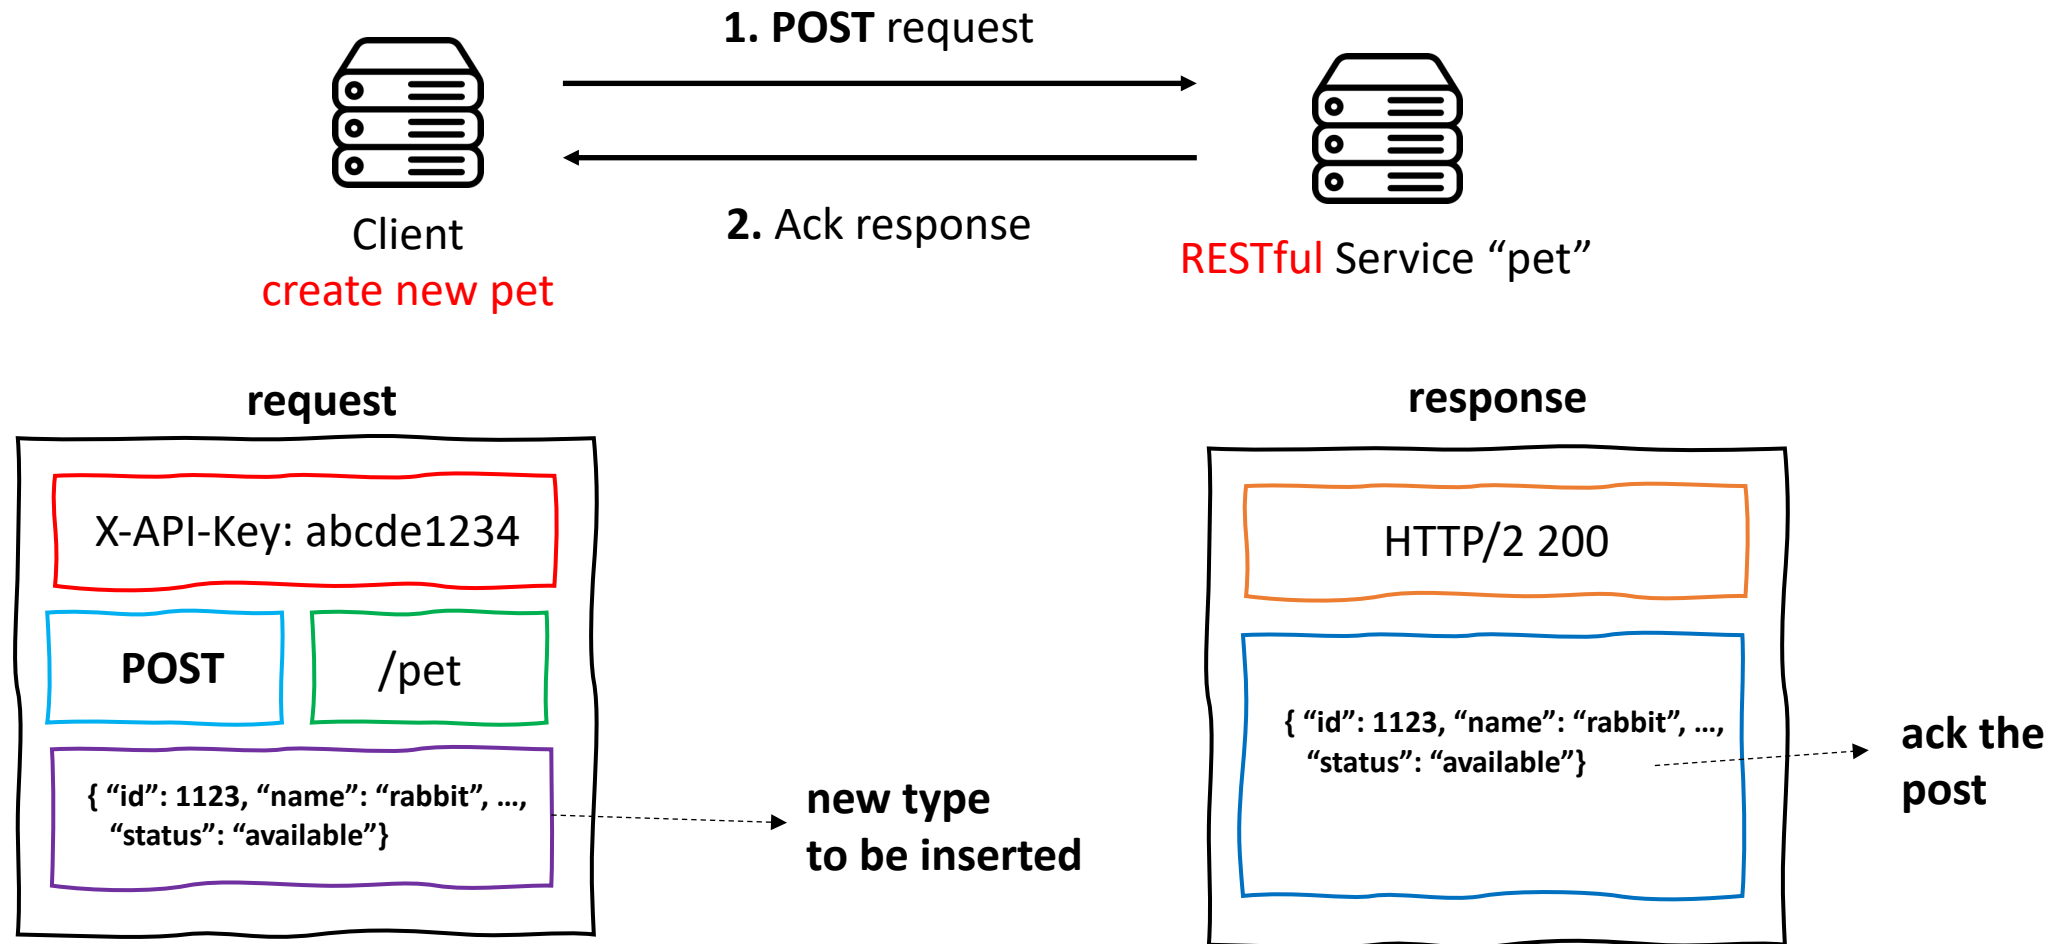
```
HTTP/2 200
date: Fri, 13 Oct 2023 16:38:33 GMT
content-type: application/json ...
```

Status code

```
[ { "id": 9222968140497191000, "category": { "id": 0, "name":
"string" }, "name": "doggie", "photoUrls": [ "string" ], "tags": [ {
"id": 0, "name": "string" } ], "status": "available" }, ... ]
```

Serialized data

REST APIs: request/response scenario (2)



REST APIs: request/response scenario (2)

Request

```
curl -X 'POST' \  
  'https://petstore.swagger.io/v2/pet' \  
  -H 'accept: application/json' \  
  -H 'Content-Type: application/json' \  
  -d { "id": 12300, "category": { "id": 0, "name": "string" },  
    "name": "rabbit", "photoUrls": [ "string" ], "tags": [ { "id": 0,  
    "name": "string" } ], "status": "available" }
```

Response

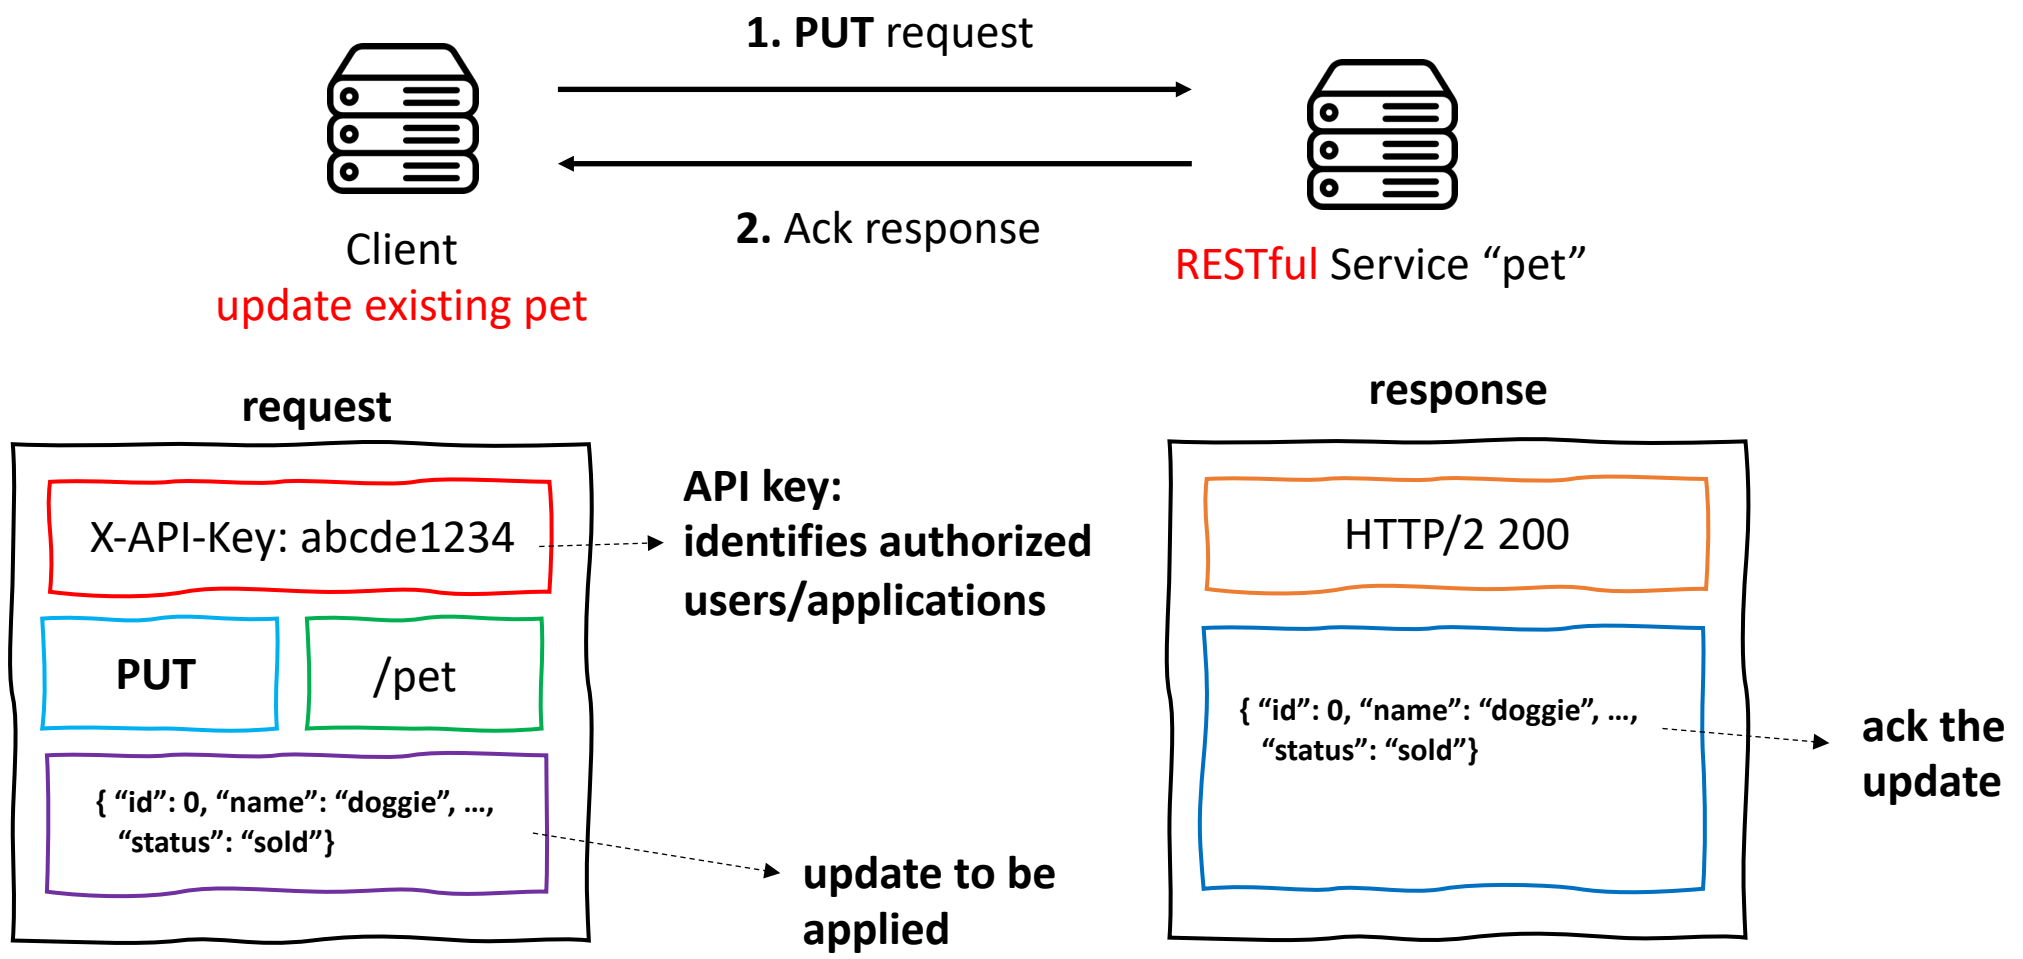
```
HTTP/2 200  
date: Fri, 13 Oct 2023 16:38:33 GMT  
content-type: application/json ...
```

```
{ "id": 12300,  
  "category": { "id": 0, "name": "string" },  
  "name": "rabbit",  
  "photoUrls": [ "string" ],  
  "tags": [ { "id": 0, "name": "string" } ],  
  "status": "available" }
```

EXTRA MATERIAL



REST APIs: request/response scenario (3)



EXTRA MATERIAL



POLITECNICO
MILANO 1863

REST APIs: request/response scenario (3)

Request


```
curl -X 'PUT' \  
  'https://petstore.swagger.io/v2/pet' \  
  -H 'accept: application/json' \  
  -H 'Content-Type: application/json' \  
  -d '{ "id": 0, "category": { "id": 0, "name": "string" }, "name":  
    "doggie", "photoUrls": [ "string" ], "tags": [ { "id": 0, "name":  
    "string" } ], "status": "sold" }'
```

Response

```
HTTP/2 200  
date: Fri, 13 Oct 2023 16:38:33 GMT  
content-type: application/json ...
```

```
{ "id": 0,  
  "category": { "id": 0, "name": "string" },  
  "name": "doggie",  
  "photoUrls": [ "string" ],  
  "tags": [ { "id": 0, "name": "string" } ],  
  "status": "sold" }
```


Representation and structure of exchanged data



- **Representation** impact on
 - Expressiveness
 - Interoperability
 - Performance
 - Transparency

Representation and structure of exchanged data – JSON, XML, Protocol Buffer



POLITECNICO
MILANO 1863

XML

```
<guests>
  <guest>
    <firstName>John</firstName> <lastName>Doe</lastName>
  </guest>
  <guest>
    <firstName>María</firstName> <lastName>García</lastName>
  </guest>
  <guest>
    <firstName>Nikki</firstName> <lastName>Wolf</lastName>
  </guest>
</guests>
```

Message instance

Interface description

JSON

```
{"guests":[
  { "firstName":"John", "lastName":"Doe" },
  { "firstName":"María", "lastName":"García" },
  { "firstName":"Nikki", "lastName":"Wolf" }
]}
```

Message instance

Proto definition

```
message Guest {
  string firstName = 1;
  string lastName = 2;
}
message Guests {
  repeated Guest guests = 3; }
```

Representation and structure of exchanged data – JSON, XML, Protobuffer



POLITECNICO
MILANO 1863

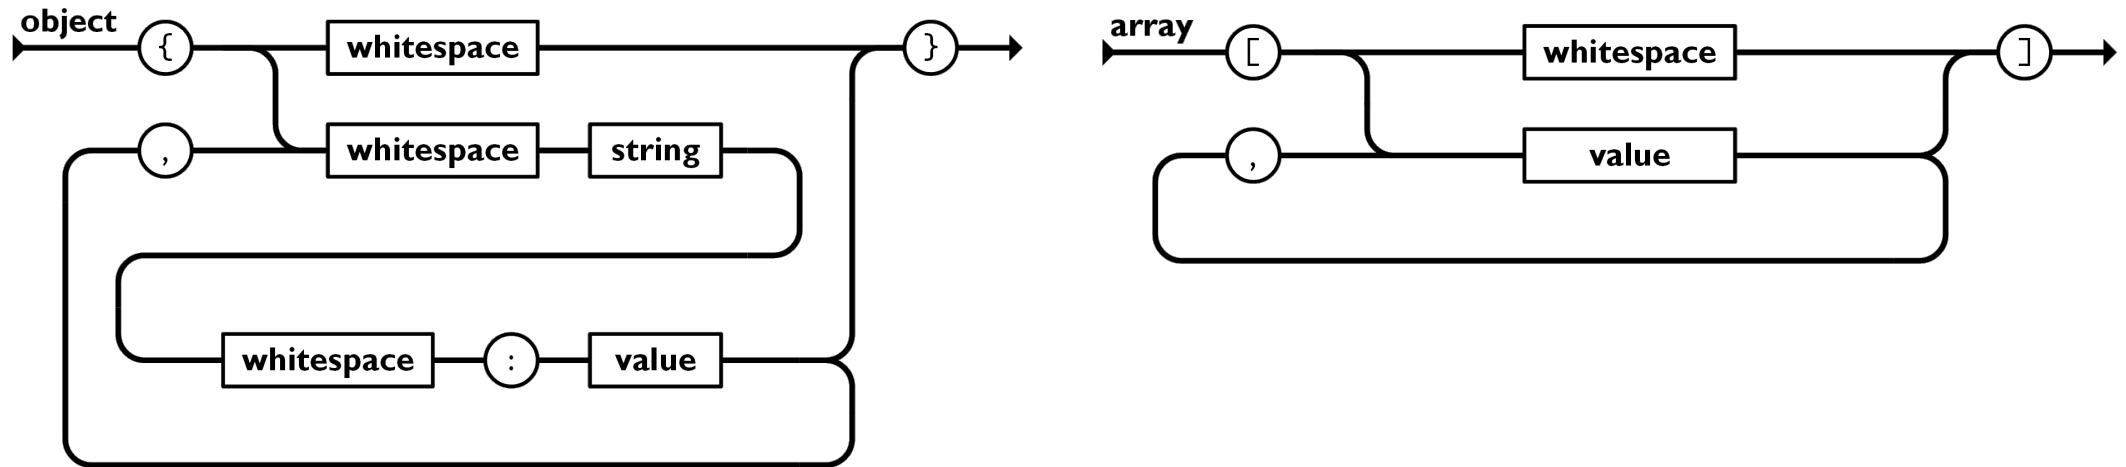
	XML	JSON	Protocol Buffers
Expressiveness	Good	Good	Good
Interoperability	Good	Good	Good
Performance	Verbose, requires multiple parsing passes	More compact than XML, single pass parsing	The most compact one
Transparency	Good, data passed as text	Good, data passed as text	Binary format

EXTRA MATERIAL



JSON standard (json.org)

- Language-independent **text format** to **encode** data
- Two main constructs: **Object** (name-value pairs), **Array** (ordered list)



EXTRA MATERIAL




JSON standard (json.org)

- Possible JSON representation describing pets in our pet store

Array of objects

```
[ { "id": 9222968140497189000,
  "category": { "id": 42, "name": "Chihuahua" },
  "name": "doggie",
  "photoUrls": [ "http://..." ],
  "tags": [ { "id": 0, "name": "small" } ],
  "status": "available"
},
{ "id": 9222968140497189111,
  "category": { ... },
  "name": "fish",
  "photoUrls": [ "http://..." ],
  "tags": [ { "id": 0, "name": "small" } ],
  "status": "sold"
}, ...
]
```

Key-value
pairs



Error handling

- **Examples of issues**

- An operation is called with invalid parameters
- The call does not return anything
 - The component cannot handle the request in the current state
 - Hardware/software errors prevent successful execution
 - Misconfiguration issues (e.g., the server is not correctly connected to the database)

- **Possible reactions**

- Raising an exception
- Returning an error code
- Log the problem



Multiple interfaces and separation of concerns

- A server can offer **multiple interfaces** at the same time
- This enables
 - Separation of concerns
 - Different levels of access rights
 - Support to **interface evolution** (see next)

Interface evolution

- Interfaces constitute the contract between servers and clients
- Sometimes interfaces need to evolve (e.g., to support new requirements)
- **Strategies** to support continuity
 - **Deprecation**: declare well in advance that an interface version will be eliminated by a certain date
 - **Versioning**: maintain multiple active versions of the interface
 - **Extension**: a new version extends the previous one




Software Design

Interface Documentation

Documenting interfaces

- Interface **documentation**
 - explain how to use it
 - should not include information about the internals of a component (otherwise users may assume behavior that could change in the future)
- **Audience** of an interface documentation
 - **Developers** and maintainers of the component **offering** the interface
 - **Developer** and maintainers of the component **using** the interface
 - **QA teams** for system **integration and testing**
 - **Software architects** (including those looking for reusable components)

OpenAPI Specification



- **Scenario**
 - **New_dev** should get started developing and maintaining the application using the “pet” service
 - How can New_dev **understand the API** and become productive **without digging** into the sources of “pet”?
 - New_dev can refer to the so-called **OpenAPI definition**

OpenAPI Specification

- **OpenAPI Specification**
 - **Defines** how to describe a REST API interface through an **OpenAPI definition**
- **OpenAPI definition**
 - File (JSON or YAML)
 - Describes what a service can do using its interface
- **Benefits**
 - Standardized format, public, and well-known
 - Readable by
 - **humans** to **understand** and **use** the REST API
 - **machines** to **automate** tasks like testing or code generation



OpenAPI Specification

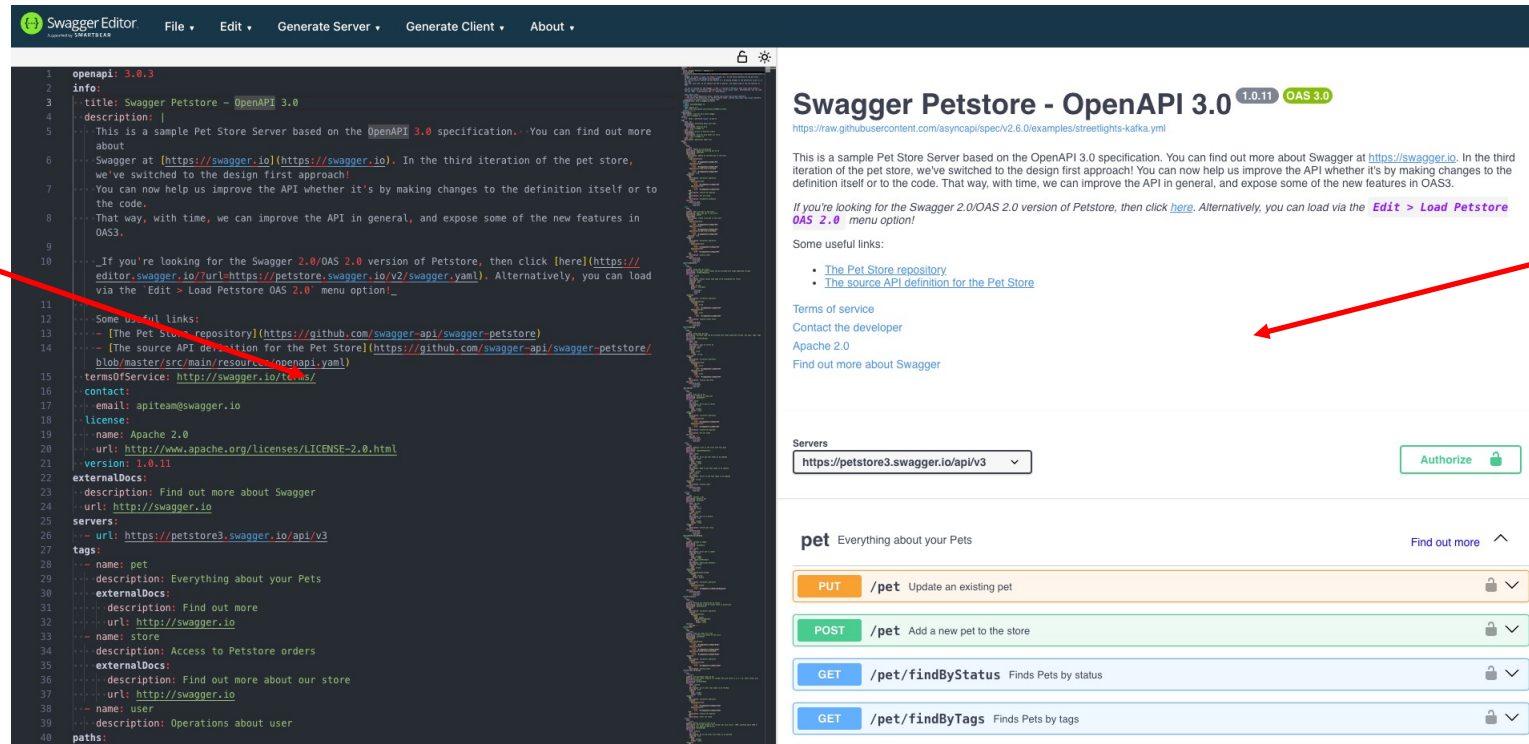
- **OpenAPI definition describes**
 - Endpoints
 - Resources
 - Operations
 - Parameters including data types
 - Authentication/authorization mechanism
- **Tool support**
 - **API validator** — conformance to the standard
 - **Documentation** generator — clear human readable description
 - **SDK** generator — automated creation of client libraries in a programming language of choice

OpenAPI Spec — petstore.yaml

Try it out at:

https://editor.swagger.io/?_ga=2.155528316.612942566.1697644409-1947143061.1697644409

OpenAPI
definition
(YAML file)



The screenshot displays the Swagger Editor interface. On the left, the OpenAPI 3.0 specification (YAML file) is shown in a code editor. The specification includes details about the Swagger Petstore, its version (1.0.11), and the API endpoints. On the right, the generated documentation is displayed, featuring the title 'Swagger Petstore - OpenAPI 3.0' and a list of API endpoints with their methods and descriptions. A red arrow points from the 'OpenAPI definition (YAML file)' text to the left pane, and another red arrow points from the 'Generated documentation' text to the right pane.

Generated
documentation

EXTRA MATERIAL



POLITECNICO
MILANO 1863

OpenAPI Spec — petstore.yaml

OpenAPI version

Basic REST API info

version
of the API
itself

```
1 openapi: 3.0.3
2 info:
3   title: Swagger Petstore - OpenAPI 3.0
4   description: |
5     This is a sample Pet Store Server based on the OpenAPI 3.0 specification. You can find out more about Swagger at
6     [https://swagger.io](https://swagger.io).
7   Some useful links:
8   - [The Pet Store repository](https://github.com/swagger-api/swagger-petstore)
9   - [The source API definition for the Pet Store](https://github.com/swagger-api/swagger-petstore/blob/master/src/main/resources/openapi.yaml)
10  version: 1.0.11
11 externalDocs:
12   description: Find out more about Swagger
13   url: http://swagger.io
14 servers:
15   - url: https://petstore3.swagger.io/api/v3
```

Textual description

Base URL of the endpoints

EXTRA MATERIAL



OpenAPI Spec — petstore.yaml

Description of API requests (“paths” section)

Resource
HTTP method

```
29 paths:
  ...
  105 .. /pet/findByStatus:
  106 ..   get:
  107 ..     tags:
  108 ..       - pet
  109 ..     summary: Finds Pets by status
  110 ..     description: Multiple status values can be provided with comma separated strings
  111 ..     operationId: findPetsByStatus
  112 ..     parameters:
  113 ..       - name: status
  114 ..         in: query
  115 ..         description: Status values that need to be considered for filter
  116 ..         required: false
  117 ..         schema:
  118 ..           type: string
  119 ..           default: available
  120 ..           enum:
  121 ..             - available
  122 ..             - pending
  123 ..             - sold
```

Textual
description

Can be also header/path

Parameters

Enumerative type

EXTRA MATERIAL



OpenAPI Spec — petstore.yaml

Possible responses

```
105 /pet/findByStatus:
...
124 responses:
125   '200':
126     description: successful operation
127     content:
128       application/json:
129         schema:
130           type: array
131           items:
132             $ref: '#/components/schemas/Pet'
133       application/xml:
134         schema:
135           type: array
136           items:
137             $ref: '#/components/schemas/Pet'
138   '400':
139     description: Invalid status value
140   security:
141     - api_key: []
```

Ok response

Either JSON
or XML format
depending on request

Each item of the array
has the referenced schema

mandatory

Error in case of bad request

Authentication with API key

Pet {

id

> [...]

name*

> [...]

category

> [...]

photoUrls*

> [...]

tags

> [...]

status

> [...]

}

array

object

Category > {...}

Generated representation
of the pet schema

EXTRA MATERIAL



OpenAPI Specification

- Automated generation of client code example
 - Codegen tool: <https://swagger.io/tools/swagger-codegen/>

```
swagger-codegen generate -i https://petstore.swagger.io/v2/swagger.json \  
-l python -o /petstore
```

```
55 configuration = swagger_client.Configuration()  
56 # create an instance of the API class  
57 api_instance = swagger_client.PetApi(swagger_client.ApiClient(configuration))  
58 status = ['pending'] # list[str] | Status values that need to be considered for filter  
59  
60 try:  
61     # Finds Pets by status  
62     api_response = api_instance.find_pets_by_status(status)  
63     pprint(api_response)  
64 except ApiException as e:  
65     print("Exception when calling PetApi->find_pets_by_status: %s\n" % e)
```




Software Design

Clinet-server style: handling multiple requests

Handling multiple requests



- The server must be able to receive and handle requests from multiple clients, example:



- How can the server serve the second `withdraw` if the first one is still ongoing?

Handling multiple requests: a first approach (forking)

- Simple approach used by **Apache Web Server**
 - **One process per request** or per client



Handling multiple requests: a first approach (forking)

- **Strengths**

- Architectural simplicity
- Isolation and protection given by the “one-connection-per-process” model
 - Slow processes do not affect other incoming connections
- Simple to program
- Effective solution till ~2000

- **Issues**

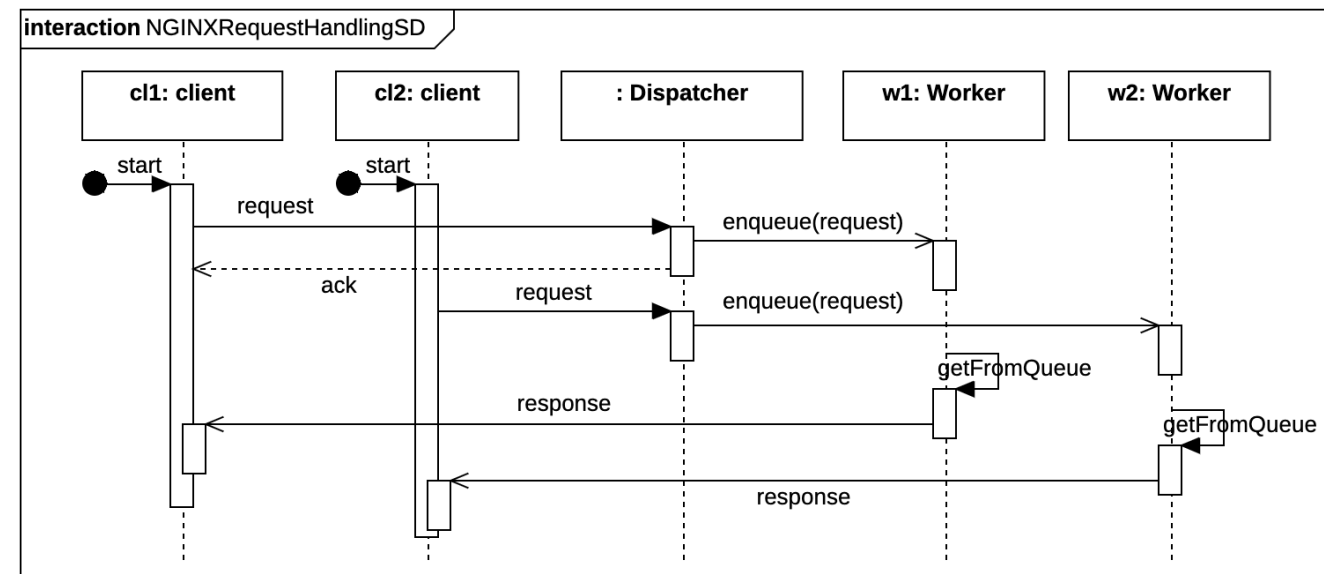
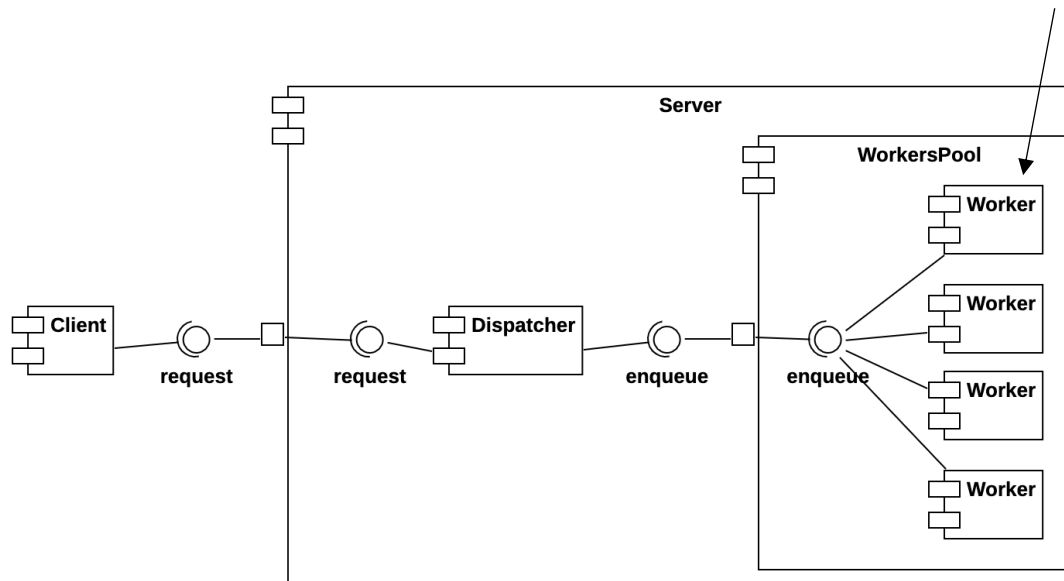
- Growth of the WWW over the past 20 years (#users and weight of web pages)
- #active processes at time t difficult to predict and may saturate the resources
- Expensive fork-kill operations for each incoming connection
- → **Scalability issues!**

Handling multiple requests: alternative approach (worker pooling)

- Alternative approach adopted by **NGINX Web Server**
- Designed for high concurrency — deals with scalability issues



Predefined range of instances (n..m)



Handling multiple requests: alternative approach (worker pooling)

- **Note (1):** NGINX tackles the previous issues by introducing a new **architectural tactic**
 - **Tactic** = Design decisions that influence the control of one or more quality attributes
- **Strengths of the new approach**
 - Number of workers is fixed → they do not saturate available resources
 - Each worker has a queue
 - When queues are full the dispatcher drops the incoming requests to keep high performance
 - Dispatcher balances the workload among available workers according to specific policies

Handling multiple requests: alternative approach (worker pooling)

- **Note (2):** architectural tactics introduce **quality attribute trade-offs**
- NGINX decided to **optimize scalability** and **performance** by **sacrificing availability (in some cases)**
 - When all worker queues are full, the dispatcher drops incoming requests rather than spawning new workers



POLITECNICO
MILANO 1863

Software Design

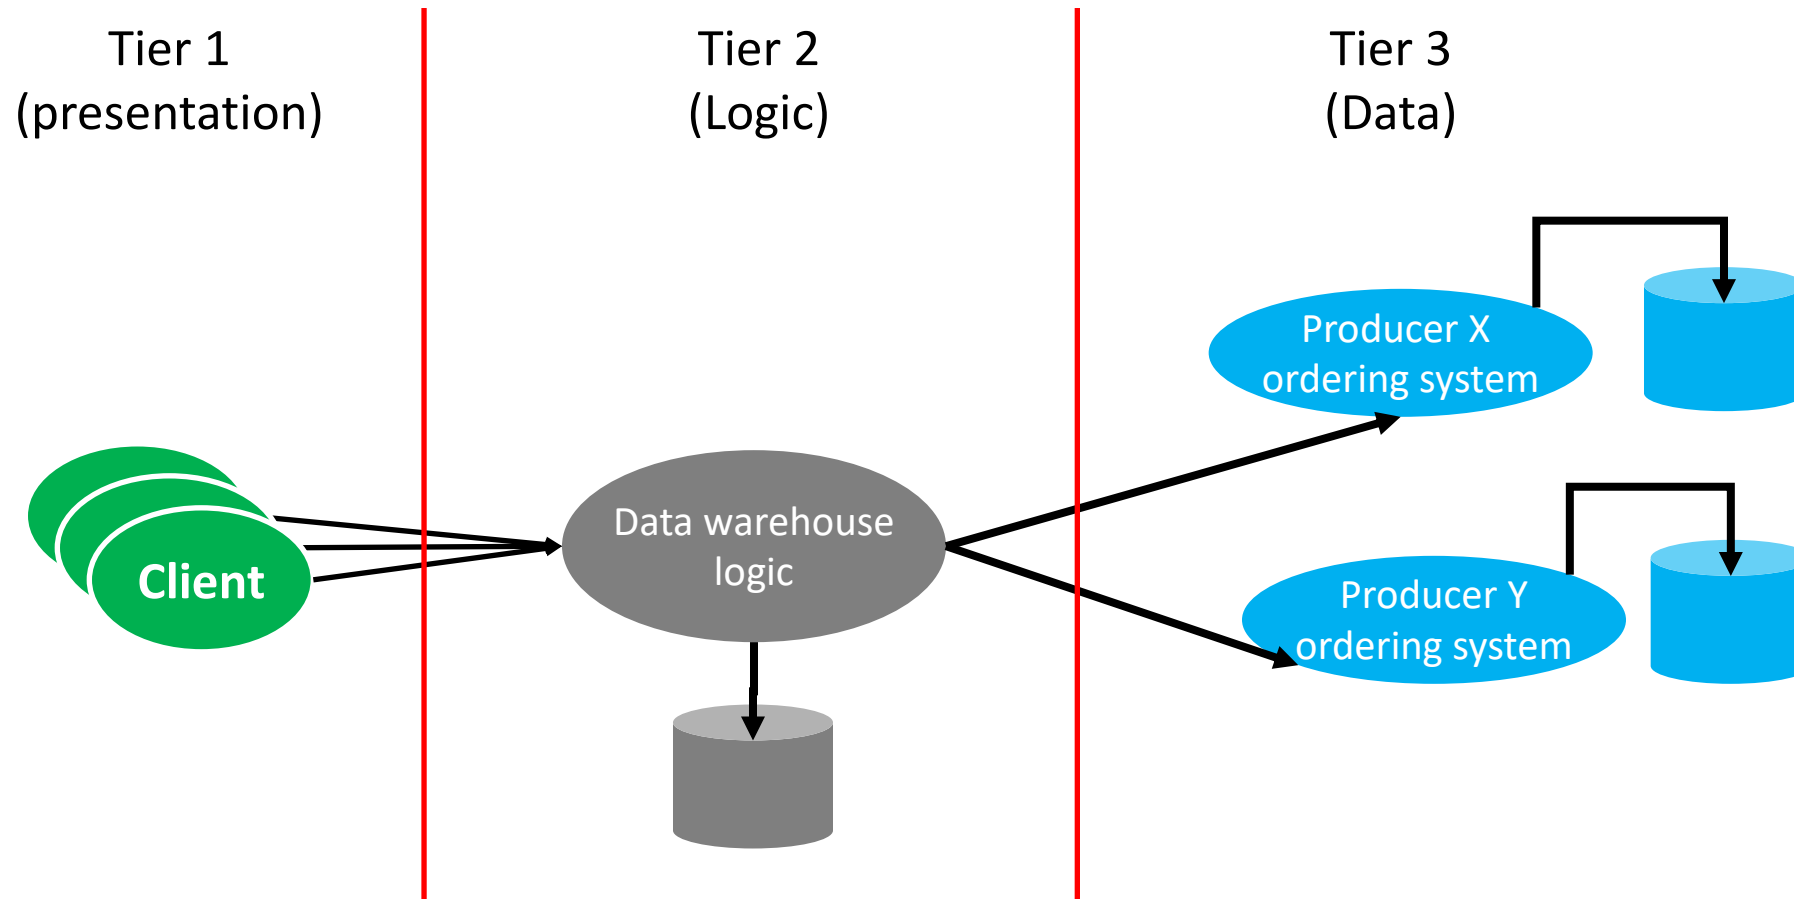
Other styles: Multi-tier, Event-driven

Three-tier Architecture – an example



POLITECNICO
MILANO 1863

Data warehouse for a supermarket




From 3 to N tiers

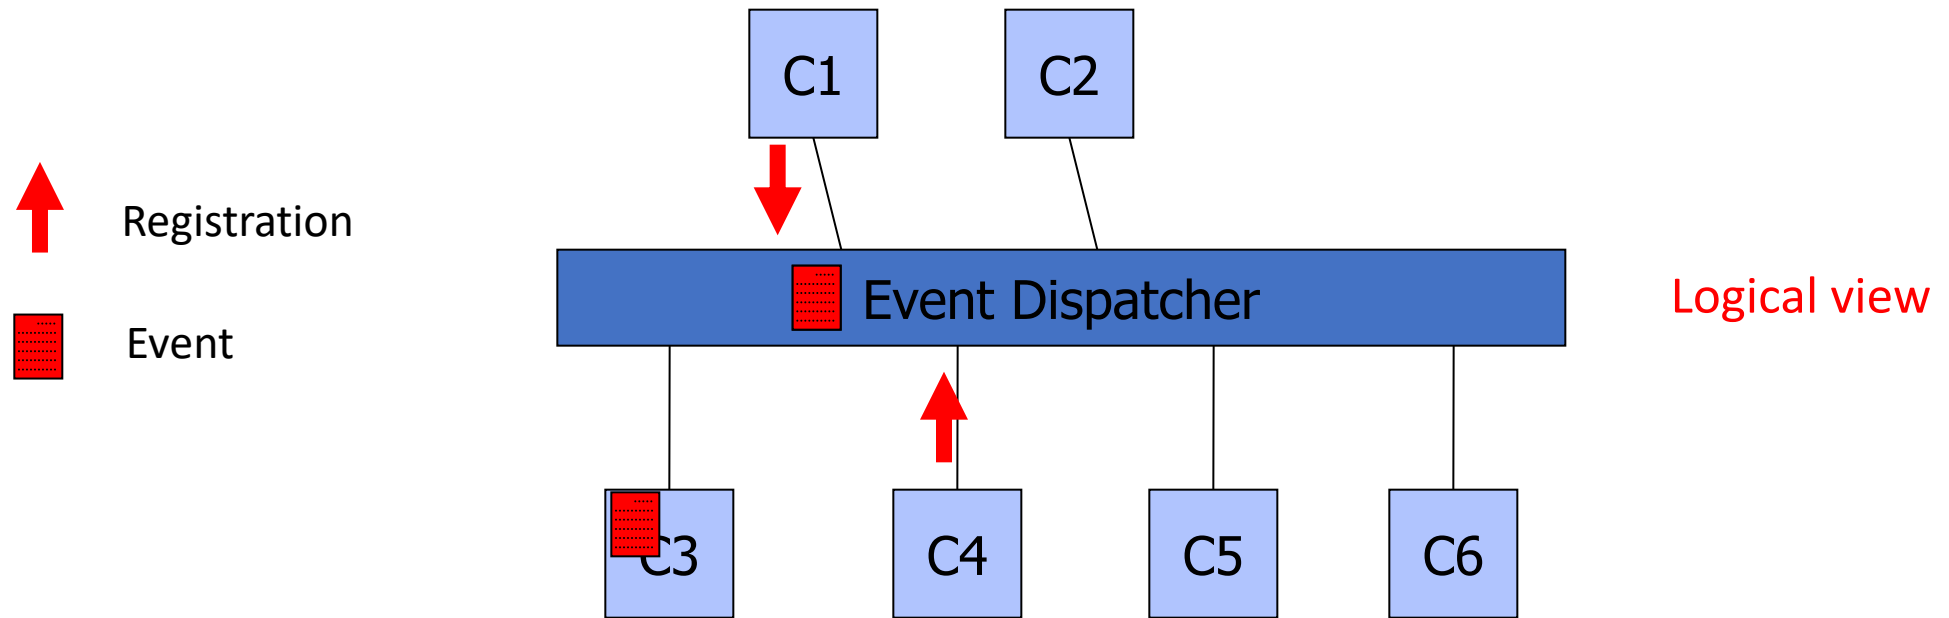


POLITECNICO
MILANO 1863

Data warehouse for a supermarket



Event-driven Architecture



- Components can **register** to / **send events**
- Events are **sent to all registered components**
- **No explicit naming** of target component



Event-driven paradigm

- Often called **publish-subscribe**
 - **Publish** = event generation
 - **Subscribe** = declaration of interest
- Different kinds of event languages possible

Event-driven systems Pros & Cons

- + **Very common in modern development practices**
 - + E.g., continuous integration / deployment (such as GitHub Actions)
- + **Easy addition/deletion of components**
 - + Publishers/subscribers are decoupled
 - + The event dispatcher handles this dynamic set
- **Potential scalability problems**
 - The event dispatcher may become a bottleneck (under high workload)
- **Ordering of events**
 - Not guaranteed, not straightforward

Other characteristics

- Messages/events are **asynchronous** (send and forget)
- Computation is **reactive** (driven by receipt of message)
- **Destination** of messages **determined by receiver**, not sender (location/identity abstraction)
- **Loose coupling** (senders/receivers added without reconfiguration)
- **Flexible** communication means (one-to-many, many-to-one, many-to-many)



Examples of relevant technologies

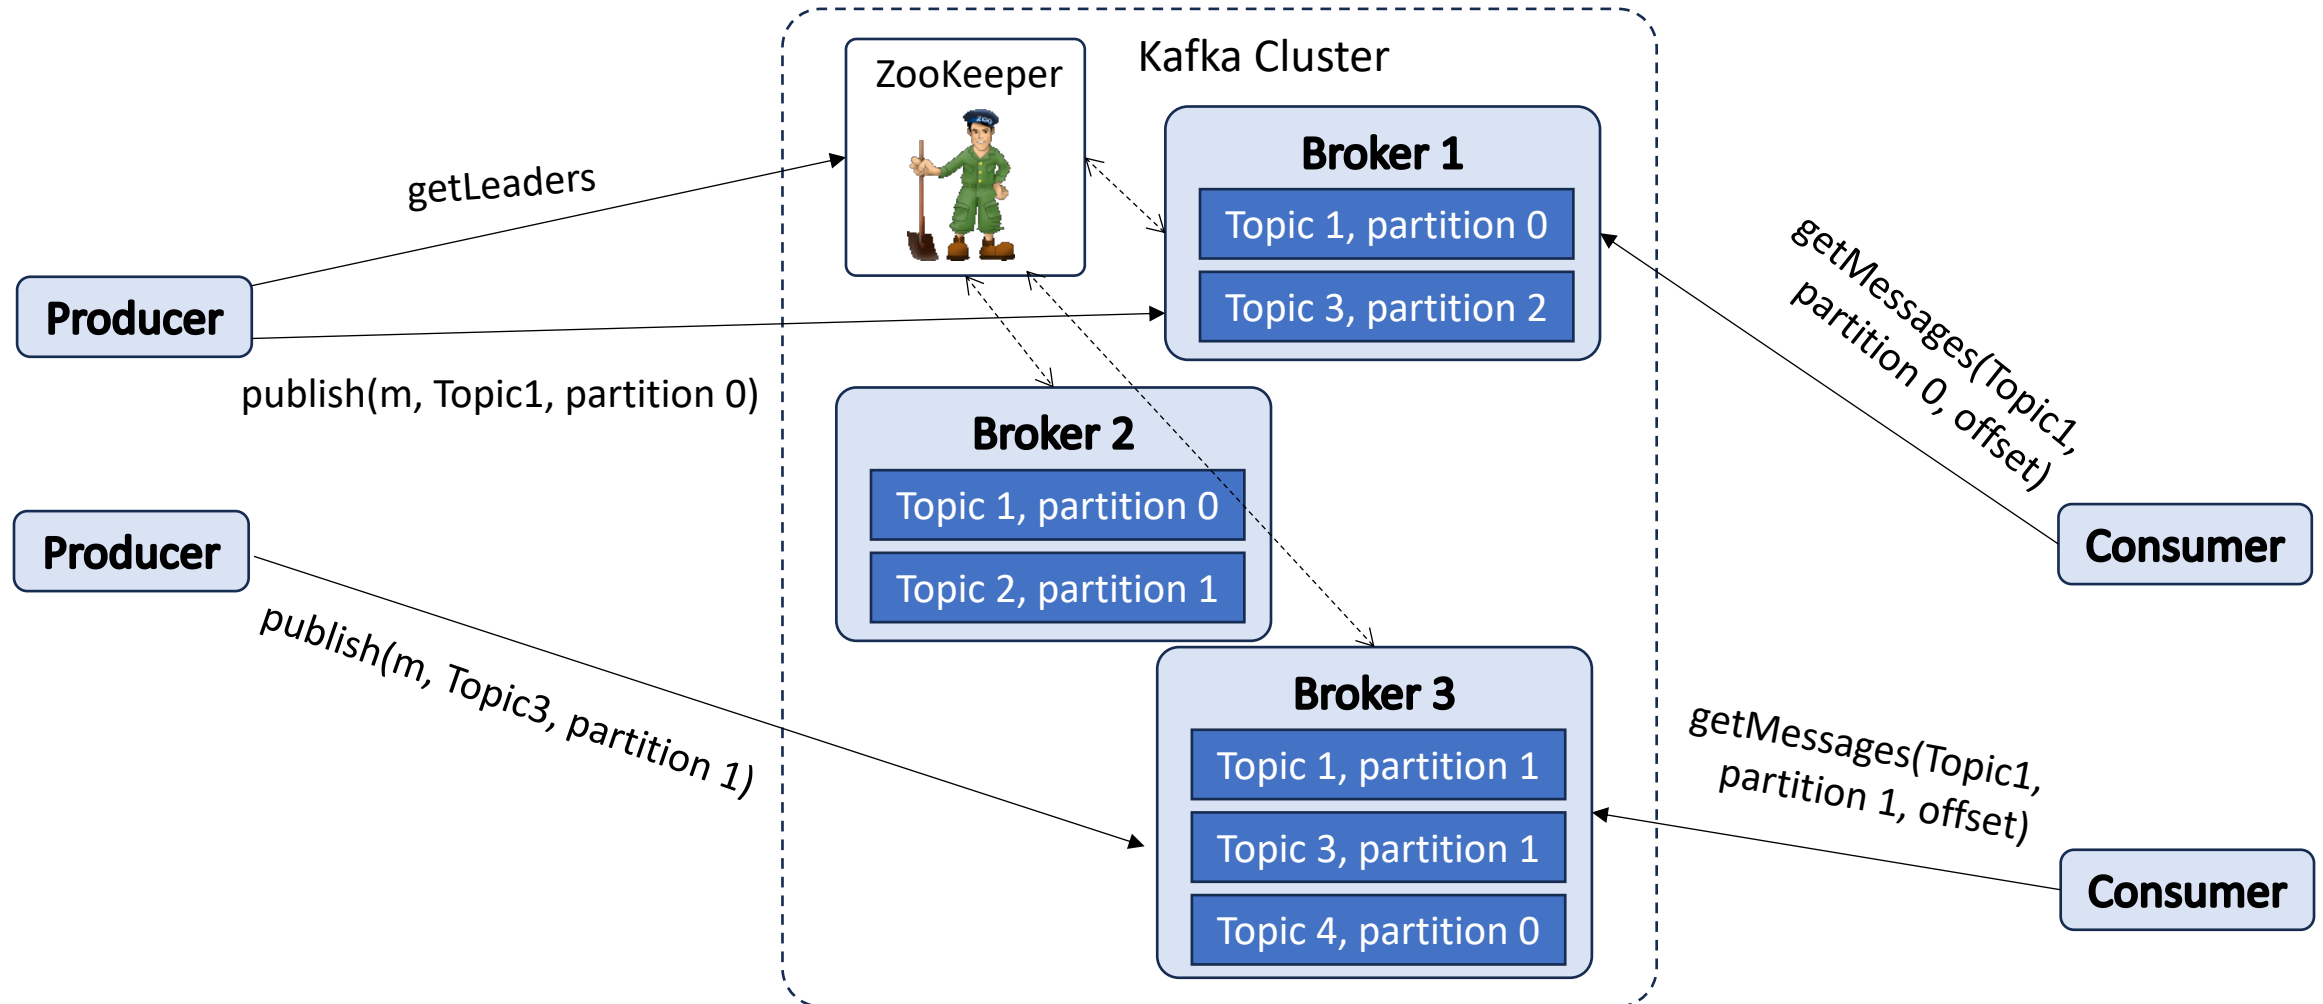
- Apache Kafka
 - <https://kafka.apache.org/>
- RabbitMQ
 - <https://www.rabbitmq.com/>

Apache Kafka

- Kafka is a **framework** for the event-driven paradigm:
 - Includes primitives to create **event producers** and **consumers** and a runtime infrastructure to handle **event transfer** from producers to consumers
 - **Stores events** durably and reliably
 - Allows consumers to **process events** as they occur or retrospectively
- These services are offered in a distributed, highly scalable, elastic, fault-tolerant, and secure manner



Kafka architecture

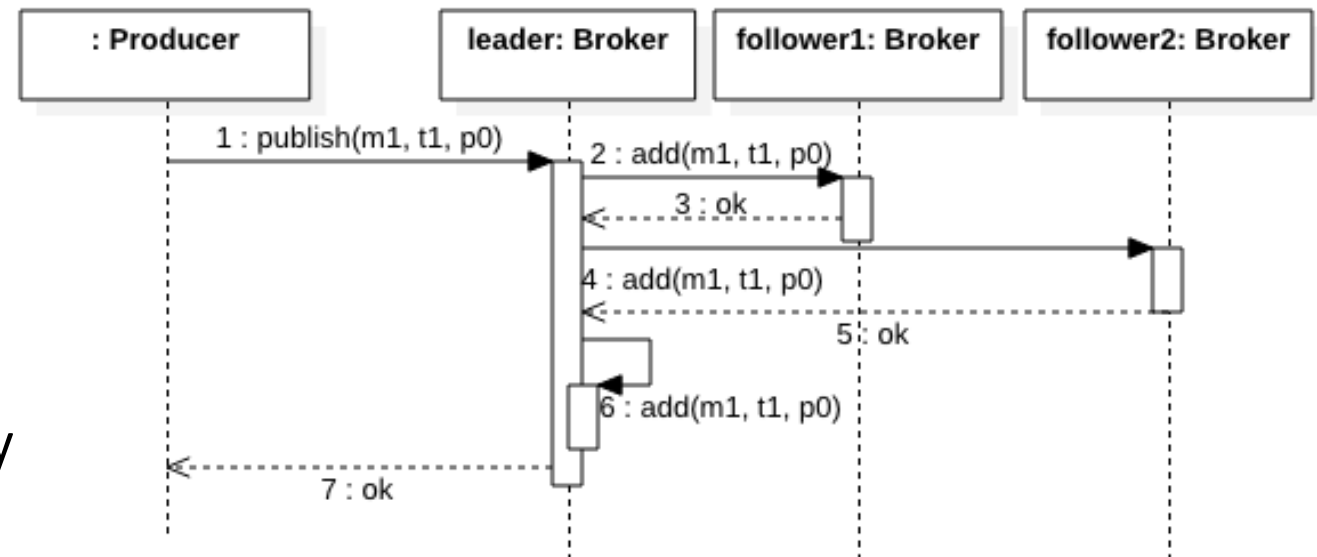


Kafka main characteristics

- Each broker handles a set of **topics** and **topic partitions**, parts including sets of messages on the topic
- Partitions are independent from each other and can be **replicated** on multiple brokers for fault tolerance
- There is one **leading broker** per partition. The other brokers containing the same partition are **followers**
- **Producers** know the available leading brokers and send messages to them
- Messages in the same topic are organized in **batches** at the producers' side and then sent to the broker when the batch size overcomes a certain threshold
- Consumers adopt a **pull approach**. They receive in a single batch all messages belonging to a certain partition starting from a specified **offset**
- Messages remain available at the brokers' side for a specified period and can be **read multiple times** in this period
- The leader keeps track of the **in-synch followers**
- **ZooKeeper** is used to oversee the correct operation of the cluster. All brokers send heartbeat to ZooKeeper. ZooKeeper will replace a failing broker by electing a new leader for all the partitions the failing broker was leading. It may also start/restart brokers


Message delivery semantics — Producer

- Brokers commit messages by storing them in the corresponding partition
- Leader adds the message to followers (replicas) if available
- **Issue:** in case of failure, the producer may not get the response (msg #7)
 - The producer has to resend the message
 - Kafka brokers can identify and eliminate duplicates
- Synchronization with replicas can be transactional
 - **Exactly-once** semantics is possible but long waiting time
 - **At-least-once** can be chosen by excluding duplicates' management
 - **At-most-once** can be chosen by publishing messages asynchronously



Message delivery semantics — Consumer

- Each **consumer** can rely on a **persistent log** to keep track of the **offset** so that it is not lost in case of failure
- If the consumer fails after having elaborated messages and before storing the new offset in the log, the same messages will be retrieved again
 - → **at-least-once semantics**
- The delivery semantics can be changed if the new offset is stored before the elaboration
 - → **at-most-once semantics** because, if failing after storing the offset, the effect of the received messages does not materialize
- Transactional management of log allows for **exactly-once semantics**



Kafka architectural tactics

- **Scalability**

- **Multiple partitions and multiple brokers** → possibility to distribute producers/consumers on different partitions handled by different brokers
- **Scale of operation:** Kafka supports the creation of clusters of brokers
 - Each clusters includes up to 1K brokers able to handle trillions of messages per day

- **Fault tolerance**

- Partitions are **stored persistently**
- **Replication** significantly reduces the chances of losing data
- Cluster management takes care of restarting brokers and setting leaders as needed