

x86 Crash Course

With a focus on Linux and a glance to x86_64

Alessandro Bertani, alessandro.bertani@polimi.it

April 3, 2024

Based on material from Mario Polino, Armando Bellante and Lorenzo Binosi

The x86 Architecture

Instruction Set Architecture (ISA)

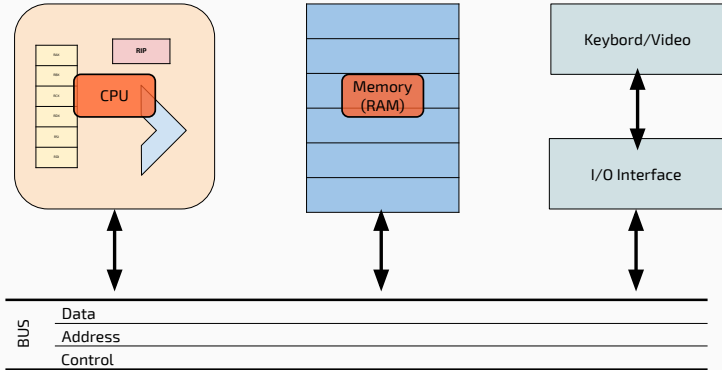
- “Logical” specification of a computer architecture
- Concerned with programming concepts
 - instructions, registers, interrupts, memory architecture, ...
- May differ (widely) from the actual microarchitecture
- Examples:
 - x86 (IA-32 and x86_64)
 - ARM (mobile devices)
 - MIPS (embedded devices, e.g., consumer routers)
 - AVR, SPARC, Power, RISC V, ...

The x86 ISA

- Born in 1978, 16-bit ISA (Intel 8086)
- Evolved to a 32-bit ISA (1985, Intel 80386)
- Evolved to a 64-bit ISA (2003, AMD Opteron)
- CISC design (e.g., string operations)
- Many legacy features (e.g, segmentation)
- We'll see the basics of the “core” ISA
 - There is also the floating point unit, processor-specific features, and extensions such as SIMD (MMX, SSE, SSE2) with their own instructions and registers¹

¹Complete reference: Intel Software Developer's Manual, about 5,000 pages (<https://software.intel.com/en-us/articles/intel-sdm>)

Von Neumann Architecture

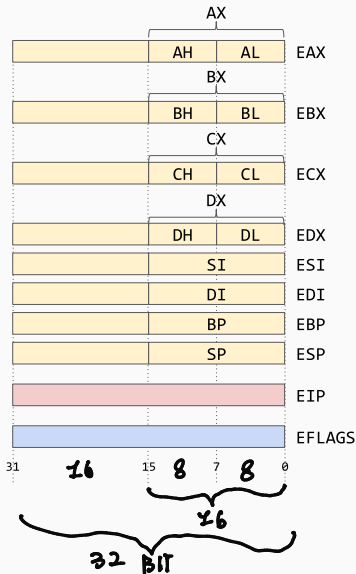


Memory

EACH CELL HAS A UNIQUE ADDRESS
IDENTIFIED AS AN HEXADECIMAL NUMBER

	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
0x8040204...	00	01	02	03	04	05	06	07	08	09	0A	0B	0C	0D	0E	0F
0x8040203...	E8	FF	FF	FF	E0	FF	FF	FF	D8	FF	FF	FF	D0	FF	FF	FF
0x8040202...	50	20	40	80	60	20	40	80	70	20	40	80	80	20	40	80
0x8040201...	48	65	6C	6C	6F	20	77	6F	72	6C	64	00	00	00	00	00
0x8040200...	'H'	'e'	'l'	'l'	'o'	' '	'w'	'o'	'r'	'l'	'd'	00	00	00	00	00

IA-32: Registers



HYPOTETICALLY USED FOR ANY OPERATION

- General-purpose registers

- EAX, EBX, ECX, EDX
- ESI, EDI (source and destination index for string operations)
- EBP (base pointer)
- ESP (stack pointer)

- Instruction pointer: EIP

"SPECIAL PURPOSE"

- No explicit access
- Modified by jmp, call, ret
- Read through the stack (saved IP)

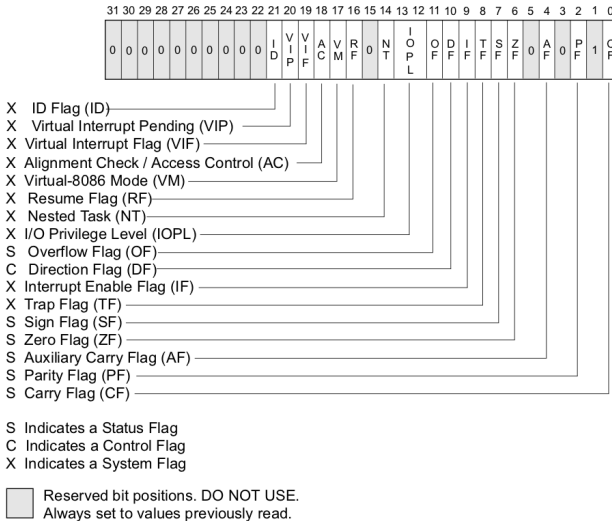
- Program status and control: EFLAGS

- (segment registers)

IA-32: EFLAGS register

- 32-bits register, boolean flags
- **Program status:** overflow, sign, zero, auxiliary carry (BCD), parity, carry
 - Indicate the result of arithmetic instructions
 - Extremely important for control flow
- **Program control:** direction flag
 - controls string instructions (auto-increment or auto-decrement)
- **System:** control operating-system operations

IA-32: EFLAGS register



4 MAIN DIMENSIONS USED TO ADDRESS MEMORY

byte 8 bits

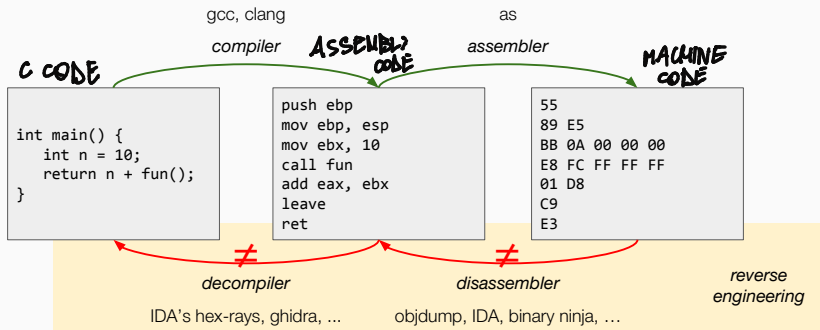
word 2 bytes

dword Doubleword, 4 bytes (32 bits)

qword Quadword, 8 bytes (64 bits)

Assembly and Machine Code

Assembly language: specific to each ISA, mapped to binary code



For simplicity, we don't deal with the *linking* process.

Two main syntaxes:

- **Intel**: default in most Windows programs (e.g., IDA)
- **AT&T**: default in most UNIX tools (e.g., gdb, objdump)

Beware: The order of the operands is **different**

We will use the Intel syntax

Assembly: Syntax

move the value 0 to EAX

Intel

`mov eax, 0h`

"0 AS **HEXADECIMAL**"

AT&T

`movl $0x0,%eax`

move the value 0 to the address contained in EBX+4

Intel

`mov [ebx+4h],0h`

AT&T

`movl $0x0,0x4(%ebx)`

Examples

Immediate to register:

`mov eax, 4h` $EAX = 4$

Register to register:

`mov eax, ebx` $EAX = EBX$

Memory to register (and register to memory):

`mov eax, [ebx]` $EAX = *EBX$

`mov eax, [ebx + 4h]` $EAX = *(EBX + 4)$

`mov eax, [edx + ebx*4 + 8]` $EAX = *(EDX + EBX * 4 + 8)$

Note: memory to memory is an **invalid** combination²

²Except in some instructions, such as `movs` (move from string to string).

x86 Assembly and Machine Code

Instruction = opcode + operand

Example

<code>mov ebp, esp</code>	<code>mov r/m r</code>	<code>ebp, esp</code>
	89	E5 10
<code>mov bl, 10</code> ↖ <i>LOWEST 8 BITS</i>	<code>mov bl</code>	
	B3	0A
<code>mov bx, 10</code> ↖ <i>LOWEST 16 BITS</i>	<i>size prefix</i>	0A 00
	66 BB	
<code>mov ebx, 10</code>	<i>mov ebx</i>	0A 00 00 00
	BB	
	<i>prefix + opcode (1-3 bytes)</i>	<i>operands</i>

Beware: in x86, instructions have **variable length**.

Basic instructions

- **Data Transfer:** `mov`, `push`, `pop`, `xchg`, `lea`
- **Integer Arithmetic:** `add`, `sub`, `mul`, `imul`, `div`, `idiv`, `inc`, `dec`
- **Logical:** `and`, `or`, `not`, `xor`
- **Control Transfer:** `jmp`, `jne`, `call`, `ret`
- and lots more...

Data Transfer: `mov`

- `mov` destination, source
source: immediate, register, memory location
destination: register or memory location
- Basic load/store operations
 - Register to register, register to memory, immediate to register, immediate to memory
 - Memory to memory is INVALID (in every instruction)

Examples

`MOV eax, ebx`

`MOV eax, FFFFFFFFh`

`MOV ax, bx`

`MOV [eax],ecx`

`MOV [eax],[ecx] NO!!!`

`MOV al, FFh`

Integer Arithmetics: add and sub

add <u>destination</u> , <u>source</u>	sub <u>destination</u> , <u>source</u>
$\text{dest} \leftarrow \text{dest} + \text{source}$	$\text{dest} \leftarrow \text{dest} - \text{source}$

- Addressing:
source: immediate, register, memory location
destination: register or memory location
(the destination has to be at least as large as the source)
- Negate a value: `neg [op]`
- Bitwise operations: `and`, `or`, `xor`, `not` work similarly

Examples

add esp, 44h	add edx, cx	add al, dh
sub esp, 33h	sub eax, ebx	sub [eax], 1h

Integer Arithmetics: unsigned multiply (`mul`)

- `mul source` **SINGLE OPERAND**
source: register or memory location
 - `dest ← implied_op × source`
 - **Implied operands** according to the size of **source**
 - **First operand:** AL, AX, or EAX
 - **Destination:** AX, DX:AX, EDX:EAX (double the size of **source**)
 - Signed multiply: `imul`
- EXAMPLE: 8 bit · 8 bit**
→ 16 bits
↓
WE USE 2 DIFFERENT REGISTERS TO STORE THE RESULT

Example

- `mul ebx: EDX:EAX ← EAX * EBX`
 - most significant bits of the result in EDX
 - least significant bits of the result in EAX
- `mul cx: DX:AX ← AX * CX`
- `mul cl: AX ← AL * CL`

Integer Arithmetics: unsigned divide (div)

- `div source`
source: register or a memory location
- Computes quotient and remainder
- Implied operand: `EDX:EAX` (according to the size of **source**)
- Signed divide: `idiv`

LOWEST
HIGHEST
32 bits

Examples

- `div ebx` (4 bytes)
 - $EAX \leftarrow EDX:EAX / EBX$
 - $EDX \leftarrow EDX:EAX \% EBX$
- `div bx` (2 bytes)
 - $AX \leftarrow DX:AX / BX$ $DX = DX:AX \% BX$
- `div bl` (1 byte)
 - $AL \leftarrow AX / BX$ $AH = AX \% BX$

Integer Arithmetics: `cmp` and `test`

`cmp` `op1`, `op2`

Computes `op1 - op2`

RESULT: WHICH ONE IS
HIGHER (OR =)

`test` `op1`, `op2`

Computes `op1 & op2`

- Sets the flags (`ZF`, `CF`, `OF`, ...)
- Discards the result

Examples

`cmp eax, ebx`

`cmp eax, 44BBCCDDh`

`cmp al, dh`

`cmp al, 44h`

`cmp ax, FFFFh`

`cmp [eax], 4h`

Control-Flow Instructions: conditional jumps

j<cc> address or offset

Jump to **address** if and only if a certain condition is verified

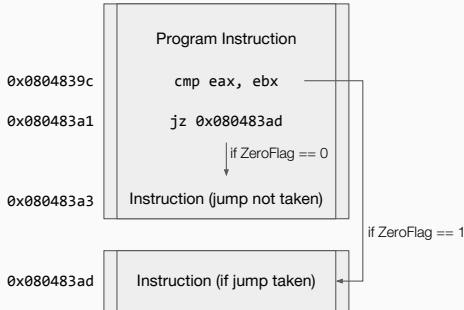
<cc>: condition

- O,NO,S,NS,E,Z,NE, ...
- based on one or more status flags of EFLAGS

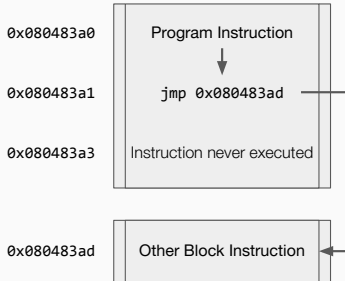
Examples:

- jz = jump if zero
- jg = jump if greater than
- jlt = jump if less than

Reference: <http://www.unixwiz.net/techtips/x86-jumps.html>



Control-Flow Instructions: unconditional jump `jmp`



- `jmp` address or offset
- Unconditional jump: just set the EIP to **address**
- Can be also *relative*: increment or decrement EIP by an offset

Exercise 1

Translate the following C code in assembly x86. Assume $EBX \leftarrow b$, $ECX \leftarrow c$. Finally, a goes in EAX .

```
if (c == 0)
    a = b;
else
    a = -b;
```


Solution

```
    mov edx, 0
    cmp ecx, edx
    jne ELSE
    mov eax, ebx
    jmp ENDIF
ELSE:
    mov eax, 0
    sub eax, ebx
ENDIF:
    nop
    ...
```

Exercise 2

Translate the following C code in assembly x86. The variable `a` goes in EAX.

```
a = 0;
for(i = 0; i < 10; i++)
    a += i;
```

Solution

```
    mov eax, 0
    mov ebx, 0
    mov ecx, 10
LOOP:
    cmp ebx, ecx
    jge END
    add eax, ebx
    inc ebx
    jmp LOOP
END:
    nop
    ...
```

A very simple example (what does it do?)

Assume that the input is in registers: ECX and EDX; output: EAX

```
    mov eax, ecx
    mov ebx, edx
    cmp ebx, 0
    jz LABEL
LOOP:
    cmp ebx, 1
    jle RET
    mul ecx
    sub ebx, 1
    jmp LOOP
LABEL:
    mov eax, 1
RET:
    ...
```

Load effective address (lea)

- lea destination, source
source: memory location
destination: register
- Like a mov, but it is storing the pointer, not the value
- It does NOT access memory

Example

Registers		Memory	
EAX	0x00000000	0x7C81776F	0x00403A40
		0x7C911000	0x00403A44
EBX	0x00403A40	0x0012C140	0x00403A48
	...	0x7FFDB000	0x00403A4C

lea eax, [ebx + 8] → EAX = 0x00403A48

N.B.: with mov eax, [ebx+8] → EAX = 0x0012C140

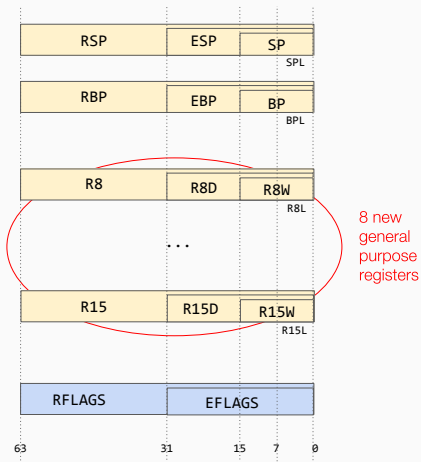
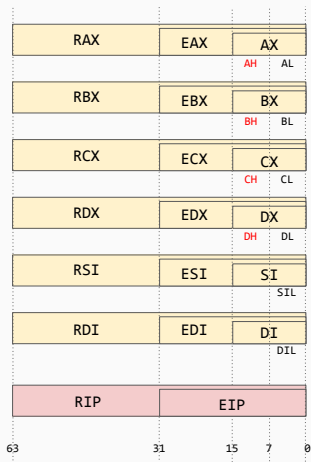
Basic Instructions: `nop`

- `nop` = **No Operation**. Just move to next instruction.
- The opcode is pretty famous and is `0x90`
- Really useful in exploitation (we will see!)

Interrupts and Syscalls

- `int value`
 - **value**: software interrupt number to generate (0-255)
 - Every OS has its set of interrupt numbers (e.g., 80h for Linux system calls)
- `syscall` used for Linux 64-bit
- `sysenter` used by Microsoft Windows

The x86_64 ISA




Endianness

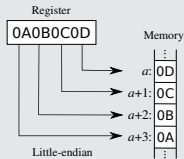
Endianness: convention that specifies in which order the bytes of a data word are lined up sequentially in memory.

Big-endian (left)

Systems in which the *most significant* byte of the word is stored in the *smallest address* given.



Little-endian



Systems in which the *least significant* byte is stored in the *smallest address*.
IA-32 is “little endian”.

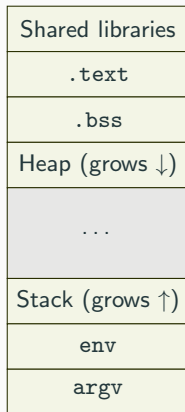
Program Layout and Functions

How an executable is mapped to memory in Linux (ELF)

Executable	Description
.plt	This section holds stubs which are responsible of external functions linking.
.text	This section holds the "text," or executable instructions, of a program.
.rodata	This section holds read-only data that contribute to the program's memory image
.data	This section holds initialized data that contribute to the program's memory image
.bss	This section holds uninitialized data that contributes to the program's memory image. By definition, the system initializes the data with zeros when the program begins to run.
.debug	This section holds information symbolic debugging.
.init	This section holds executable instructions that contribute to the process initialization code. That is, when a program starts to run, the system arranges to execute the code in this section before calling the main program entry point (called main for C programs).
.got	This section holds the global offset table.

Simplified program memory layout

Low addresses (0x80000000)



High addresses (0xbfffffff)

The Stack

- LIFO (last in first out) data structure
- Used to manage functions
 - local variables
 - return addresses
 - ...
- Handled through the register ESP (stack pointer) *IT ALWAYS POINTS AT THE TOP OF THE STACK*
- Remember: the stack grows **toward lower addresses** (downward the address space)

Stack Management Instructions: push

PUSH: ADD VALUE ON THE TOP OF THE STACK
push immediate or register

Stores the immediate or register value at the top of the stack and decrements the ESP of the operand size

Example

Initial condition: EAX = 30


push eax

is equivalent to:

sub esp, 4

mov DWORD PTR [esp], eax

Low addresses (0x80000000)



High addresses (0xbfffffff)

Stack Management Instructions: `push`

`push` immediate or register

Stores the immediate or register value at the top of the stack and decrements the ESP of the operand size

Example

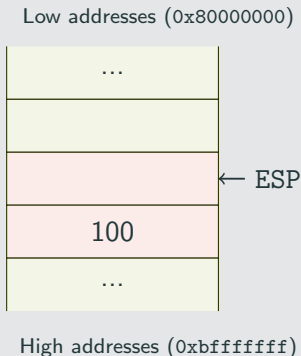
Initial condition: `EAX = 30`

`push eax`

is equivalent to:

`sub esp, 4`

`mov DWORD PTR [esp], eax`



Stack Management Instructions: `push`

`push` immediate or register

Stores the immediate or register value at the top of the stack and decrements the ESP of the operand size

Example

Final condition: `EAX = 30`


`push eax`

is equivalent to:

`sub esp, 4`

`mov DWORD PTR [esp], eax`

Low addresses (0x80000000)



High addresses (0xbfffffff)

Stack Management Instructions: pop

POP: TAKE A VALUE FROM THE STACK AND INCREASE THE
pop destination STACK POINTER VALUE MOVING IT DOWN AND
"DISCARDING THE VALUE"

Loads to the destination a word off the top of the stack, then it increases ESP of the operand's size.

Example

Initial condition: EAX = ???

pop eax


is equivalent to:

mov eax, DWORD PTR [esp]

add esp, 4

**EVERYTHING
ABOVE IS
IGNORED**

Low addresses (0x80000000)



High addresses (0xbfffffff)

Stack Management Instructions: pop

pop destination

Loads to the destination a word off the top of the stack, then it increases ESP of the operand's size.

Example

Initial condition: EAX = 30

pop eax

is equivalent to:

```
mov eax, DWORD PTR [esp]
```

```
add esp, 4
```



Stack Management Instructions: pop

pop destination

Loads to the destination a word off the top of the stack, then it increases ESP of the operand's size.

Example

Final condition: EAX = 30

pop eax

is equivalent to:

mov eax, DWORD PTR [esp]

add esp, 4



Calling a function

Instruction call:

- Push to the stack the address of the next instruction
 - Move the address of the first instruction of the callee into EIP
- //SAFE-TO-TURN ADDRESS*
↓
USEFUL FOR KNOWING WHERE TO RESUME THE FUNCTION'S EXECUTION

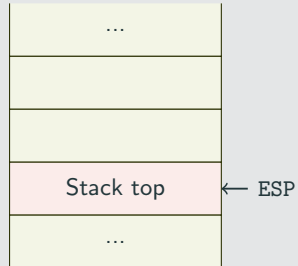
Example: Let's call func, located at 0x800bff00

Equivalent to:

- push address(of the instruction after the call!)
- jmp func

(reminder: we can't read or set EIP directly!)

Low addresses (0x80000000)



High addresses (0xbfffffff)

EIP = 0x8001020

Calling a function

Instruction `call`:

- Push to the stack the address of the next instruction
- Move the address of the first instruction of the callee into EIP

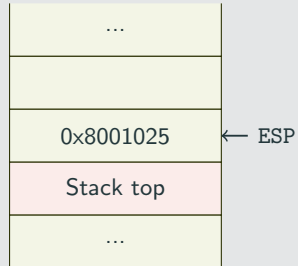
Example: Let's call `func`, located at `0x800bff00`

Equivalent to:

- `push address(of the instruction after the call!)`
- `jmp func`

(reminder: we can't read or set EIP directly!)

Low addresses (`0x80000000`)



High addresses (`0xbfffffff`)

EIP = `0x8001020`

Calling a function

Instruction `call`:

- Push to the stack the address of the next instruction
- Move the address of the first instruction of the callee into EIP

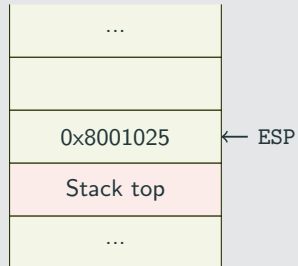
Example: Let's call `func`, located at `0x800bff00`

Equivalent to:

- `push address(of the instruction after the call!)`
- `jmp func`

(reminder: we can't read or set EIP directly!)

Low addresses (`0x80000000`)



High addresses (`0xbfffffff`)

EIP = `0x800bff00`

Returning from a function

Instruction `ret`:

- Restores the return address saved by `call` from the top of the stack

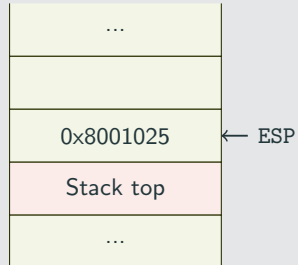
Example: let's return from `func`

Equivalent to:

- `pop eip`

(reminde: we can't read or set EIP directly!)

Low addresses (0x80000000)



High addresses (0xbfffffff)

EIP = 0x800bff00

Returning from a function

Instruction `ret`:

- Restores the return address saved by `call` from the top of the stack

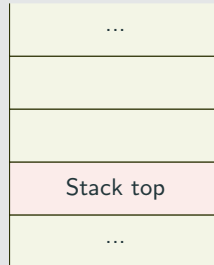
Example: let's return from `func`

Equivalent to:

- `pop eip`

(reminds: we can't read or set EIP directly!)

Low addresses (0x80000000)



← ESP

High addresses (0xbfffffff)

EIP = 0x8001025


Functions and Stack Frames

At the beginning of a function, the function itself must reserve space for its local variables.

```
void foo() {  
    int a;  
    int b;  
    int c;  
    b = 0;  
}
```

USED TO MAKE SPACE
FOR ALL THE LOCAL VARIABLES

Low addresses (0x80000000)



High addresses (0xbfffffff)

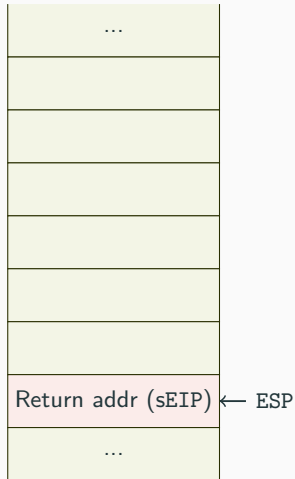
Functions and Stack Frames

At the beginning of a function, the function itself must reserve space for its local variables.

F00:

```
mov ebp, esp
sub esp, 0xc
mov [ebp - 0x8], 0x0
add esp, 0xc
ret
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

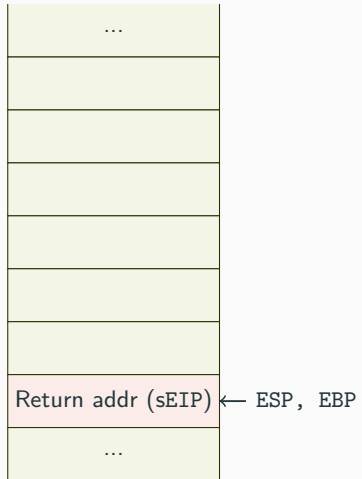
Functions and Stack Frames

At the beginning of a function, the function itself must reserve space for its local variables.

F00:

```
mov ebp, esp ←  
sub esp, 0xc  
mov [ebp - 0x8], 0x0  
add esp, 0xc  
ret
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

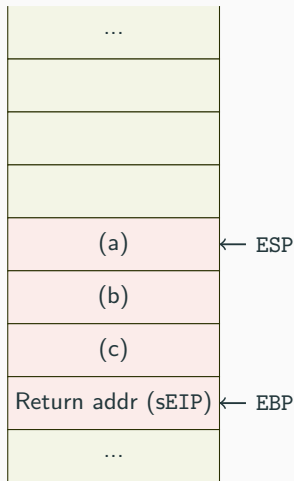
Functions and Stack Frames

At the beginning of a function, the function itself must reserve space for its local variables.

F00:

```
mov ebp, esp
sub esp, 0xc ←
mov [ebp - 0x8], 0x0
add esp, 0xc
ret
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

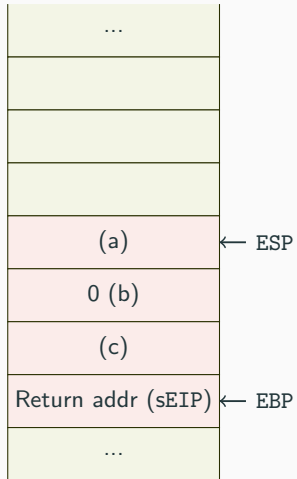
Functions and Stack Frames

At the beginning of a function, the function itself must reserve space for its local variables.

F00:

```
mov ebp, esp
sub esp, 0xc
mov [ebp - 0x8], 0x0 ←
add esp, 0xc
ret
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

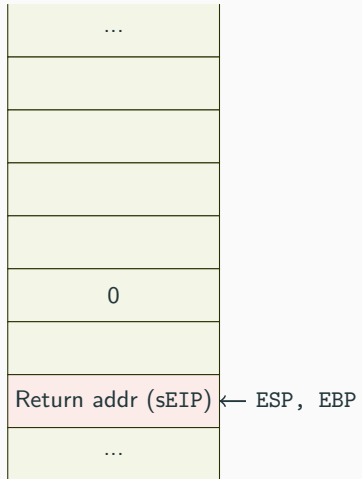
Functions and Stack Frames

At the beginning of a function, the function itself must reserve space for its local variables.

F00:

```
mov ebp, esp
sub esp, 0xc
mov [ebp - 0x8], 0x0
add esp, 0xc ←
ret
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

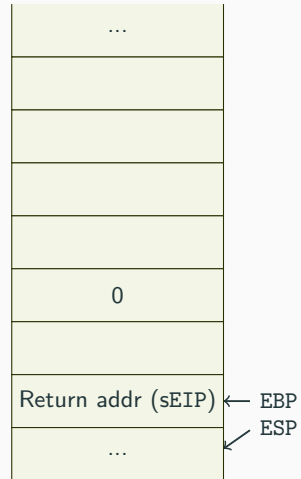
Functions and Stack Frames

At the beginning of a function, the function itself must reserve space for its local variables.

F00:

```
mov ebp, esp
sub esp, 0xc
mov [ebp - 0x8], 0x0
add esp, 0xc
ret ←
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

Functions and Stack Frames

That works!!! But what if foo calls bar.

```
void foo() {  
    int a;  
    int b;  
    int c;  
  
    bar();  
    b = 0;  
}  
  
void bar() {  
    int d;  
  
    d = 1;  
}
```


Functions and Stack Frames


F00:

```
mov ebp, esp
sub esp, 0xc
call BAR
mov [ebp - 0x8], 0x0
add esp, 0xc
ret
```

BAR:

```
mov ebp, esp
sub esp, 0x4
mov [ebp - 0x4], 0x1
add esp, 0x4
ret
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

Functions and Stack Frames

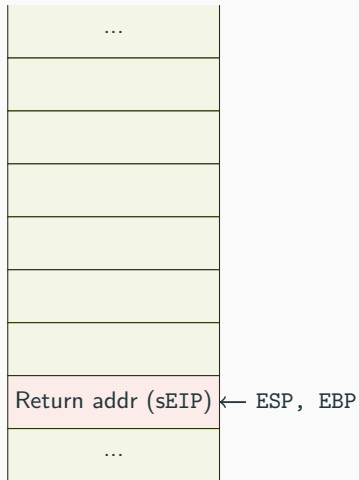
F00:

```
mov ebp, esp ←  
sub esp, 0xc  
call BAR  
mov [ebp - 0x8], 0x0  
add esp, 0xc  
ret
```

BAR:

```
mov ebp, esp  
sub esp, 0x4  
mov [ebp - 0x4], 0x1  
add esp, 0x4  
ret
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

Functions and Stack Frames

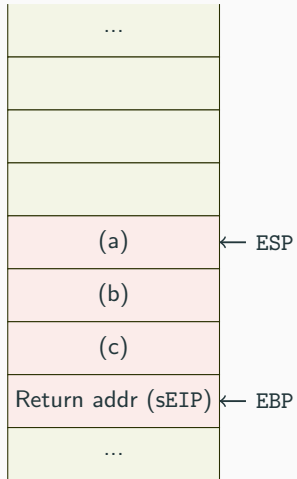
F00:

```
mov ebp, esp
sub esp, 0xc ←
call BAR
mov [ebp - 0x8], 0x0
add esp, 0xc
ret
```

BAR:

```
mov ebp, esp
sub esp, 0x4
mov [ebp - 0x4], 0x1
add esp, 0x4
ret
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

Functions and Stack Frames

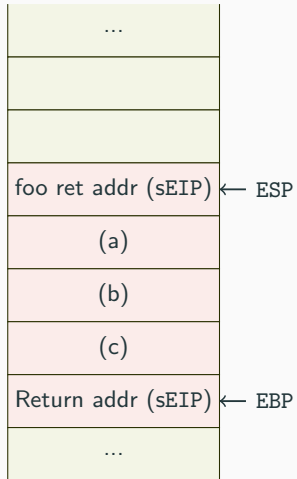
F00:

```
mov ebp, esp
sub esp, 0xc
call BAR ←
mov [ebp - 0x8], 0x0
add esp, 0xc
ret
```

BAR:

```
mov ebp, esp
sub esp, 0x4
mov [ebp - 0x4], 0x1
add esp, 0x4
ret
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

Functions and Stack Frames

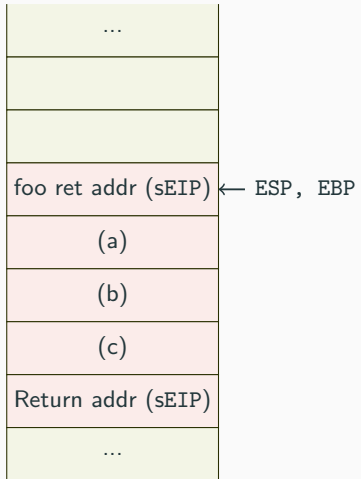
F00:

```
mov ebp, esp
sub esp, 0xc
call BAR
mov [ebp - 0x8], 0x0
add esp, 0xc
ret
```

BAR:

```
mov ebp, esp ←
sub esp, 0x4
mov [ebp - 0x4], 0x1
add esp, 0x4
ret
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

Functions and Stack Frames

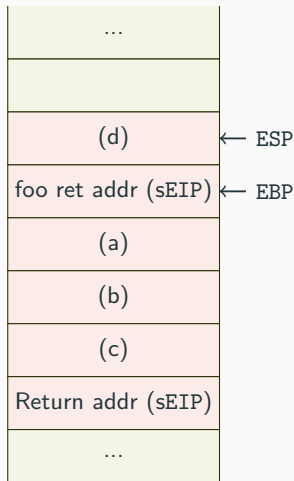
F00:

```
mov ebp, esp
sub esp, 0xc
call BAR
mov [ebp - 0x8], 0x0
add esp, 0xc
ret
```

BAR:

```
mov ebp, esp
sub esp, 0x4 ←
mov [ebp - 0x4], 0x1
add esp, 0x4
ret
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

Functions and Stack Frames

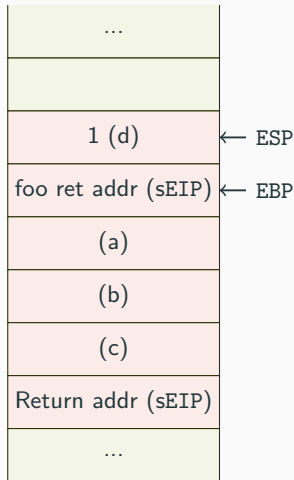
F00:

```
mov ebp, esp
sub esp, 0xc
call BAR
mov [ebp - 0x8], 0x0
add esp, 0xc
ret
```

BAR:

```
mov ebp, esp
sub esp, 0x4
mov [ebp - 0x4], 0x1 ←
add esp, 0x4
ret
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

Functions and Stack Frames

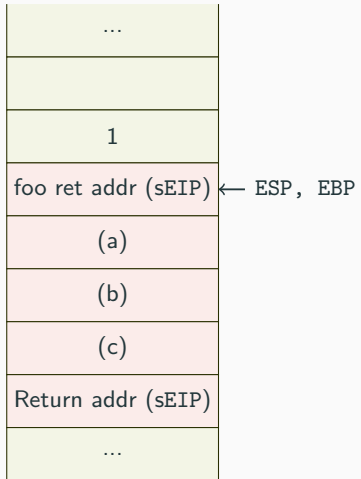
F00:

```
mov ebp, esp
sub esp, 0xc
call BAR
mov [ebp - 0x8], 0x0
add esp, 0xc
ret
```

BAR:

```
mov ebp, esp
sub esp, 0x4
mov [ebp - 0x4], 0x1
add esp, 0x4 ←
ret
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

Functions and Stack Frames

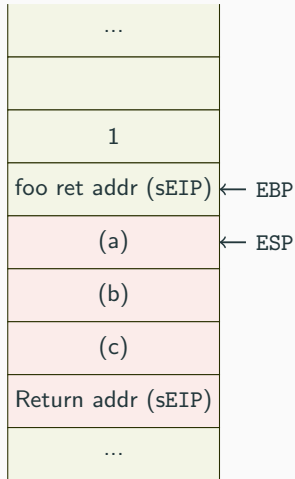
F00:

```
mov ebp, esp
sub esp, 0xc
call BAR
mov [ebp - 0x8], 0x0
add esp, 0xc
ret
```

BAR:

```
mov ebp, esp
sub esp, 0x4
mov [ebp - 0x4], 0x1
add esp, 0x4
ret ←
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

Functions and Stack Frames

Wrong memory access!!! EBP changed and we lost the old value!

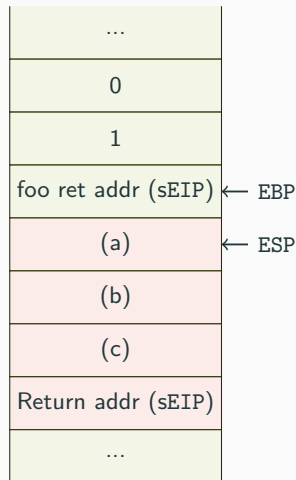
F00:

```
mov ebp, esp
sub esp, 0xc
call BAR
mov [ebp - 0x8], 0x0 ←
add esp, 0xc
ret
```

BAR:

```
mov ebp, esp
sub esp, 0x4
mov [ebp - 0x4], 0x1
add esp, 0x4
ret
```

Low addresses (0x80000000)



High addresses (0xbfffffff)

Functions and Stack Frames

Redacted!

F00:

```
push ebp
mov ebp, esp
sub esp, 0xc
call BAR
mov [ebp - 0x8], 0x0
leave → USED TO STORE THE OLD VALUE OF EBP
ret
```

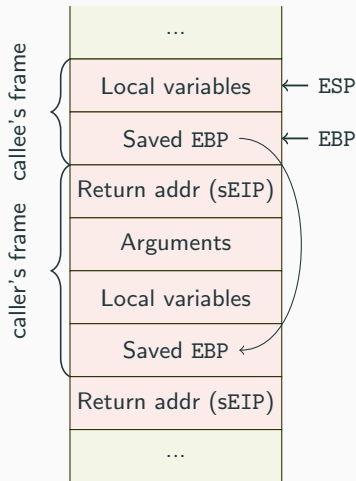
BAR:

```
push ebp
mov ebp, esp
sub esp, 0x4
mov [ebp - 0x4], 0x1
leave
ret
```

Functions and Stack Frames

- Stack frame = stack area allocated to a function
- EBP register: pointer to the beginning (base) of a function's frame
- At the beginning of a function:
 - Save EBP to stack
 - Set EBP to the address of the function's frame

Low addresses (0x80000000)



High addresses (0xbfffffff)


Entering a function

Example: We've just called `func`, located at `0x800bff00`

Setup the stack frame

- `push ebp`
- `mov ebp, esp`

Low addresses (`0x80000000`)



High addresses (`0xbfffffff`)

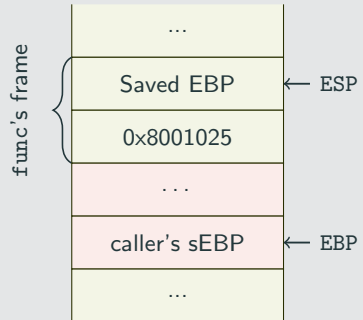
Entering a function

Example: We've just called `func`, located at `0x800bff00`

Setup the stack frame

- `push ebp`
- `mov ebp, esp`

Low addresses (`0x80000000`)



High addresses (`0xbfffffff`)

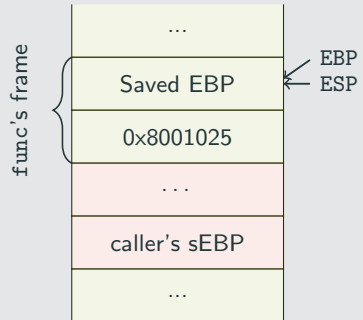
Entering a function

Example: We've just called `func`, located at `0x800bff00`

Setup the stack frame

- `push ebp`
- `mov ebp, esp`

Low addresses (`0x80000000`)



High addresses (`0xbfffffff`)

Leaving a function

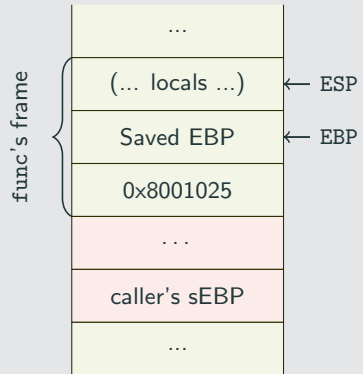
Instruction `leave`: restores the caller's base pointer

Example: We're about to return from `func`

Equivalent to:

- `mov esp, ebp`
- `pop ebp`

Low addresses (0x80000000)



High addresses (0xbfffffff)

Leaving a function

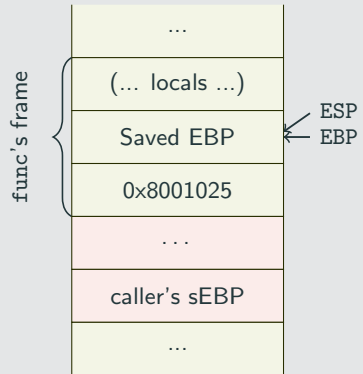
Instruction `leave`: restores the caller's base pointer

Example: We're about to return from `func`

Equivalent to:

- `mov esp, ebp`
- `pop ebp`

Low addresses (0x80000000)



High addresses (0xbfffffff)

Leaving a function

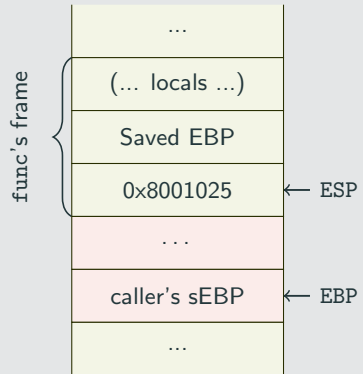
Instruction `leave`: restores the caller's base pointer

Example: We're about to return from `func`

Equivalent to:

- `mov esp, ebp`
- `pop ebp`

Low addresses (0x80000000)



High addresses (0xbfffffff)

Calling Conventions

DIFFERENT CONVENTIONS BASED
ON SPECIFIC NEEDS

- Defines
 - how to pass parameters (stack, registers or both, and who is responsible to clean them up)
 - how to return values
 - caller-saved or callee-saved registers
- The high-level language, the compiler, the OS, and the target architecture all together “implement” and “agree upon” a certain calling convention
 - it's part of the ABI, the Application Binary Interface

Calling Conventions: cdecl (C declaration)

- Default calling convention used by most x86 C compilers
 - Can be forced with the modifier `_cdecl`
- Arguments: passed **through the stack**, right to left order
- Cleanup: the **caller removes** the parameters from the stack *after* the called function completes
- Return: register EAX
- Caller-saved registers: EAX, ECX, EDX (other are callee-saved)

cdecl: Example

```
void demo_cdecl(int a, int b, int c, int z);
```

```
//...
```

```
demo_cdecl(1, 2, 3, 4); //calling
```

```
; ...
```

```
push 4 ; push last parameter value
```

```
push 3 ; push third parameter value
```

```
push 2 ; ...
```

```
push 1
```

```
call demo_cdecl ; call the subroutine
```

```
add esp, 16 ; clean up the stack
```

Calling Conventions: `fastcall` (C declaration)

- Default calling convention used by most x86_64 C compilers
 - Can be forced with the modifier `_fastcall`
- Parameters passed **in registers**: `rdi`, `rsi`, `rdx`, `rcx`, `r8`, `r9`, subsequent ones on the stack (reverse order, caller cleanup)
- Callee-saved registers: `rbx`, `rsp`, `rbp`, `r12`, `r13`, `r14`, and `r15`
- Caller-saved registers (scratch): `rax`, `rdi`, `rsi`, `rdx`, `rcx`, `r8`, `r9`, `r10`, `r11`
- Return value: `rax` (if 128-bit: `rax` and `rdx`)

Linux x86-64 calling convention: Example

main:

```
push rbp
mov rbp, rsp
sub rsp, 16
mov DWORD PTR [rbp-4], edi
mov QWORD PTR [rbp-16], rsi
mov esi, 2 ; Second parameter
mov edi, 1 ; First parameter
call function
mov esi, eax ; Return value -> first param
mov edi, OFFSET FLAT:.LCO ; "The return ...
mov eax, 0
call printf
leave
ret
```

function:

```
push rbp
mov rbp, rsp
mov DWORD PTR [rbp-4], edi
mov DWORD PTR [rbp-8], esi
mov edx, DWORD PTR [rbp-4]
mov eax, DWORD PTR [rbp-8]
add eax, edx
pop rbp
ret
```

```
int function (int a, int b) {
    return a + b;
}

int main (int argc, char** argv) {
    return printf("The return value is %d\n",
        function(1,2));
}
```

Low addresses

...
2
1
Saved RBP
Return address
...

High addresses

Tooling

Shell for Dummies³

important paths	/	#root (first) directory
	~/	#your home directory
	.	#current directory
	..	#parent directory
filesystem utils	pwd	#show current directory
	cd <i>path</i>	#change directory to <i>path</i>
	ls <i>path</i>	#list files in the directory at <i>path</i>
	cp <i>path_src path_dst</i>	#copy <i>path_src</i> to <i>path_dst</i> , -r if copying directories
basic text editor	mv <i>path_src path_dst</i>	#move <i>path_src</i> to <i>path_dst</i>
	nano <i>path/file</i>	#opens/create <i>file</i> in <i>path</i> (ctrl+x to exit)
	vim <i>path/file</i>	#opens/create <i>file</i> in <i>path</i> (i to edit; esc, :wq to save and quit)
remote actions	ssh <i>user@server_addr</i>	#ssh to <i>server_addr</i> as <i>user</i>
	scp [<i>u@s:</i>] <i>p_src</i> [<i>u@s:</i>] <i>p_dst</i>	#cp to/from remote server

³cmd --help or cmd -h to get the available options

Shell for Dummies ⁴

file visualization	cat <i>file</i> less <i>file</i> hexdump -Cv <i>file</i>	#print <i>file</i> to stdout #visualize <i>file</i> better, q to quit #visualize raw bytes
redirections	command > <i>file</i> command >> <i>file</i> command < <i>file</i> cmd1 cmd2	#write stdout of command to <i>file</i> #append stdout of command to <i>file</i> #use <i>file</i> as stdin of command #stdout of cmd1 as stdin cmd2
argv from command	cmd `cmd2` cmd \$(cmd2)	#executes cmd2 first and uses the output to eval the next command
other useful commands	chmod +x <i>file</i> grep <i>expression</i> python -c 'cmd1;cmd2;...'	#give exec permission to <i>file</i> #search for <i>expr</i> in stdin #executes python commands

⁴cmd --help or cmd -h to get the available options

- `man objdump`
objdump **displays information** about one or more **object files**.
- `-x` all-headers
- `-d` disassemble
- `-M intel` intel syntax (default is AT&T)

- **What is GDB?**

GDB is GNU Project's Debugger: allows to follow, step by step, at assembler-level granularity, a running program, or what a program was doing right before it crashed.⁵

⁵<http://www.gnu.org/software/gdb/>

Start, break and navigate the execution with gdb

- Suppose you have an executable binary and want run it
 - **`gdb /path/to/executable`** loads the binary in gdb
- Now you decide to start the program with two parameters
 - **`run 1 "abc"`** passes 1 via `argv[1]` and "abc" as `argv[2]`
 - **`run 'printf "AAAAAAAAAAAAAA"'`** (with the back ticks)
we're passing the output of the print (very useful when you need to pass non printable characters such as raw bytes)
- Suppose you want to stop the execution at the address of a certain instruction
 - **`break *0xDEADBEAF`** places a break point at that address
 - **`break *main+1`** with debugging symbols this can be less painful
 - **`catch syscall`** block the execution when a syscall happens

Start, break and navigate the execution with gdb

- Now the execution stops at our break point. Here we can do several things
- Examples:
 - **ni** allows to proceed instruction per instruction
 - **next 4** moves 4 lines ahead (if you have the line-numbers information in the binary)
 - **si** step into function
 - **finish** run until the end of current function
 - **continue** runs until the next break point (if any)
- To see info about the execution state:
 - **info registers** to inspect the content of the registers
 - **info frame** to see the values of the stack frame related to the function where we are in
 - **info file** print the information about the sections of the binary

Navigate the stack

- Suppose we're stopped somewhere in the code and want to inspect the stack
- Some useful view of the stack is achievable with:
 - **x/100wx \$esp** prints 100 words of memory from the address found in the ESP to ESP+100 (x = hexadecimal formatting)
 - **x/10wo \$ebp-100** prints 10 words of memory from EBP-100 to EBP-100+10 (o = octal formatting)
 - **x/s \$eax** prints the elements pointed by EAX (s = string formatting)
- Do you have debug symbols? (i.e., gcc -ggdb)
 - **print args** prints info about the main parameters
 - **print a** prints the content of variable 'a'
 - **print *b** prints the value pointed by 'b'

Our friend gdb

- **The '~/.gdbinit' file**

Gdb is a command line tool and it supports the configuration script as almost all the *nix software.

Some options that you may want to tune are:

- **set history save on**

To have the latest commands always available also when we re-open gdb

- **set follow-fork-mode child**

Allows you, if the process spawns children, to follow them and not only wait their end.

- **set disassembly-flavor [intel | att]**

This option sets in which predefined syntax your disassembled will be showed up. The default one is at&t

- Highly recommended to install pwndbg
<https://github.com/pwndbg/pwndbg>

GDB - How to Survive ⁶

start

`gdb -q program`

#starts gdb silently for program

disassemble

`set disassembly-flavor intel`
`disass *address (or f-name)`

#sets intel syntax
#disassemble from given address

run program

`run (r)`
`start`
`run arg1`
`run <<< arg1`

#runs the program
#runs the program and imm. stops
#runs program with *arg1* in argv
#runs program with *arg1* in stdin

memory layout

`vmmap`

#show memory layout

⁶CTRL + C to Break and Debug

GDB - How to Survive ⁷

execution	stepi (si) nexti (ni) finish (f) continue (c)	#exec next inst - enters a function #exec next inst - skips the function #exec till next return statement #continue exec till next break/ watch
breakpoints	b *address b *address if \$reg==val del br_num	#set software breakp at address #set conditional breakp #remove breakpoint br_num
watchpoints	w *address rw *address	#set watch for write at address #set watch for read at address
examine	x/numF *address search string p symbol	#show num data of type F (useful Fs are bx, wx, gx, c, s, i) #search for string in memory #print address of symbol

⁷CTRL + D to Exit

- Intercepts and records system calls and signals
- Dumps to standard error name, argument and return value of each system call

Useful options

- `-p <pid>` attach to existing process
- `-f` trace child process
- `-o <filename>` output to file
- `-e <expr>` modifies which events to trace (see manpage)

- Intercepts and records dynamic library calls
- Similar to strace, but at a different layer

Questions?