

## Chapter 3

# Processes

---

October, 1st

# Introduction

- The notion of process is fundamental
- Originally introduced in operating systems
- “A process is a program in execution”
- Processes and threads
- Processes vs processors
- Related concepts: virtualization

# Processes and Threads

---

## Introduction to processes

- Virtual processors created by the OS
- Process table (per virtual processor)
  - Process context: CPU register values, memory maps, open files, ...
- Process is a program executed by a virtual processor
- Concurrent processes are *isolated* from each other
  - Special hardware enabling protection
  - ✓ Sharing of CPU and hardware resources (RAM) is transparent
  - ✗ High performance price

## Introduction to processes

The price of process concurrency

- Creating a new process
  - OS must create an independent address space
- Switching between two processes
  - Save CPU context (registers, program counter, stack pointer, ...)
  - Modify registers of the memory management unit (MMU)
  - Invalidate translation lookaside buffer (TLB)
  - Maybe swap processes between main memory and disk

## Introduction to threads

The consequences of thread concurrency:

- Less transparency between threads (than between processes)
- Thread context
  - CPU context and some other management data (**not memory**)
  - e.g., mutex variables
- Implications
  - ✓ multithreaded applications are highly efficient
  - ✗ more difficult to program *correctly* and *securely*


## Threads usage in “non-distributed systems”

- ✓ Better performance from non-blocking requests
  - Reading from I/O and computing
- ✓ Better performance in parallel systems
  - Each thread is assigned to a processor
- ✓ Some applications are simpler to program
  - e.g., word processor (Microsoft Word):
    - thread for handling input
    - thread for spellchecking
    - thread for document layout
  - e.g., integrated development environment (VSCode):
    - thread for handling input
    - thread for analyzing the source code
    - thread for highlighting
    - thread for executing tests

# Threads

## Threads usage in “non-distributed systems”

- Threads reduce the number of context switches in complex applications using IPC (e.g., pipes, message queues)
- Multithread communication by means of shared memory



Context switching as the result of IPC




## Thread implementation

- User-level threads: threads exist in user space only
  - ✓ Creating and destroying threads is cheap
  - ✓ Fast thread switching (only CPU registers)
  - ✗ **Process blocks on blocking system calls**
  - e.g., GNU Portable Threads
- Operating system threads
  - More similar to processes
  - Native POSIX Thread Library (NPTL) for Linux (POSIX Threads)

# Threads

## Thread implementation

- Combining user-level and kernel-level threads
  - Kernel-level threads execute user-level threads



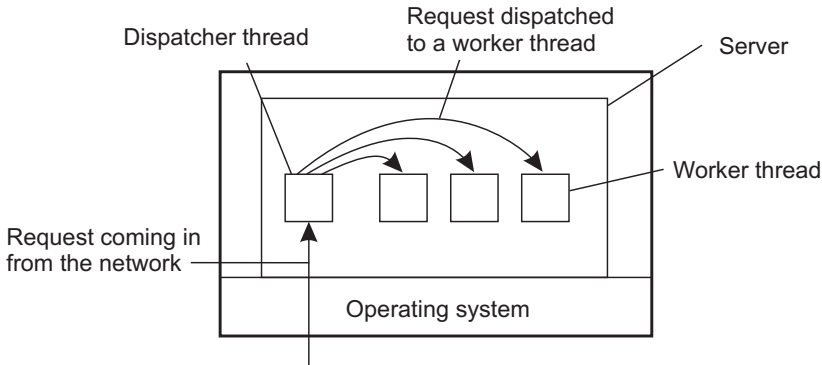
## Threads in distributed systems

- Main advantages:
  - ✓ not blocking on *communication* calls
  - ✓ exploiting parallelism in multi-processor systems
- Multithreaded clients
  - hide large latencies in wide-area communication
  - e.g., web browser:
    - display layout before fetching media
    - fetch multiple objects at the same time
- Multithreaded servers...

# Threads

## Multithreaded servers

- Simplifies server design
- Easier to develop high performance servers (concurrency)



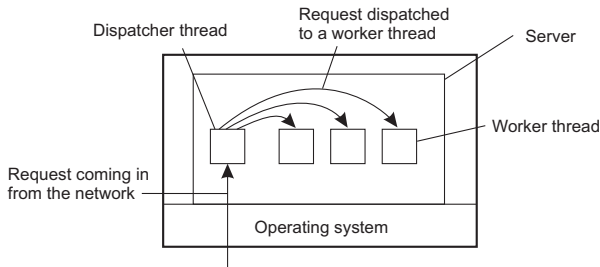
A multithreaded server organized in a dispatcher/worker model

# Threads

## Multithreaded servers

### Dispatcher/worker model workflow

- **Dispatcher** (one thread) reads incoming requests for the service, examines the request and chooses an idle worker thread
- **Worker** threads execute requests and reply to the clients
- Worker thread blocks  $\Rightarrow$  others workers continue to execute
- How to implement the server without threads?



## Multithreaded servers

How to implement the server without threads?

- Single-thread model
  - One thread receives requests, executes them and replies
  - Sequential processing
- Finite-state machine model (simplified by async programming)
  - One thread executes requests and replies, if execution is quick
  - Blocking system calls are replaced by non-blocking calls
  - Multiple requests are handled simultaneously

# Threads

## Multithreaded servers

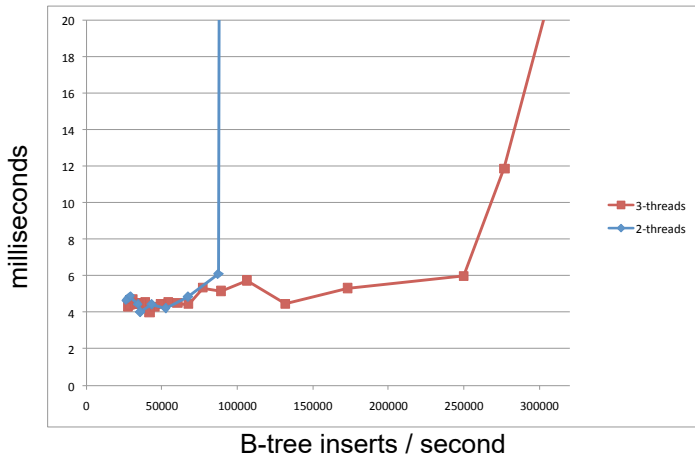
How to implement the server without threads?

- Single-thread model
  - One thread receives requests, executes them and replies
  - Sequential processing
- Finite-state machine model (simplified by async programming)
  - One thread executes requests and replies, if execution is quick
  - Blocking system calls are replaced by non-blocking calls
  - Multiple requests are handled simultaneously

Model	Characteristics	Performance	Simplicity
Threads	Parallelism, blocking system calls	✓	
Single-threaded process	No parallelism, blocking system calls		✓
Finite-state machine	Parallelism, nonblocking system calls	✓	✓

# Threads

## Multithreaded servers






# Virtualization

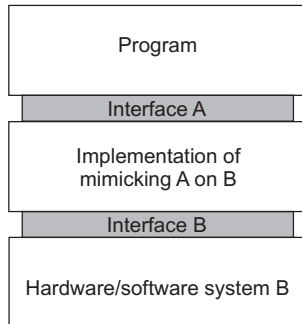
---

# The role of virtualization in distributed systems

- Resource virtualization (not only CPU)
- Virtualization extends or replaces existing interfaces  
e.g., to mimic the behavior of other systems



(a)



(b)

- (a) General organization between a program, interface, and system.  
(b) General organization of virtualizing system A on top of system B.


# The role of virtualization in distributed systems

## Reasons for virtualization

- Allow legacy software to run on expensive mainframe hardware (e.g., IBM 370 in the 1970s)
- Porting legacy interfaces to new platforms  
e.g., compatibility in Windows
- Make hardware platforms *appear* more homogeneous
- Security (by isolating systems running on the same servers)

# Architectures of virtual machines

## Interfaces offered by computer systems



Various interfaces offered by computer systems

# Architectures of virtual machines


## Interfaces offered by computer systems

- hardware/software: CPU instructions invokable by any program
- hardware/software: CPU instructions invokable only by privileged programs (e.g., operating systems)
- operating system/software: system calls
- library/software: functions generally forming what is known as an application programming interface (API)

# How to implement virtualization

## Run-time system with abstract instruction set

- Process virtual machine
- Instructions are interpreted (e.g., Java) or emulated (e.g., Windows on Unix)




A process virtual machine

# How to implement virtualization

## Layer shielding original hardware

- Virtual machine monitor (VMM), (e.g., VMware)
- Offer complete set of instructions of same or different hardware




A virtual machine monitor

# How to implement virtualization

## A hosted virtual machine monitor

- Virtual machine monitor (VMM), (e.g., VirtualBox)
- Uses interfaces provided by the host operating system

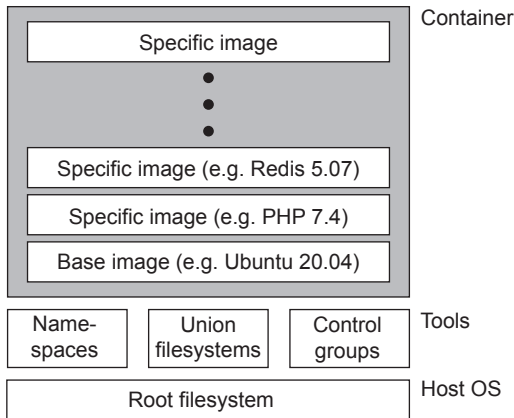


A (hosted) virtual machine monitor



# Containers

- applications stable in operating system and CPU architecture
- need *isolated software environment*



Structure of a typical running container

# **Clients and Servers**


---

## Networked user interfaces

- Client machines
  - Provide the means for users to *interact with remote servers*
- Application-specific protocol
  - For each remote service, the client machine has a separate counterpart that contacts the service over the network
- Application-independent protocol
  - Direct access to remote services through a convenient interface
  - Client machine used only as a terminal, no local state

## Application-specific protocol


- Some application state resides on the client machine
- e.g., calendar app synchronizes with a remote agenda



A networked application with its own protocol

## Application-independent protocol

- Thin-client approach
- Everything processed and stored on the server




A general solution to allow access to remote applications

## Client-side solution for distributed transparency

- In many applications, clients are more than user interfaces
  - e.g., embedded devices (ATMs, TV set-top boxes, ...)
  - e.g., web applications
  - e.g., smartphones
- Mechanism for achieving distribution transparency
  - e.g., users should not be able to tell that service is remote
- Client-side solution for replication transparency
  - i.e., users cannot tell that multiple replicas of a server exist

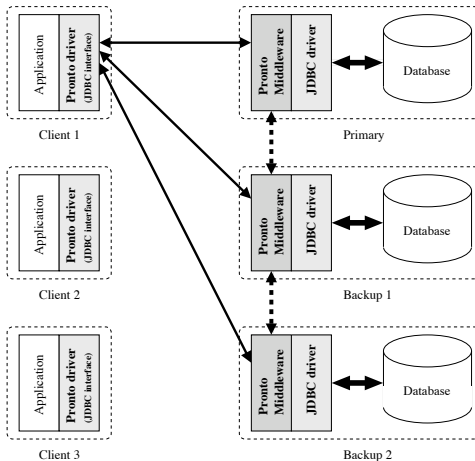
# Clients

## Client-side solution for replication transparency



Transparent replication of a server using a client-side solution

## Example: Pronto architecture





## General design issues

- Typical server organization
  - Waits for an incoming request from a client
  - Takes care of the request (e.g., executes the request)
  - Sends the response back to the client

## General design issues

- Typical server organization
  - Waits for an incoming request from a client
  - Takes care of the request (e.g., executes the request)
  - Sends the response back to the client
- Iterative versus concurrent servers

## General design issues


- Typical server organization
  - Waits for an incoming request from a client
  - Takes care of the request (e.g., executes the request)
  - Sends the response back to the client
- Iterative versus concurrent servers
  - Iterative: one request handled at a time
  - Concurrent: a separate thread or process executes the request
    - e.g., by forking a separate process

## General design issues

- Typical server organization
  - Waits for an incoming request from a client
  - Takes care of the request (e.g., executes the request)
  - Sends the response back to the client
- Iterative versus concurrent servers
  - Iterative: one request handled at a time
  - Concurrent: a separate thread or process executes the request
    - e.g., by forking a separate process
- How clients contact a server?
  - Requests are sent to an end point (a port) at the server machine
  - Finding the end point
    - “Well-known” ports (e.g, FTP, TCP port 21, HTTP, TCP port 80)
    - Special service provides port number (daemon)

# Servers

## Finding the end point: through a daemon




Client-to-server binding using a daemon

# Servers

## Finding the end point: superservers

- Listen on many ports and fork a process to execute service
- e.g., inetd in Unix



Client-to-server binding using a superserver

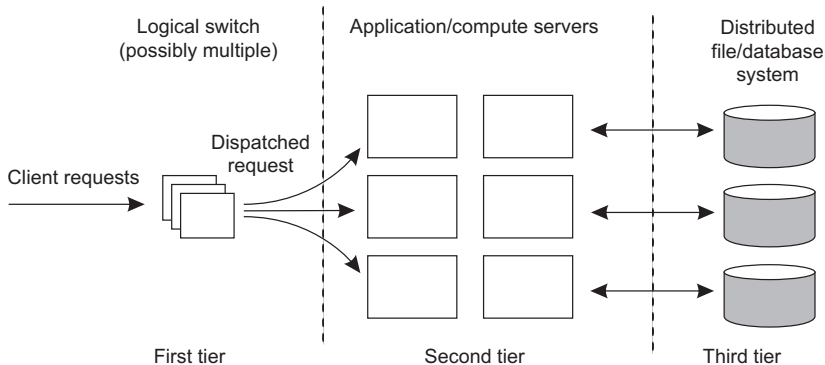
## State management

- Stateless server
  - does not keep information about the state of clients  
e.g., web server: executes request and forgets about client
  - sometimes server keeps information that can be recreated  
cf. RESTful APIs
- Stateful server
  - Maintains persistent information about clients (e.g., file server)
  - ✓ Better performance (e.g., data pre-fetching)
  - ✗ Problematic in case of failures

# Servers

## Server clusters

Collection of machines connected through a network



The general organization of a three-tiered server cluster



## Server clusters

General organization:


- First tier: logical switch
  - Receives client requests and routes them to servers
- Second tier: application servers
  - High-end servers, if service is computationally intensive
  - Low-end servers if storage is the bottleneck
- Third tier: data-processing servers
  - e.g., file/database systems

# Servers

## Server clusters

Hiding the existence of multiple servers:

- Unique access point (logical switch)
- Switch accepts TCP connection and hands it off to server
- Server replies using *switch's* IP address
- How to do load balancing?



The principle of TCP handoff

## PlanetLab

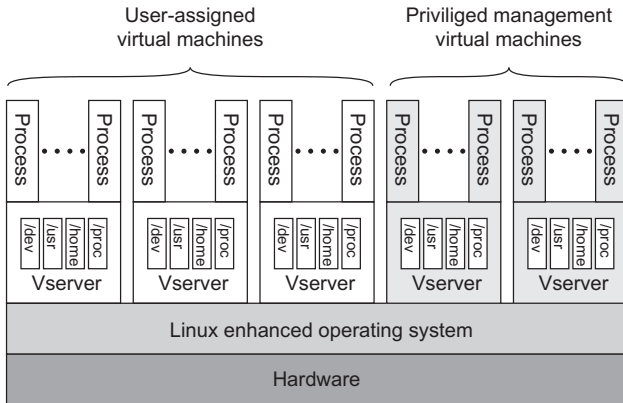
- Collaborative distributed system
- Different organizations donate servers to the main system
- Testbed for geographically distributed systems

Basic entities:

- **Virtual machine monitor (VMM)**: enhanced Linux OS
- **Vserver**: isolated environment (like a container)
- **Slice**: set of Vservers, each running on a different node

# Servers

## PlanetLab



The basic organization of a PlanetLab node

## PlanetLab

Management issues:

- Nodes belong to different organizations. Organization should be allowed to specify who can run applications on their nodes.
- Monitoring tools assume a very specific combination of hardware and software. All tailored to be used within a single organization.
- Programs from different slices but running on the same node should not interfere with each other.